

Applied Analytics: Frameworks and Methods 2

Time Series Analysis

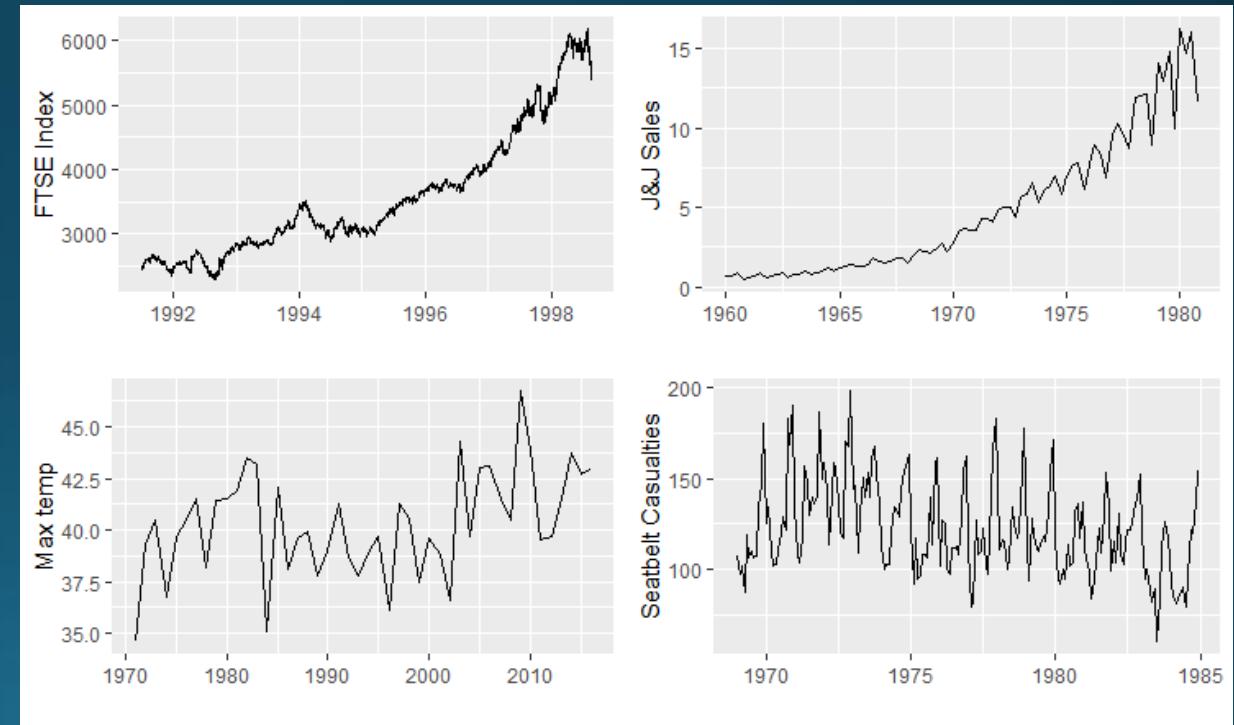
Outline

- Examine time-series data
- Explore and visualize time-series data
- Review basic/simple forecasting methods
- Explore exponential smoothing
- Examine ARIMA models
- Construct forecasts using simple methods, exponential smoothing, and ARIMA.

Time Series Data: Examine, Explore, and Visualize

Time Series

- A series of data in order of time
- A lot of data around us is measured over a span of time
 - Stock prices
 - Federal funds rate
 - Sales
 - Weather
- Time-series data is indexed on time.



Simplifying Assumptions

We make some simplifying assumptions about time-series

- Consecutive observations are equally spaced
- Apply a discrete-time index (even though this may only hold approximately)
 - E.g., daily stock returns are only available for weekdays, not weekends and holidays.

- Indexed by time/date
- We use special data structures
 - ts
 - xts (stands for eXtensible time series): matrix with a time index
 - zoo
- Use of these data structures offer access to useful functions for manipulating, analyzing and visualizing time-series data.

Time Series Patterns

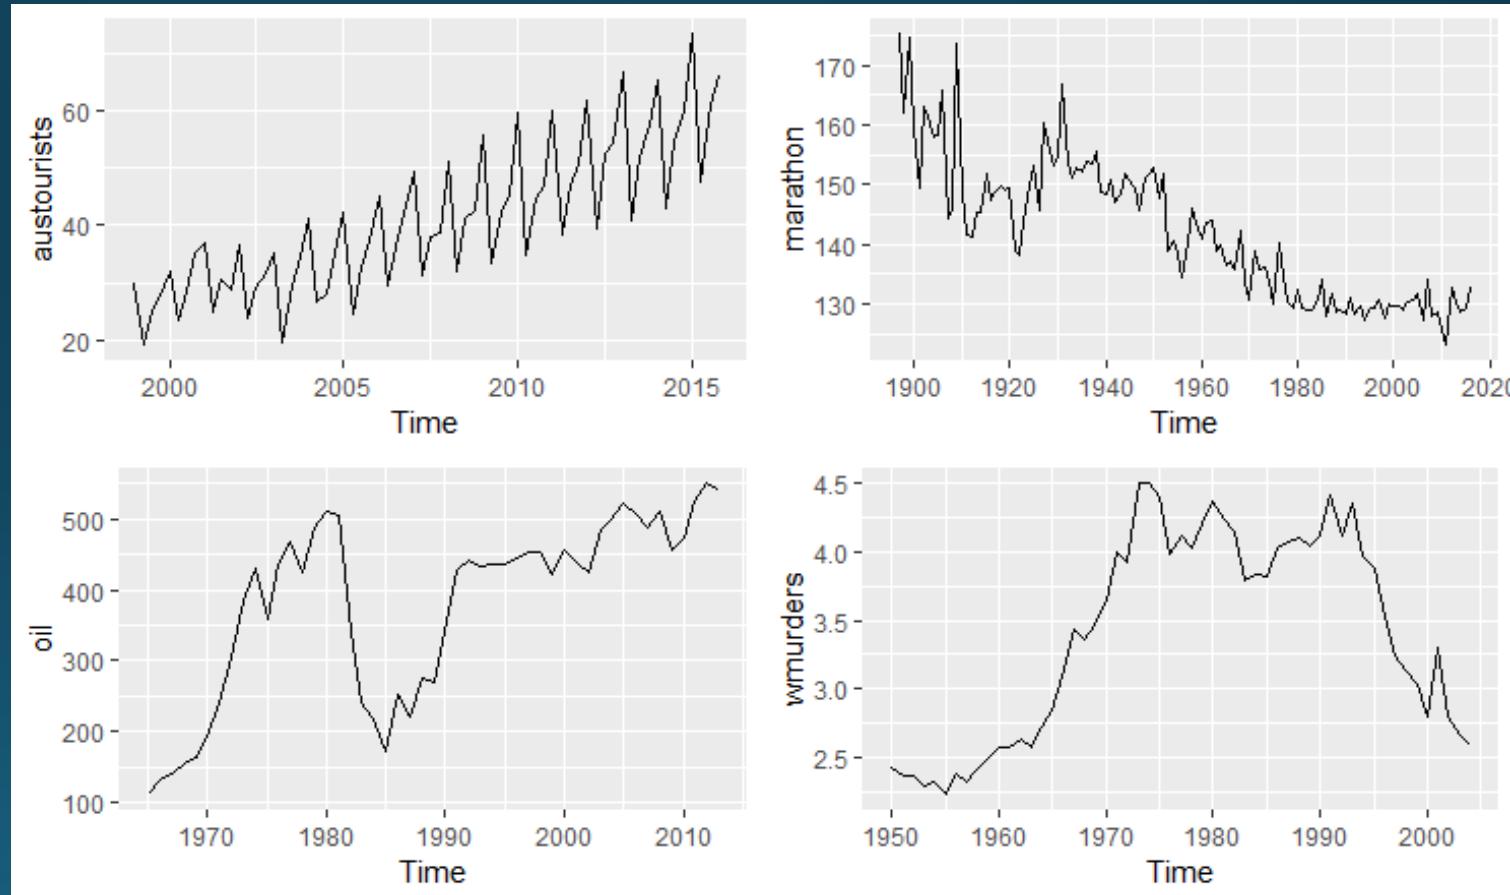
- A common observation across time-series is that prior observations influence later observations.
- The influence may be from the previous observation (trend), an observation that is a fixed number of periods ago (seasonality), or an observation that is not a fixed number of periods ago (cycle).
- Influence of prior observations can be examined by looking at lagged correlations or autocorrelations.

- Trend
 - A long term increase or decrease in the data. Could be linear, non-linear, and may go from an increasing trend to a decreasing trend
- Seasonal
 - Pattern repeats itself every season. Here season may be period of the year or day of the week. Seasonality is always of a fixed or known frequency
- Cyclic
 - Data exhibit rises or falls that are not of a fixed frequency
- On the other hand, a time-series data that does not exhibit any pattern is a Stationary Process. White Noise is the simplest Stationary Process.

Time Series Patterns

- Which is the *dominant* time series pattern in these charts?

- Trend
- Seasonal
- Cyclic



Time Series Components

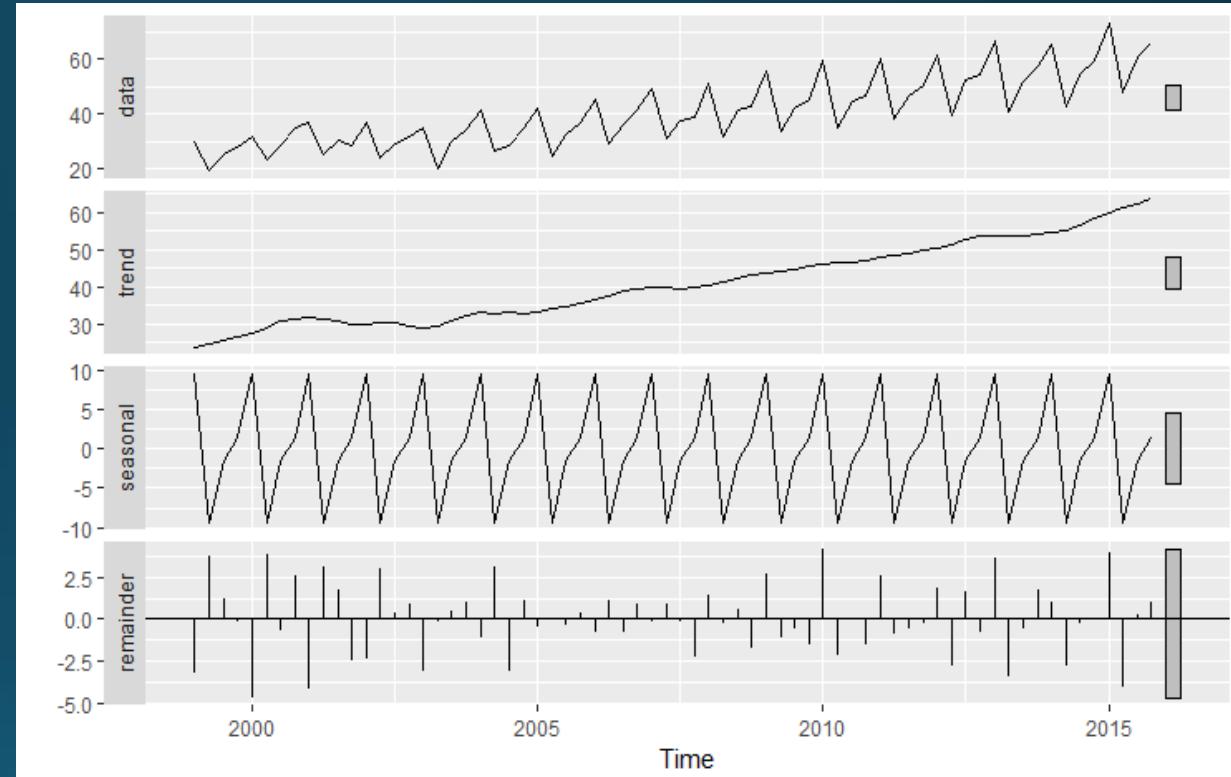
- Time-series may have multiple components in varying strengths. To better understand the series it is useful to divide it into its components: :

- Trend components
- Seasonal component
- Remainder

- Assume an additive decomposition, then we can write

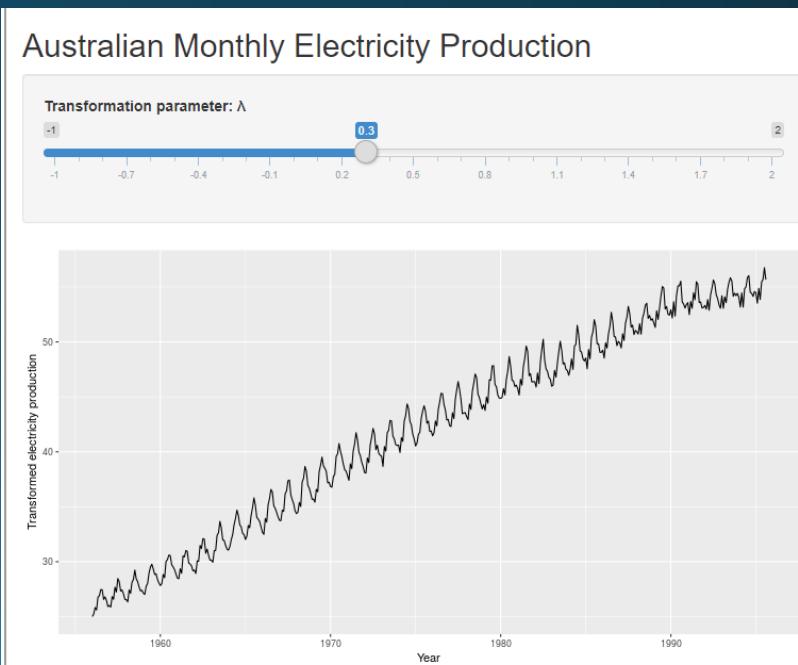
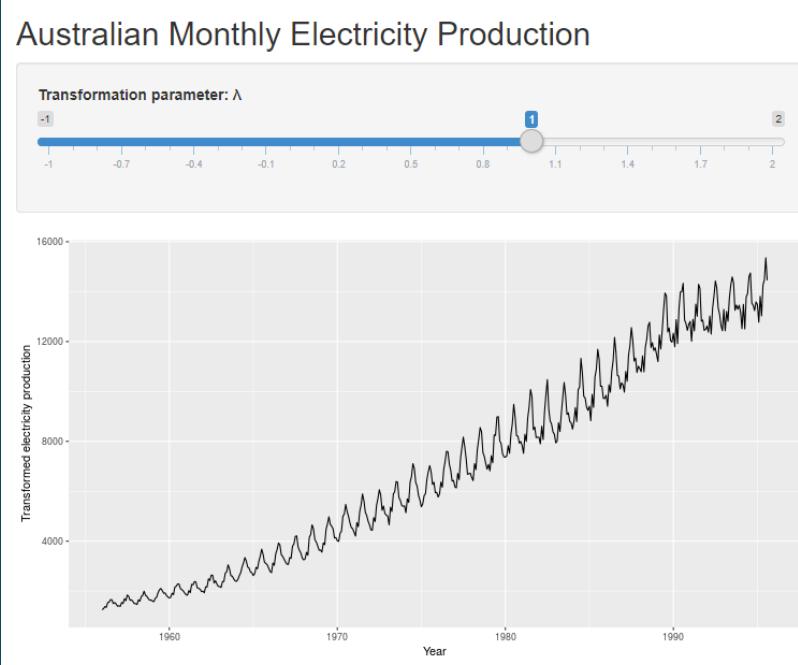
$$y_t = T_t + S_t + R_t$$

- Seasonality and Trend Decomposition using Loess (STL) is a versatile and robust method for decomposing a time series
- Decomposition can be used to measure the strength of trend and seasonality



Transformations and Adjustments

- Calendar adjustments
 - Sometimes variation seen in seasonal data may be due to a simple calendar effect such as differences in number of business days in a month.
- Population adjustments
 - A simple adjustment for population is to look at per-capita measures
- Inflation adjustments
 - Data involving money could be adjusted for changes in the time value of money
 - Adjust for inflation, CPI, etc.
- Mathematical transformations
 - When data shows change in variance with time, the variance can be stabilized by a suitable transformation. A common transformation is to use log. A family of transformations that include both log and power transformations are Box-Cox transformations which depend on the parameter lambda.
 - A good value of λ is one which makes the size of the seasonal variation about the same across the whole series



Time Series Regression

- Forecast time series of interest with another time-series
- Important to check for assumptions
 - Model is a reasonable approximation of reality
 - The errors must
 - have a mean of 0
 - are not auto-correlated
 - unrelated to predictors
 - Useful properties of errors:
 - be normally distributed
 - have constant variance (no heteroscedasticity)

- Predictions with time-series data are called forecasts
- Of course, not all forecasts are derived from time-series data. Judgmental forecasts are often used when
 - There is no historical data
 - Data is incomplete or available after a delay
- Three approaches to incorporating judgmental forecasts
 - Rely entirely on judgment when there is no available data
 - When data are available, generate statistical forecasts and adjust using judgment
 - When data are available, generate statistical and judgmental forecasts independently and combine the two.

- Delphi method
 - the aim of the Delphi method is to construct consensus forecasts from a group of experts in a structured iterative manner
- Forecasting by analogy
 - example is the pricing of a house through an appraisal process
- New product forecasting
 - Sales force composite, executive opinion, customer intentions

Basic Forecasting Methods

- Average Method
 - A very simple prediction, often the baseline in linear regression, is to use the average.
- Naïve Method
 - Future will be the same as the last observation
 - Since this is the best prediction for a random walk, these are also called random walk forecasts.
- Seasonal Naïve Method
 - Forecast is equal to the last observed value from the same season.
- Drift Method
 - Allow forecasts to increase or decrease over time, where the amount of change over time (called the drift) is set to be average change seen in historical data.

Exponential Smoothing Models

Exponential Smoothing

- Forecasts are weighted averages of past observations with the weights decaying exponentially such that recent observations get weighted more than distant observations.

- Simple exponential smoothing
 - Forecasts are calculated using weighted averages, where the weights decrease exponentially.
Most recent observations get the heaviest weight.
 - Simplest of exponential smoothing methods.
 - Suitable for forecasting data with no clear trend or seasonal pattern.
- Holt's Method
 - Extends simple exponential smoothing to allow the forecasting of data with a trend.
- Holt's Method with Damping
 - Forecasts generally display a constant trend (increasing or decreasing) indefinitely into the future.
For this reason, a damping parameter is usually included.

Types

- Holt-Winter's seasonal method
 - Extends Holt's method to capture seasonality. Two types:
 - Additive method is used when seasonal variations are roughly constant
 - Multiplicative method is used when seasonal variations change in proportion to the level of the series
- But, exponential smoothing methods are not limited to these. By considering variations of trend (none, additive, additive damped) and seasonal components (non, additive, multiplicative), there are nine exponential smoothing methods
- For each of these methods, errors may be additive or multiplicative.

Exponential Smoothing Models: A Taxonomy

Trend Component	Seasonal Component		
	N	A	M
	(None)	(Additive)	(Multiplicative)
N (None)	(N,N)	(N,A)	(N,M)
A (Additive)	(A,N)	(A,A)	(A,M)
A _d (Additive damped)	(A _d ,N)	(A _d ,A)	(A _d ,M)

Some of these methods we have already seen using other names:

Short hand	Method
(N,N)	Simple exponential smoothing
(A,N)	Holt's linear method
(A _d ,N)	Additive damped trend method
(A,A)	Additive Holt-Winters' method
(A,M)	Multiplicative Holt-Winters' method
(A _d ,M)	Holt-Winters' damped method

State Space Models

- Each model consists of

- a measurement equation that describes the observed data and
- some state equations that describe how the unobserved components or states (level, trend, seasonal) change over time.

- Hence, these are referred to as State Space Models

Table 7.7: State space equations for each of the models in the ETS framework.

ADDITIONAL ERROR MODELS

	Trend	Seasonal	
	N	A	M
N	$y_t = \ell_{t-1} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \alpha \varepsilon_t$	$y_t = \ell_{t-1} + s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \alpha \varepsilon_t$ $s_t = s_{t-m} + \gamma \varepsilon_t$	$y_t = \ell_{t-1} s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \alpha \varepsilon_t / s_{t-m}$ $s_t = s_{t-m} + \gamma \varepsilon_t / \ell_{t-1}$
A	$y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$ $b_t = b_{t-1} + \beta \varepsilon_t$	$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$ $b_t = b_{t-1} + \beta \varepsilon_t$ $s_t = s_{t-m} + \gamma \varepsilon_t$	$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t / s_{t-m}$ $b_t = b_{t-1} + \beta \varepsilon_t / s_{t-m}$ $s_t = s_{t-m} + \gamma \varepsilon_t / (\ell_{t-1} + b_{t-1})$
A_d	$y_t = \ell_{t-1} + \phi b_{t-1} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha \varepsilon_t$ $b_t = \phi b_{t-1} + \beta \varepsilon_t$	$y_t = \ell_{t-1} + \phi b_{t-1} + s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha \varepsilon_t$ $b_t = \phi b_{t-1} + \beta \varepsilon_t$ $s_t = s_{t-m} + \gamma \varepsilon_t$	$y_t = (\ell_{t-1} + \phi b_{t-1}) s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha \varepsilon_t / s_{t-m}$ $b_t = \phi b_{t-1} + \beta \varepsilon_t / s_{t-m}$ $s_t = s_{t-m} + \gamma \varepsilon_t / (\ell_{t-1} + \phi b_{t-1})$

MULTIPLICATIVE ERROR MODELS

	Trend	Seasonal	
	N	A	M
N	$y_t = \ell_{t-1}(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1}(1 + \alpha \varepsilon_t)$	$y_t = (\ell_{t-1} + s_{t-m})(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1} + \alpha(\ell_{t-1} + s_{t-m})\varepsilon_t$ $s_t = s_{t-m} + \gamma(\ell_{t-1} + s_{t-m})\varepsilon_t$	$y_t = \ell_{t-1} s_{t-m}(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1}(1 + \alpha \varepsilon_t)$ $s_t = s_{t-m}(1 + \gamma \varepsilon_t)$
A	$y_t = (\ell_{t-1} + b_{t-1})(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1})\varepsilon_t$	$y_t = (\ell_{t-1} + b_{t-1} + s_{t-m})(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha(\ell_{t-1} + b_{t-1} + s_{t-m})\varepsilon_t$ $b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1} + s_{t-m})\varepsilon_t$ $s_t = s_{t-m} + \gamma(\ell_{t-1} + b_{t-1} + s_{t-m})\varepsilon_t$	$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m}(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1})\varepsilon_t$ $s_t = s_{t-m}(1 + \gamma \varepsilon_t)$
A_d	$y_t = (\ell_{t-1} + \phi b_{t-1})(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + \phi b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = \phi b_{t-1} + \beta(\ell_{t-1} + \phi b_{t-1})\varepsilon_t$	$y_t = (\ell_{t-1} + \phi b_{t-1} + s_{t-m})(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha(\ell_{t-1} + \phi b_{t-1} + s_{t-m})\varepsilon_t$ $b_t = \phi b_{t-1} + \beta(\ell_{t-1} + \phi b_{t-1} + s_{t-m})\varepsilon_t$ $s_t = s_{t-m} + \gamma(\ell_{t-1} + \phi b_{t-1} + s_{t-m})\varepsilon_t$	$y_t = (\ell_{t-1} + \phi b_{t-1}) s_{t-m}(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + \phi b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = \phi b_{t-1} + \beta(\ell_{t-1} + \phi b_{t-1})\varepsilon_t$ $s_t = s_{t-m}(1 + \gamma \varepsilon_t)$

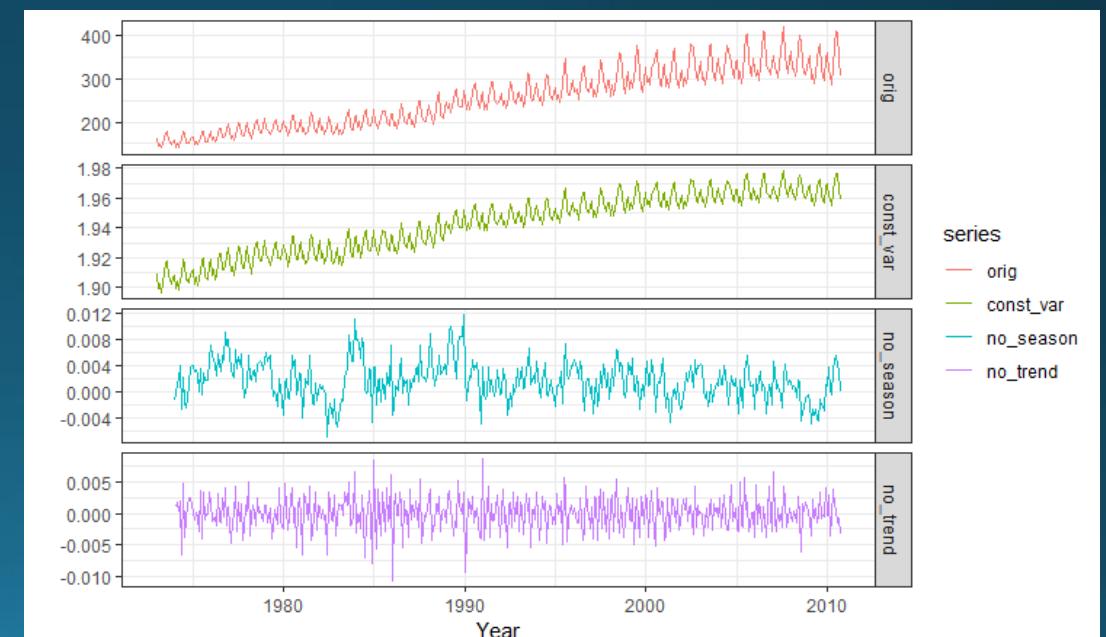
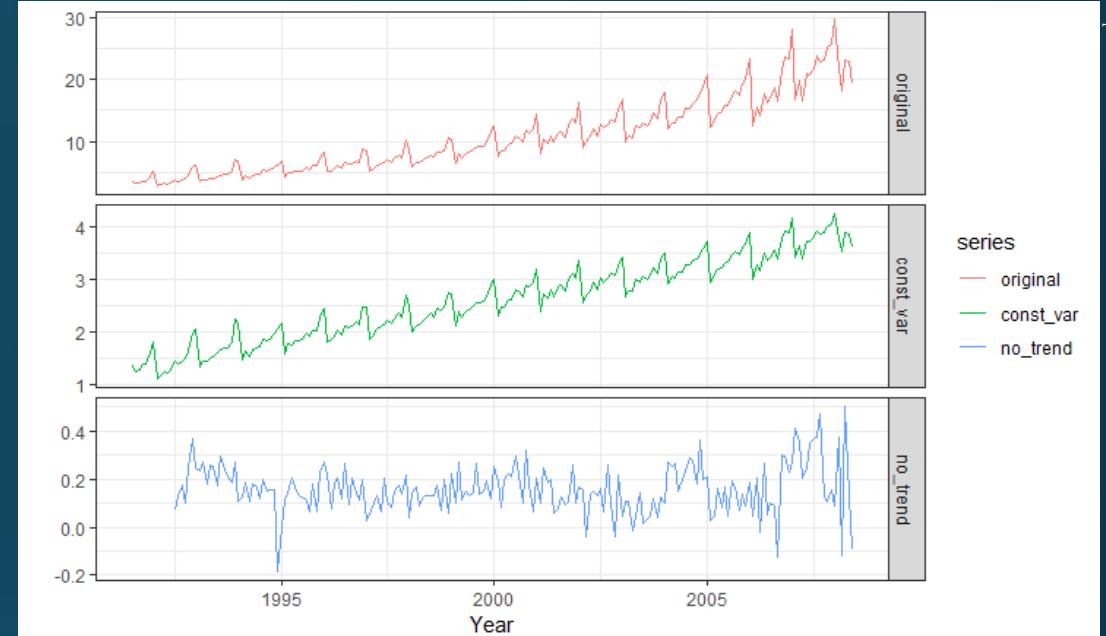
ARIMA

- Exponential Smoothing and ARIMA are the two most widely used approaches to time-series forecasting and provide complementary approaches to the problem.
- While Exponential Smoothing models are based on a description of trend and seasonality in the data, ARIMA models aim to describe autocorrelations in the data.

- A stationary time series is one whose properties do not depend on the time at which the series is observed. Thus, time series with trends, or with seasonality, are not stationary
- Stationarity is an assumption of many time-series analysis procedures.
- However, most time-series are non-stationary.
- Unit Root Tests: Statistical Test of Stationarity designed to see if differencing is required. E.g., KPSS Test.

Stationary Process

- Fortunately, a non-stationary process can be transformed into a (weakly) stationary process through transformations that remove trend and stabilize variance.
 - Stabilize variance: Box-Cox Transformation
 - Remove seasonality and trend: Differencing



- AutoRegressive Integrated Moving Average
- ARIMA = AR + I + MA

Differencing (I)

- Differencing can help stabilize the mean of a time series by removing changes in the level of a time series, and so eliminating trend and seasonality.

Autoregressive Models (AR)

- Forecast a variable using a linear combination of past values of the variable
- $y_t = c + \varphi_1 y_{t-1} + \varphi_2 y_{t-2} + \varphi_3 y_{t-3} + \dots + \varepsilon_t$
 - where ε_t is white noise

- A moving average model uses past forecast errors in a regression-like model
- $y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}$
 - where ε_t is white noise

- $y'_t = c + \varphi_1 y'_{t-1} + \cdots + \varphi_p y'_{t-p} + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t$
 - where y'_t is the differenced series
- ARIMA(p,d,q)
 - p = Number of ordinary Autoregressive lags
 - d = Number of ordinary differences
 - q = Number of ordinary Moving average lags

Seasonal ARIMA

- Non-Seasonal ARIMA + seasonal terms = Seasonal ARIMA
- ARIMA(p,d,q)(P,D,Q)m
 - (P,D,Q) is seasonal component
 - m is seasonal period

where

- p = Number of ordinary Autoregressive lags
- d = Number of ordinary differences
- q = Number of ordinary Moving average lags
- P = Number of seasonal Autoregressive lags
- D = Number of seasonal differences
- Q = Number of seasonal Moving Average lags
- m = Seasonal period or number of observations per year

- ARIMA model is estimated using maximum likelihood estimation. Technique finds the values of the parameters which maximize the probability of obtaining the data that we have observed
- Information Criteria
 - AIC
 - BIC
 - AICc
- Information are useful for selecting the best values of p and q but not for selecting the appropriate order of differencing (d). Differencing changes the data on which likelihood is computed making models with different d not comparable.

- Automatic Process
 - `auto.arima()`
 - Uses computational shortcuts, so solution may not be optimal
 - Ensure more thorough examination by
 - `stepwise=F`
 - `approximation=F`

- Steps to select ARIMA model
 - 1. Plot data to identify any unusual observations
 - 2. If necessary, transform data to stabilize variance
 - 3. If the data are non-stationary, take first differences until the data are stationary
 - 4. Examine ACF and PACF to decide on AR (ARIMA(p,d,o) and/or MA terms (ARIMA(o,d,q))
 - 5. Try chosen model and use AICc to search for a better model
 - 6. Check residuals by plotting. If they do not look like white noise, try a modified model
 - 7. Once the residuals look like white noise, calculate forecasts

- Dynamic regression models
- Hierarchical and grouped time series
- TBATS model for complex seasonality
- Vector autoregressions
- Neural network autoregression

Illustration - R

About Time Series

Simplifying Assumptions

Time Series Classes

Nature of Time Series

Time Series Patterns

Time Series Components

Simple Forecasting Methods

Exponential Smoothing Models

ETS Models

ARIMA

Comparing Forecasting Models

Simple Forecasting Methods

Average Method

Naive Method

Seasonal Naive Method

Drift Method

Exponential Smoothing Models

Simple exponential smoothing

Holt's Method

Holt's Method with Damping

Holt-Winter's seasonal method

Holt_Winter's Additive

Holt_Winter's Multiplicative

ETS Models

ETS: AAA

ETS: Automatic Selection

ARIMA

Stationary Process

ARIMA Model

ARIMA - Automatic Model

Selection

Time Series

Packages Used

```
library(ggplot2);library(ggthemes);library(gridExtra) # For plots
library(quantmod);library(xts);library(zoo) # For using xts class objects
library(forecast) # Set of forecasting functions
library(fpp); library(fpp2) # Datasets from Forecasting text by Rob Hyndman
library(tseries) # for a statistical test
library(dplyr) # Data wrangling
```



The screenshot shows the R documentation interface for the `auto.arima` function. The URL in the address bar is `R: Fit best ARIMA model to univariate time series`. The main content area displays the following text:

`auto.arima {forecast}`

Fit best ARIMA model to univariate time series

Description

Returns best ARIMA model according to either AIC, AICc or BIC value. The function conducts a search over possible model within the order constraints provided.

About Time Series

A lot of data around us is measured over a span of time

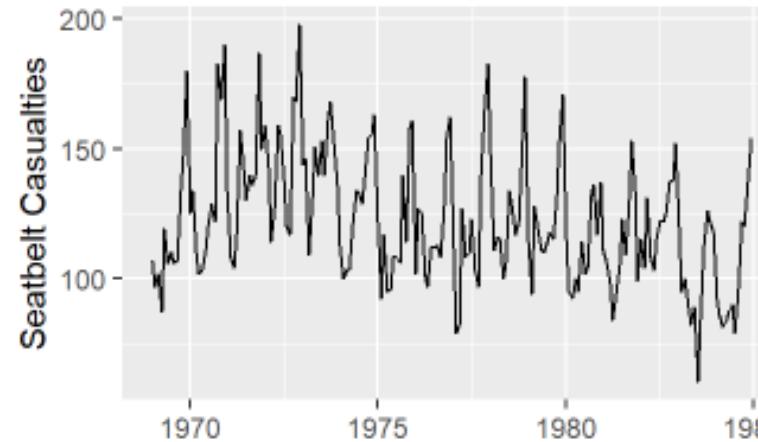
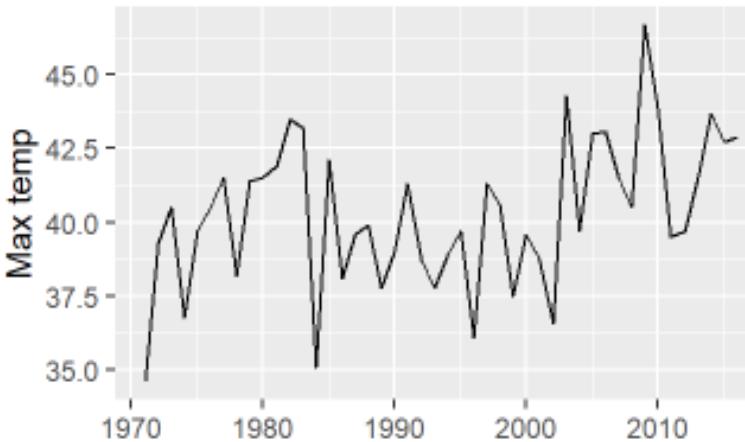
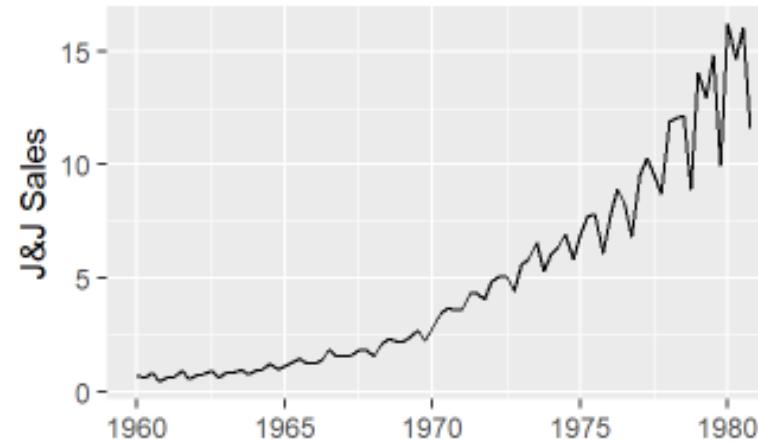
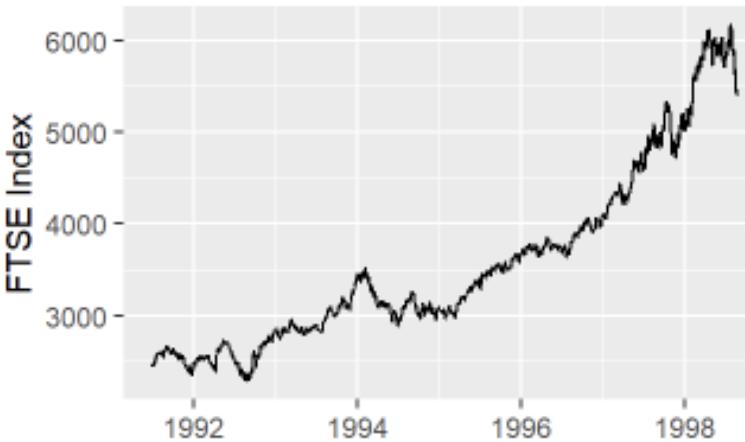
- Stock prices
- Federal funds rate
- Sales
- Weather

```
# ?EuStockMarkets  
# ?JohnsonJohnson  
# ?maxtemp  
# ?Seatbelts
```

```

library(gridExtra)
grid.arrange(autoplot(EuStockMarkets[, 'FTSE'])+xlab('')+ylab('FTSE Index'),
             autoplot(JohnsonJohnson)+xlab('')+ylab('J&J Sales'),
             autoplot(maxtemp)+xlab('')+ylab('Max temp'),
             autoplot(Seatbelts[, 'DriversKilled'])+xlab('')+ylab('Seatbelt Casualties'))

```



Console Terminal × Jobs ×

~ / . . .

> View(EuStockMarkets)

> View(JohnsonJohnson)

EuStockMarkets × timeSeries Sp2021_kkc.R

Filter

	DAX	SMI	CAC	FTSE
1	1628.75	1678.1	1772.8	2443.6
2	1613.63	1688.5	1750.5	2460.2
3	1606.51	1678.6	1718.0	2448.2
4	1621.04	1684.1	1708.1	2470.4
5	1618.16	1686.6	1723.1	2484.7
6	1610.61	1671.6	1714.3	2466.8
7	1630.75	1682.9	1734.5	2487.9
8	1640.17	1703.6	1757.4	2508.4

JohnsonJohnson

Filter

	V1
1	0.71
2	0.63
3	0.85
4	0.44
5	0.61
6	0.69
7	0.92
8	0.55

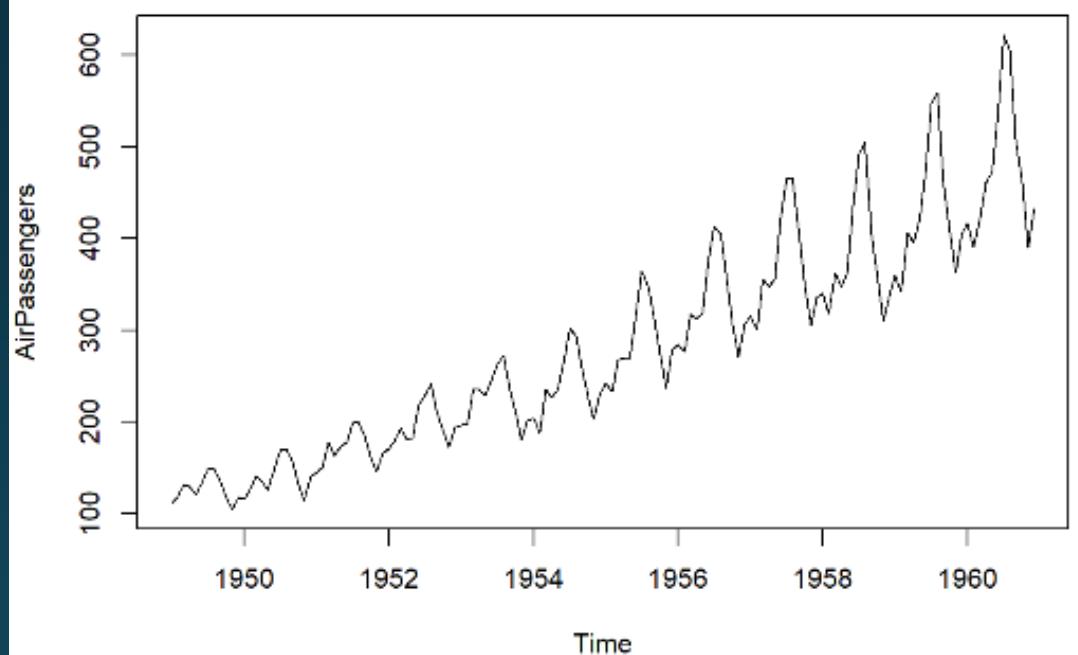
Simplifying Assumptions

We make some simplifying assumptions about time-series

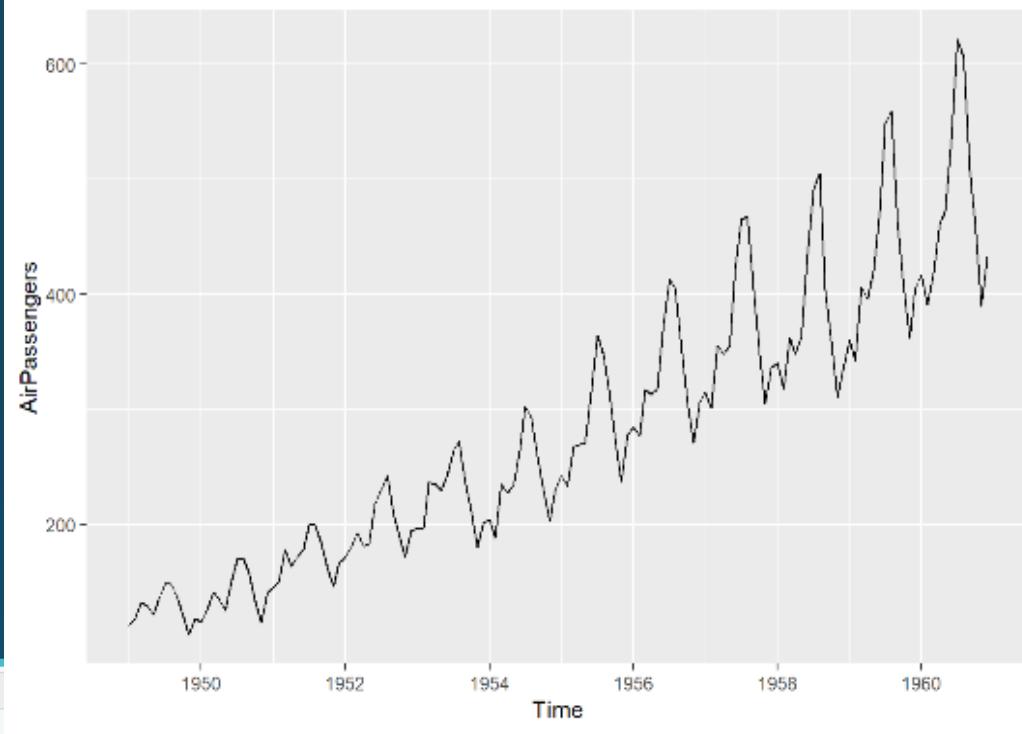
- Consecutive observations are equally spaced
- Apply a discrete-time index (even though this may only hold approximately).
E.g., daily stock returns are only available for weekdays, not weekends and holidays.

Lets see if AirPassengers data meets these assumptions

```
plot(AirPassengers)
```



```
autoplot(AirPassengers)
```



Files Plots Packages Help

R: Monthly Airline Passenger Numbers 1949-1960 ▾ Find in Topic

AirPassengers {datasets}

R Documentation

Monthly Airline Passenger Numbers 1949-1960

Description

The classic Box & Jenkins airline data. Monthly totals of international airline passengers, 1949 to 1960.

Usage

AirPassengers

Format

A monthly time series, in thousands.

Time Series Classes

- Time series data needs a special class because it is indexed by time/date
- Data structures designed for time series data: ts, xts, zoo, ..
- Use of these data structures offers access to useful functions for manipulating, analyzing and visualizing time-series data.

ts methods/functions

Let us take a look at the ts structure first.

AirPassengers

```
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 1949 112 118 132 129 121 135 148 148 136 119 104 118
## 1950 115 126 141 135 125 149 170 170 158 133 114 140
## 1951 145 150 178 163 172 178 199 199 184 162 146 166
## 1952 171 180 193 181 183 218 230 242 209 191 172 194
## 1953 196 196 236 235 229 243 264 272 237 211 180 201
## 1954 204 188 235 227 234 264 302 293 259 229 203 229
## 1955 242 233 267 269 270 315 364 347 312 274 237 278
## 1956 284 277 317 313 318 374 413 405 355 306 271 306
## 1957 315 301 356 348 355 422 465 467 404 347 305 336
## 1958 340 318 362 348 363 435 491 505 404 359 310 337
## 1959 360 342 406 396 420 472 548 559 463 407 362 405
## 1960 417 391 419 461 472 535 622 606 508 461 390 432
```

What type of data structure is AirPassengers?

Class

```
class(AirPassengers)
```

```
## [1] "ts"
```

First period of the time series

```
start(AirPassengers)
```

```
## [1] 1949 1
```

Last Period of the time series

```
end(AirPassengers)
```

```
## [1] 1960 12
```

time - creates the vector of times at which a time series was sampled.

All time periods

```
time(AirPassengers)
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
## 1949	1949.000	1949.083	1949.167	1949.250	1949.333	1949.417	1949.500	1949.583
## 1950	1950.000	1950.083	1950.167	1950.250	1950.333	1950.417	1950.500	1950.583
## 1951	1951.000	1951.083	1951.167	1951.250	1951.333	1951.417	1951.500	1951.583
## 1952	1952.000	1952.083	1952.167	1952.250	1952.333	1952.417	1952.500	1952.583
## 1953	1953.000	1953.083	1953.167	1953.250	1953.333	1953.417	1953.500	1953.583
## 1954	1954.000	1954.083	1954.167	1954.250	1954.333	1954.417	1954.500	1954.583
## 1955	1955.000	1955.083	1955.167	1955.250	1955.333	1955.417	1955.500	1955.583
## 1956	1956.000	1956.083	1956.167	1956.250	1956.333	1956.417	1956.500	1956.583
## 1957	1957.000	1957.083	1957.167	1957.250	1957.333	1957.417	1957.500	1957.583
## 1958	1958.000	1958.083	1958.167	1958.250	1958.333	1958.417	1958.500	1958.583
## 1959	1959.000	1959.083	1959.167	1959.250	1959.333	1959.417	1959.500	1959.583
## 1960	1960.000	1960.083	1960.167	1960.250	1960.333	1960.417	1960.500	1960.583
	Sep	Oct	Nov	Dec				
## 1949	1949.667	1949.750	1949.833	1949.917				
## 1950	1950.667	1950.750	1950.833	1950.917				
## 1951	1951.667	1951.750	1951.833	1951.917				
## 1952	1952.667	1952.750	1952.833	1952.917				
## 1953	1953.667	1953.750	1953.833	1953.917				
## 1954	1954.667	1954.750	1954.833	1954.917				
## 1955	1955.667	1955.750	1955.833	1955.917				
## 1956	1956.667	1956.750	1956.833	1956.917				
## 1957	1957.667	1957.750	1957.833	1957.917				
## 1958	1958.667	1958.750	1958.833	1958.917				
## 1959	1959.667	1959.750	1959.833	1959.917				
## 1960	1960.667	1960.750	1960.833	1960.917				

Period	
1	0.083
2	0.167
3	0.250
4	0.333
5	0.417
6	0.500
7	0.583
8	0.667
9	0.750
10	0.833
11	0.917
12	1.000

```
deltat(AirPassengers) # one month is 1/12 of a year
```

```
## [1] 0.08333333
```

```
cycle(AirPassengers)
```

```
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 1949   1   2   3   4   5   6   7   8   9   10  11  12
## 1950   1   2   3   4   5   6   7   8   9   10  11  12
## 1951   1   2   3   4   5   6   7   8   9   10  11  12
## 1952   1   2   3   4   5   6   7   8   9   10  11  12
## 1953   1   2   3   4   5   6   7   8   9   10  11  12
## 1954   1   2   3   4   5   6   7   8   9   10  11  12
## 1955   1   2   3   4   5   6   7   8   9   10  11  12
## 1956   1   2   3   4   5   6   7   8   9   10  11  12
## 1957   1   2   3   4   5   6   7   8   9   10  11  12
## 1958   1   2   3   4   5   6   7   8   9   10  11  12
## 1959   1   2   3   4   5   6   7   8   9   10  11  12
## 1960   1   2   3   4   5   6   7   8   9   10  11  12
```

deltat - the time interval between observations

cycle gives the positions in the cycle of each observation.

head and tail work in the same manner as on a data frame

```
head(AirPassengers)
```

```
##      Jan Feb Mar Apr May Jun
## 1949 112 118 132 129 121 135
```

```
tail(AirPassengers)
```

```
##      Jul Aug Sep Oct Nov Dec
## 1960 622 606 508 461 390 432
```

Explore data as an xts object

Convert AirPassengers to an xts object

```
AP_xts <- as.xts(AirPassengers)
```

Number of passengers for June, 1960

```
AP_xts['1960-06',]
```

```
##          [,1]
## Jun 1960 535
```

Average number of passengers for the year 1960

```
mean(AP_xts['1960',])
```

```
## [1] 476.1667
```

Number of months of data are included in this dataset?

```
nmonths(AP_xts)
```

```
## [1] 144
```

Calculate correlation between number of passengers and one-month lagged number of passengers

```
cor(AP_xts,lag.xts(AP_xts),use = 'complete.obs')
```

```
##          [,1]
## [1,] 0.9601946
```

Create a ts object

Lets create a ts object, even though in most cases we will work with data that is already a time series object. To define a ts object, we specify the data, start/end and frequency. Here, frequency is the number of observations per seasonal pattern

- annual is 1 (default)
- quarterly is 4
- monthly is 12
- weekly is 52
- other special ways to handle seasonality less than a week

```
sales_in_millions = sample(seq(1e2,1e4,1)/1e2,size = 96,replace = T)
sales_ts = ts(data=sales_in_millions,start = 2010,frequency = 12)
sales_ts
```

```
##      Jan   Feb   Mar   Apr   May   Jun   Jul   Aug   Sep   Oct   Nov   Dec
## 2010 19.28 3.79 63.88 94.65 52.21 57.91 59.36 61.65 24.37 28.73 20.64 90.32
## 2011 99.24 28.63 47.02 64.01 87.75 86.45 75.24 47.76 59.31 88.15 74.86 19.15
## 2012 48.07 19.70 46.42 27.58 71.14 67.60 84.03 17.52 99.20 80.80 45.57 67.89
## 2013 45.95 58.94 81.49 4.76 97.67 74.95 60.34 24.48 50.62 31.84 54.23 86.61
## 2014 30.67 69.23 55.78 22.58 97.06 69.66 95.45 9.91 44.06 83.33 57.91 68.62
## 2015 19.67 33.71 74.64 20.33 85.50 9.82 41.50 89.00 19.69 23.14 5.04 43.74
## 2016 17.52 96.59 60.61 21.96 20.08 36.83 31.90 72.53 86.14 96.17 60.69 67.83
## 2017 18.85 10.37 72.92 71.72 82.14 67.02 6.67 11.15 36.68 51.35 49.33 8.92
```

Verify that sales_ts is a ts object

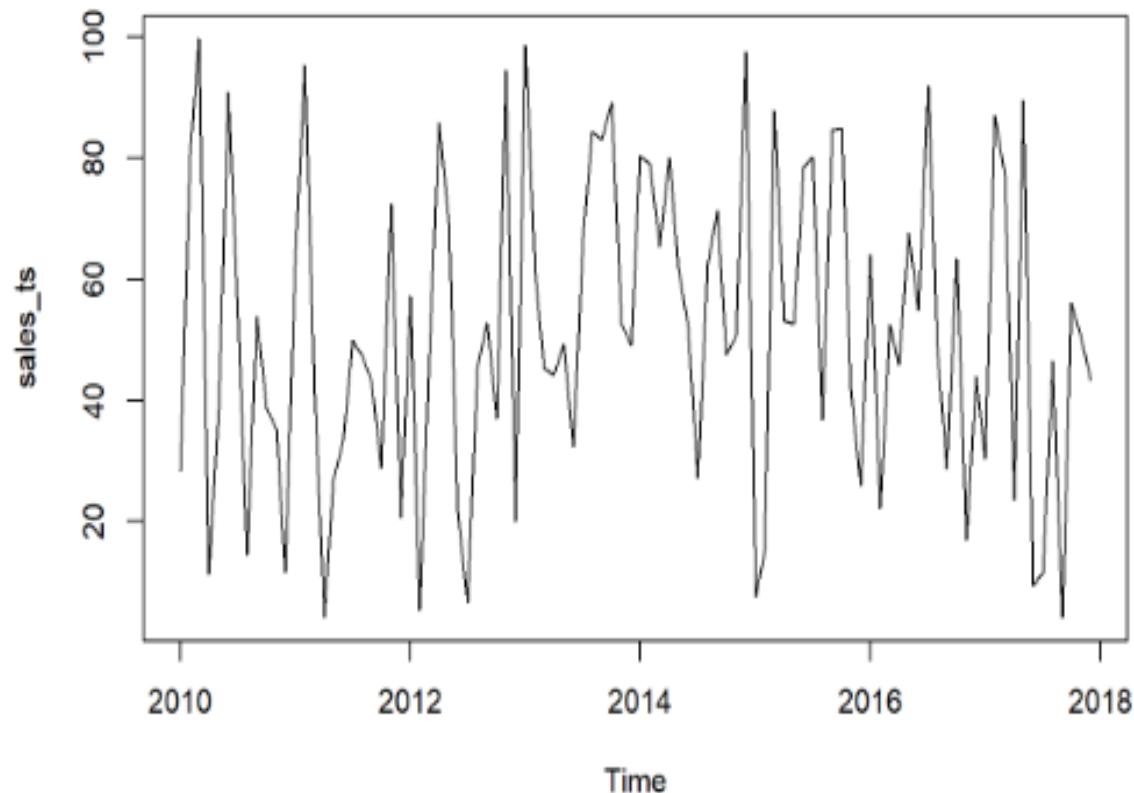
```
is.ts(sales_ts)
```

```
## [1] TRUE
```

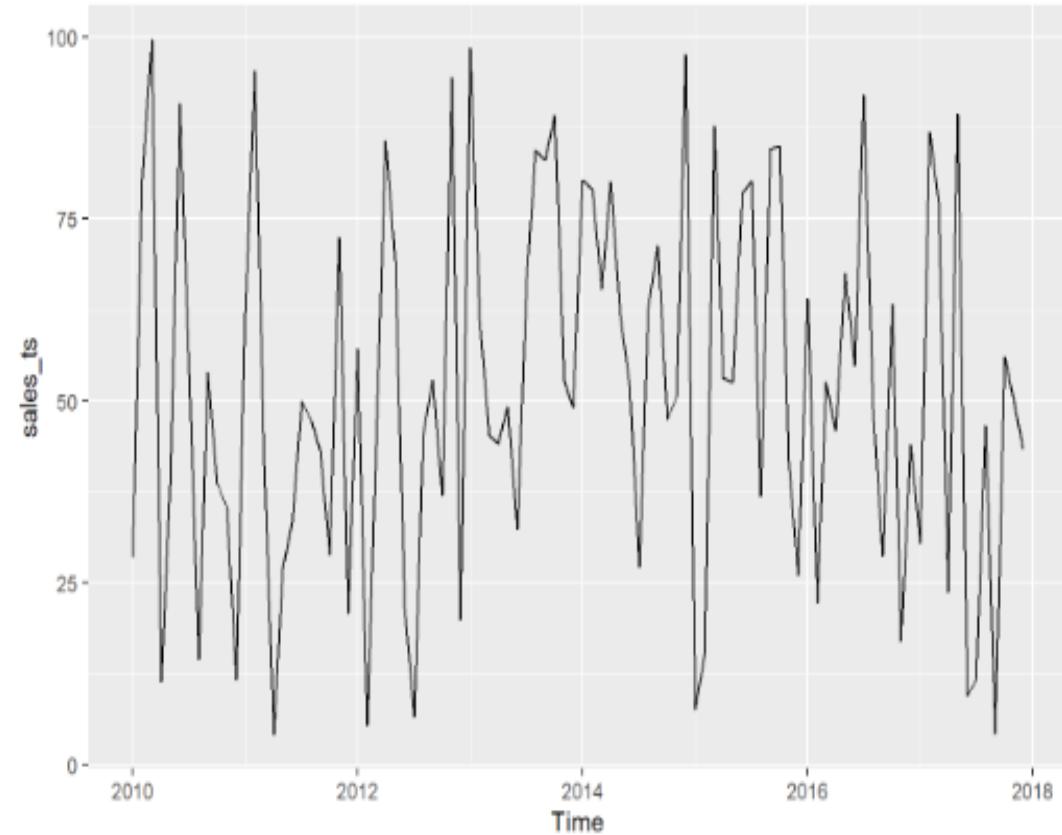
Plotting

ts objects can be plotted by using `plot()` or using `autoplot()` from library(`ggplot2`). The latter will generate a `ggplot2` object that generally looks prettier.

```
plot(sales_ts)
```



```
library(ggplot2)  
autoplot(sales_ts)
```



Nature of Time Series

A common observation across time-series is that prior observations influence later observations. The influence may be from the previous observation (trend), an observation that is a fixed number of periods ago (seasonality), or an observation that is not a fixed number of periods ago (cycle). Influence of prior observations can be examined by looking at lagged correlations or autocorrelations.

Lag 1 autocorrelation

Use AirPassengers data

```
acf(x = AirPassengers, lag.max = 1, plot=F)
```

no lag

```
##  
## Autocorrelations of series 'AirPassengers', by lag  
##  
## 0.0000 0.0833  
## 1.000 0.948
```

1 lag

This is the time stamp

Lag 1 and 2 Autocorrelation

```
acf(x=AirPassengers, lag.max=2, plot=F)
```

```
##  
## Autocorrelations of series 'AirPassengers', by lag  
##  
## 0.0000 0.0833 0.1667  
## 1.000 0.948 0.876
```

- Just as correlation measures the extent of a linear relationship between two variables, autocorrelation measures the linear relationship between lagged values of a time series.

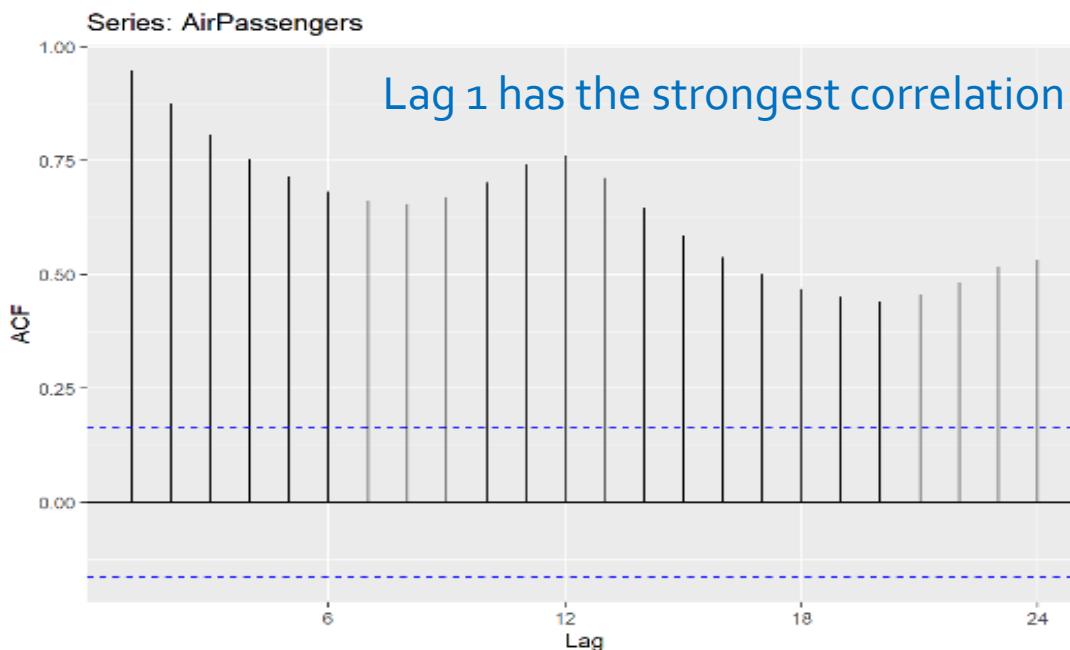
Examine the top few autocorrelations

```
acf(x = AirPassengers, plot=F)
```

```
##  
## Autocorrelations of series 'AirPassengers', by lag  
##  
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333  
## 1.000 0.948 0.876 0.807 0.753 0.714 0.682 0.663 0.656 0.671 0.703  
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667 1.7500  
## 0.743 0.760 0.713 0.646 0.586 0.538 0.500 0.469 0.450 0.442 0.457
```

Plot the top few autocorrelations but using ggAcf() instead of acf()

```
library(forecast)  
ggAcf(x = AirPassengers)
```



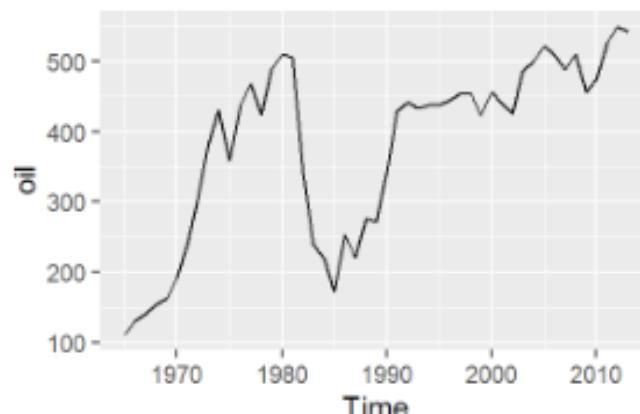
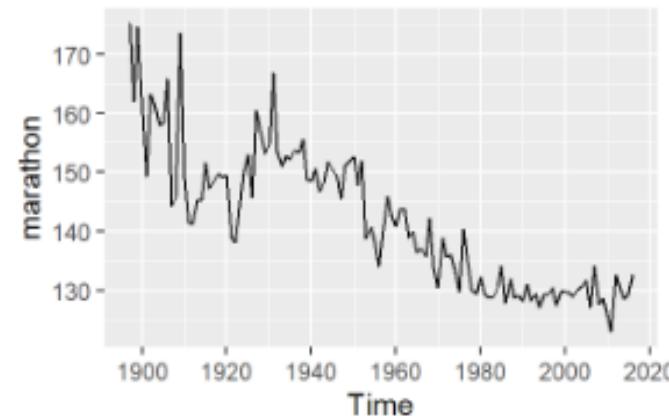
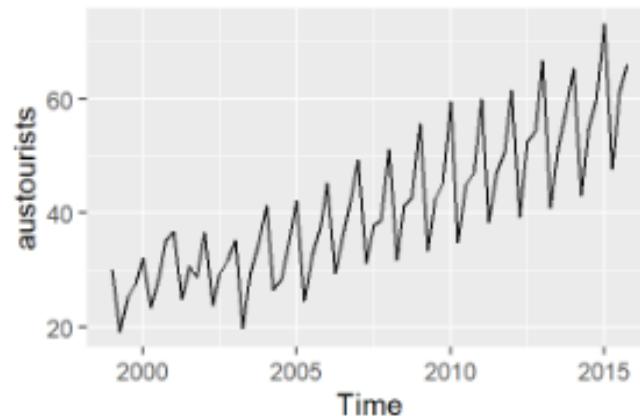
Time Series Patterns

- Trend: A long term increase or decrease in the data. Could be linear, non-linear, and may go from an increasing trend to a decreasing trend
- Seasonal: Pattern repeats itself every season. Here season may be period of the year or day of the week. Seasonality is always of a fixed or known frequency
- Cyclic: Data exhibit rises or falls that are not of a fixed frequency

On the other hand, a time-series data that does not exhibit any pattern is a Stationary Process. White Noise is the simplest Stationary Process

Which is the dominant time series pattern in these charts?

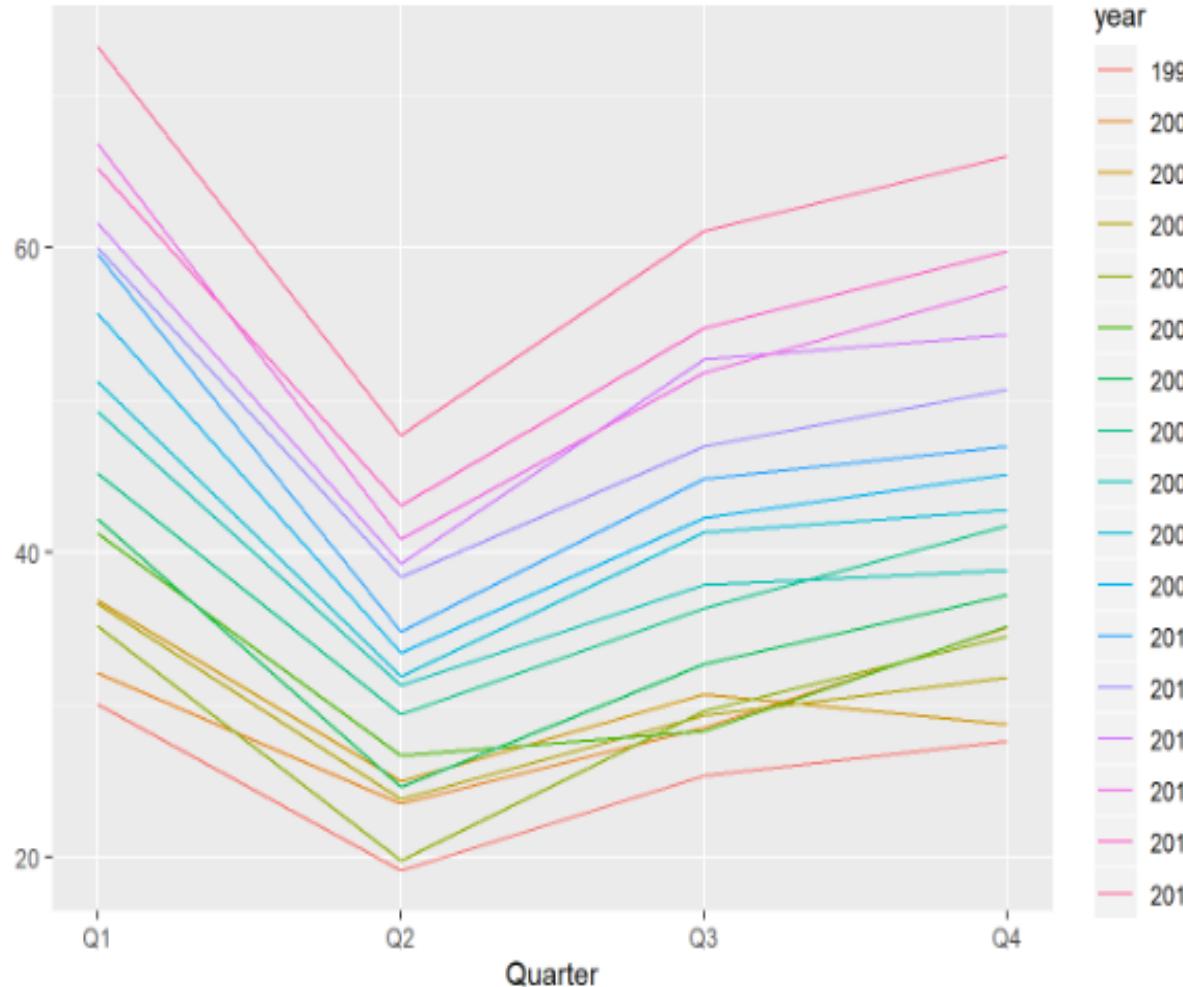
```
library(fpp)
library(gridExtra)
grid.arrange(autoplot(austourists), # International tourists to Australia
             autoplot(marathon),   # Marathon times
             autoplot(oil),        # Saudi Arabia oil production
             autoplot(wmurders))  # female murder rate
```



To more clearly spot seasonality, one can examine a seasonal plot.

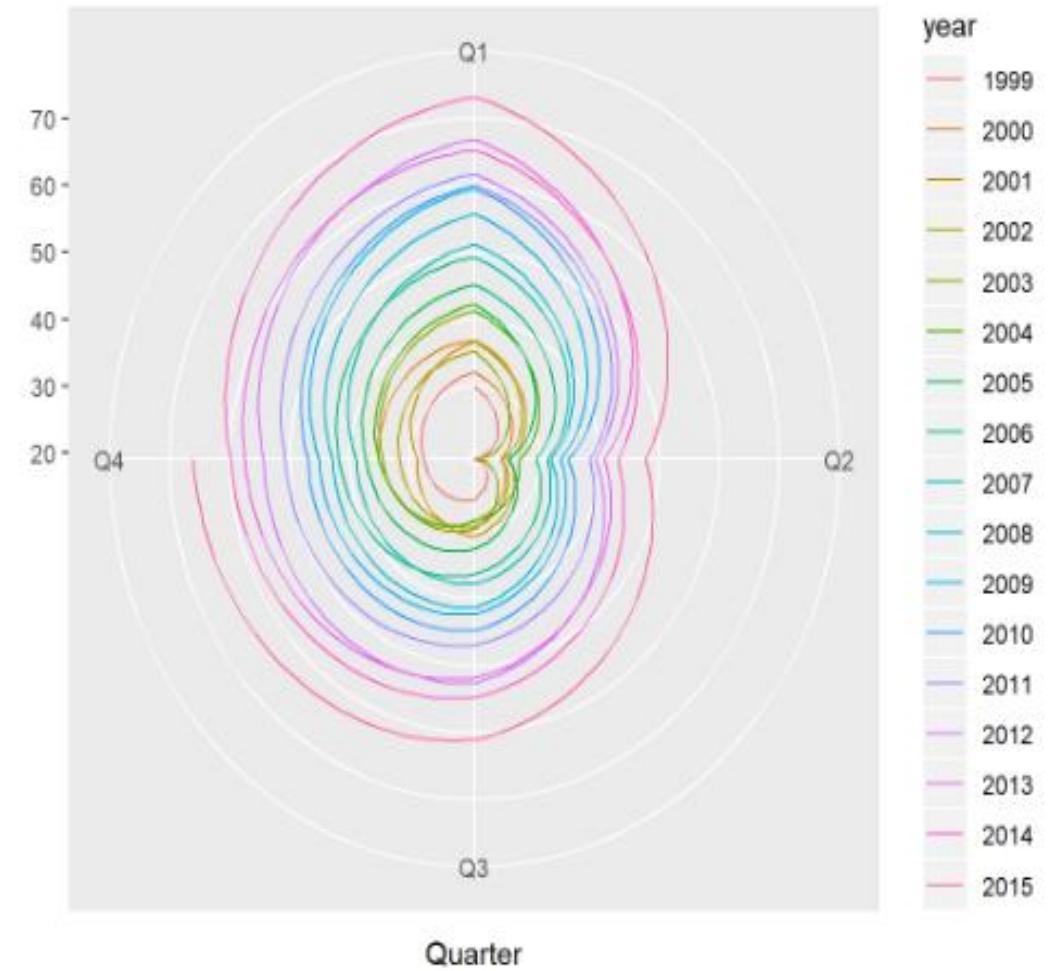
`ggseasonplot(austourists)`

Seasonal plot: austourists



`ggseasonplot(austourists,polar=T)`

Seasonal plot: austourists



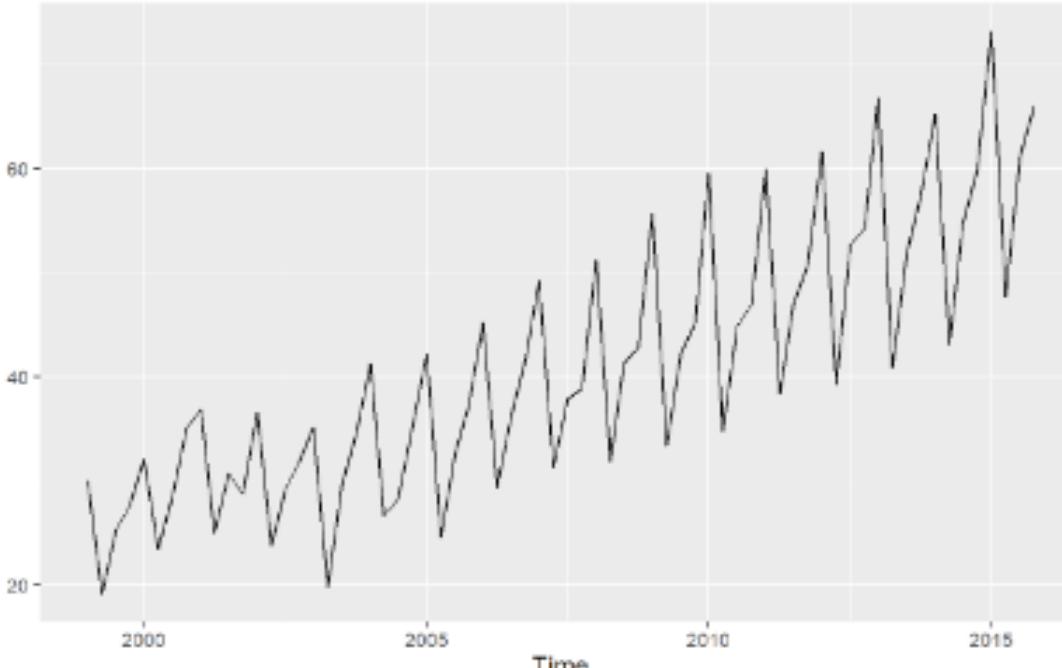
Time Series Components

Time-series may have multiple components in varying strengths. To better understand the series it is useful to divide it into its components:

- * Trend components
- * Seasonal component
- * Remainder

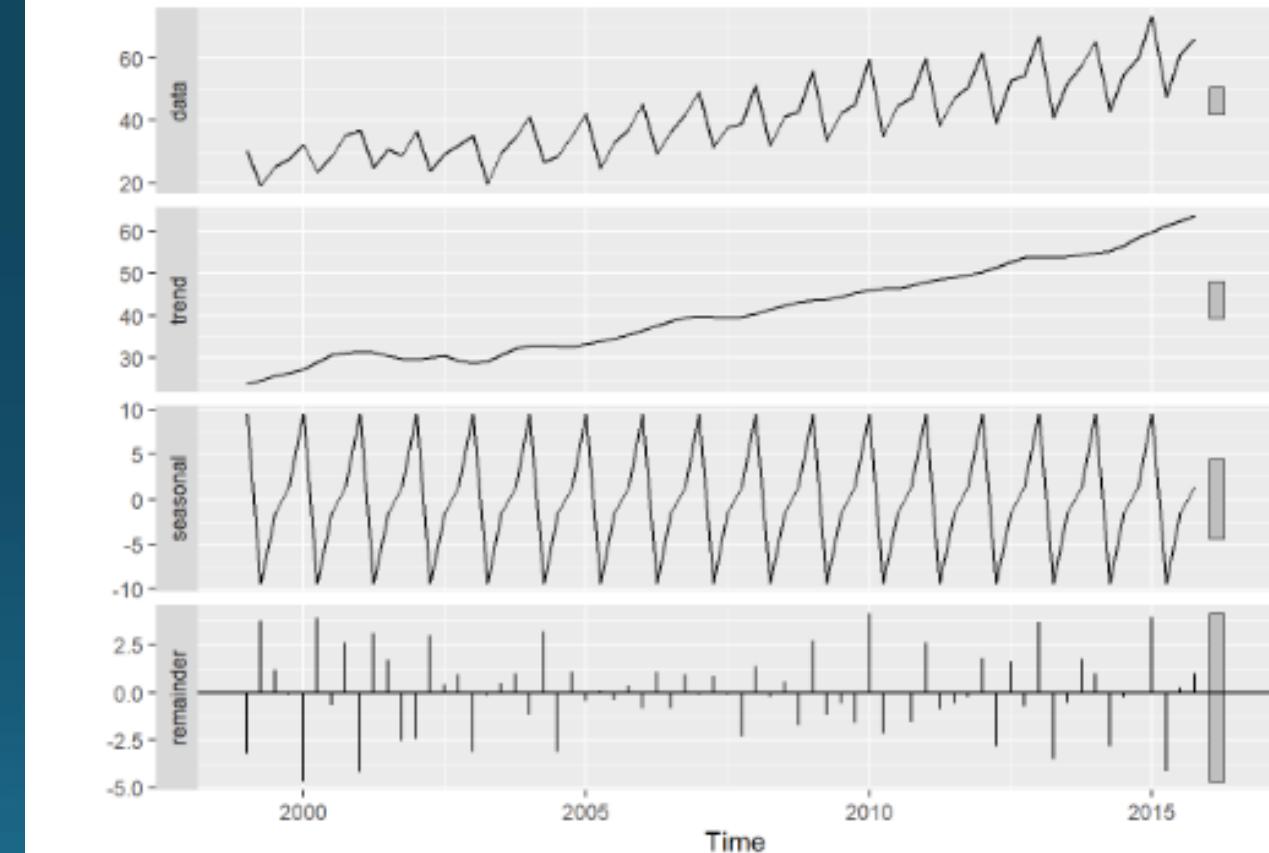
Here is a plot of quarterly visitor nights (in millions) spent by international tourists to Australia. 1999-2015

```
austourists %>%  
  autoplot()
```



Now, let us decompose this time series using STL (Seasonality and Trend Decomposition using Loess), a versatile and robust method for decomposing a time series. In fact, decomposition can be used to measure the strength of trend and seasonality.

```
austourists %>%  
  stl(s.window = 'periodic') %>%  
  autoplot()
```



Simple Forecasting Methods

Let us forecast quarterly australian beer production.

```
library(fpp2) ←
```

```
ausbeer
```

```
##      Qtr1 Qtr2 Qtr3 Qtr4
## 1956   284  213  227  308
## 1957   262  228  236  320
## 1958   272  233  237  313
## 1959   261  227  250  314
## 1960   286  227  260  311
## 1961   295  233  257  339
## 1962   279  250  270  346
## 1963   294  255  278  363
## 1964   313  273  300  370
## 1965   331  288  306  386
## 1966   335  288  308  402
## 1967   353  316  325  405
## 1968   393  319  327  442
## 1969   383  332  361  446
## 1970   387  357  374  466
## 1971   410  370  379  487
## 1972   419  378  393  506
## 1973   458  387  427  565
## 1974   465  445  450  556
## 1975   500  452  435  554
##      1991  464  424  436  574
##      1992  443  410  420  532
##      1993  433  421  410  512
##      1994  449  381  423  531
##      1995  426  408  416  520
##      1996  409  398  398  507
##      1997  432  398  406  526
##      1998  428  397  403  517
##      1999  435  383  424  521
##      2000  421  402  414  500
##      2001  451  380  416  492
##      2002  428  408  406  506
##      2003  435  380  421  490
##      2004  435  390  412  454
##      2005  416  403  408  482
##      2006  438  386  405  491
##      2007  427  383  394  473
##      2008  420  390  410  488
##      2009  415  398  419  488
##      2010  414  374
```

- Use the ausbeer data set

Split the data into train and test. Since the test data contains 42 quarters, we will be constructing forecasts for 42 periods.

```
train = window(ausbeer, end=c(1999,04))
test = window(ausbeer, start=c(2000,01))
length(test)
```

```
## [1] 42
```

Average Method

A very simple prediction, often the baseline in linear regression, is to use the average.

Model and Forecast

```
average_model = meanf(train,h = 42)  
average_model
```

80% and 95% confidence bands

#		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
#	2000 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
#	2000 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
#	2000 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
#	2000 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
#	2001 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
#	2001 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
#	2001 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
#	2001 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
#	2002 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
#	2002 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
#	2002 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
#	2002 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
#	2003 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
#	2003 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
#	2003 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
#	2003 Q4	412.4318	291.6172	533.2464	227.0775	597.7861

- forecast for beer production over the next 42 months of the test sample.
- the point forecast of Australian beer production for Q2 2003 is 412.4318

##	2004 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2004 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
##	2004 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
##	2004 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
##	2005 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2005 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
##	2005 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
##	2005 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
##	2006 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2006 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
##	2006 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
##	2006 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
##	2007 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2007 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
##	2007 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
##	2007 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
##	2008 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2008 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
##	2008 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
##	2008 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
##	2009 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2009 Q2	412.4318	291.6172	533.2464	227.0775	597.7861
##	2009 Q3	412.4318	291.6172	533.2464	227.0775	597.7861
##	2009 Q4	412.4318	291.6172	533.2464	227.0775	597.7861
##	2010 Q1	412.4318	291.6172	533.2464	227.0775	597.7861
##	2010 Q2	412.4318	291.6172	533.2464	227.0775	597.7861

The model object contains a point forecast, and 80% and 95% confidence bands for the estimate. Let us extract the estimate for the last observation, 2010, Q2

```
average_model$mean
```

```
##          Qtr1      Qtr2      Qtr3      Qtr4
## 2000 412.4318 412.4318 412.4318 412.4318
## 2001 412.4318 412.4318 412.4318 412.4318
## 2002 412.4318 412.4318 412.4318 412.4318
## 2003 412.4318 412.4318 412.4318 412.4318
## 2004 412.4318 412.4318 412.4318 412.4318
## 2005 412.4318 412.4318 412.4318 412.4318
## 2006 412.4318 412.4318 412.4318 412.4318
## 2007 412.4318 412.4318 412.4318 412.4318
## 2008 412.4318 412.4318 412.4318 412.4318
## 2009 412.4318 412.4318 412.4318 412.4318
## 2010 412.4318 412.4318
```

```
window(average_model$mean, c(2010, 02))
```

```
##          Qtr2
## 2010 412.4318
```

Accuracy

Let us examine the accuracy of the above prediction from `average_model` on the train sample. `accuracy()` from library(`forecast`) is a handy function that automatically computes a set of error indices.

```
accuracy(average_model)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -5.118499e-15 93.38414 74.56612 -6.210103 20.74216 4.668865
##                               ACF1
## Training set 0.7089738
```

Let us examine the accuracy on the test sample. To get test-set performance, set the `x` argument as the entire dataset rather than just the test set.

```
accuracy(average_model, x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -5.118499e-15 93.38414 74.56612 -6.210103 20.742160 4.668865
## Test set      1.525866e+01 40.10097 28.70346  2.876448  6.348359 1.797232
##                               ACF1 Theil's U
## Training set  0.70897377     NA
## Test set     -0.05298108 0.7832398
```

Now, let's set `x` argument as the test set and compare the results.

```
accuracy(average_model, x = test)
```

the RMSE of the prediction in the train sample

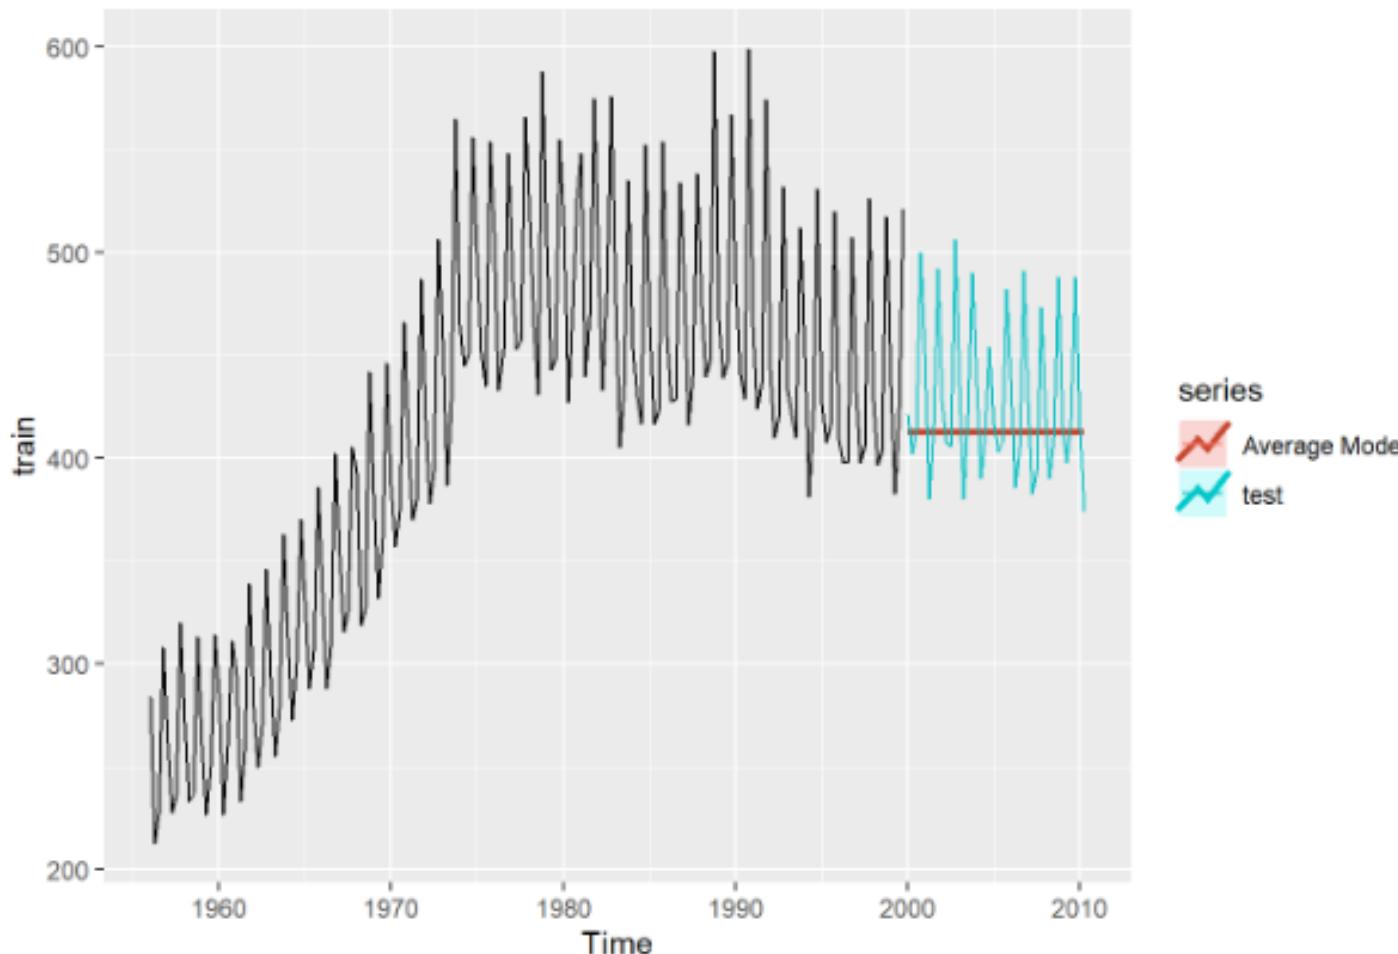
```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -5.118499e-15 93.38414 74.56612 -6.210103 20.742160 4.668865
## Test set      1.525866e+01 40.10097 28.70346  2.876448  6.348359 1.797232
##                               ACF1 Theil's U
## Training set  0.70897377     NA
## Test set     -0.05298108 0.7832398
```

the RMSE of the `average_model` on the test sample

Visualize Forecast

autoplot() uses a layered approach to plotting like ggplot2 and plot().

```
autoplot(train) +  
  autolayer(average_model, PI = F, size=1.1, series = 'Average Model') +  
  autolayer(test)
```



Naive Method

- Future will be the same as the last observation
- Since this is the best prediction for a random walk, these are also called random walk forecasts

Model and Forecast

```
naive_model = naive(train,h=42)
```

Let's see the point forecast for 2010 Q2.

```
naive_model$mean
```

```
##      Qtr1 Qtr2 Qtr3 Qtr4
## 2000   521   521   521   521
## 2001   521   521   521   521
## 2002   521   521   521   521
## 2003   521   521   521   521
## 2004   521   521   521   521
## 2005   521   521   521   521
## 2006   521   521   521   521
## 2007   521   521   521   521
## 2008   521   521   521   521
## 2009   521   521   521   521
## 2010   521   521
```

- the point forecast of beer production for Q2, 2010 is 521

Note, for a naive model, the forecast is the same as the last observation in the train sample

```
last(train)
```

```
## [1] 521
```

Use naive_model to construct a forecast for beer production over the next 42 months of the test sample.

Accuracy

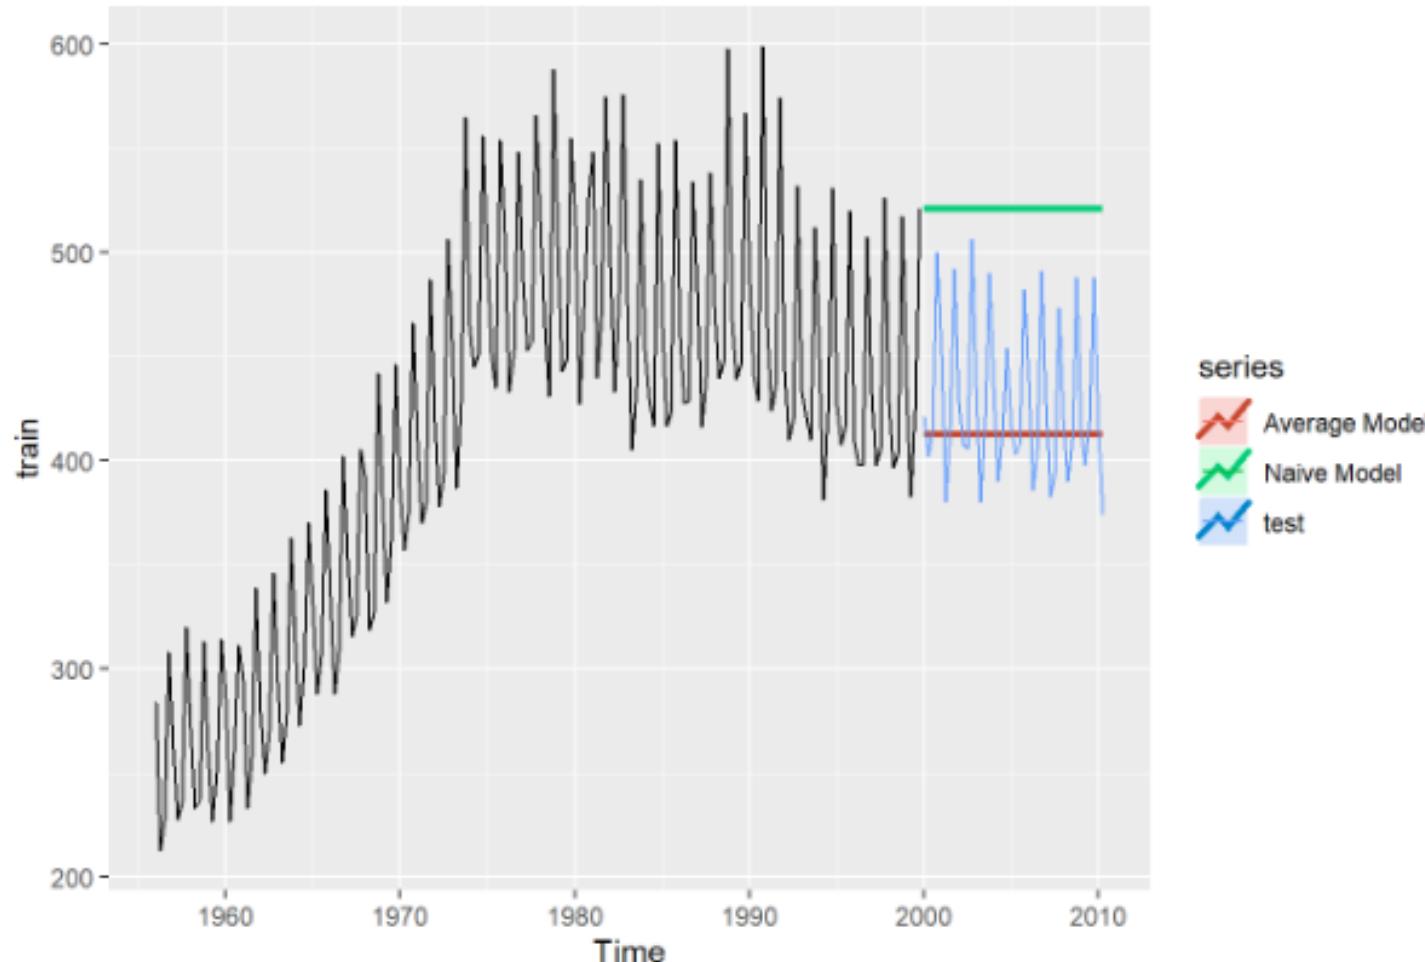
```
accuracy(naive_model,x=ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 1.354286 70.30818 58.73714 -0.967862 13.94342 3.677753
## Test set     -93.309524 100.40881 -93.30952 -22.690269 22.69027 5.842460
##                   ACF1 Theil's U
## Training set -0.22478344      NA
## Test set     -0.05298108 1.898409
```

RMSE of the naive_model on the test sample

Visualize Forecast

```
autoplot(train)+  
  autolayer(average_model,PI = F,size=1.1,series = 'Average Model')+  
  autolayer(naive_model,PI=F,size=1.1, series='Naive Model')+  
  autolayer(test)
```



Seasonal Naive Method

Forecast is equal to the last observed value from the same season

Model and Forecast

```
seasonal_naive_model = snaive(train,h=42)
seasonal_naive_model
```

```
##          Point Forecast    Lo 80     Hi 80    Lo 95     Hi 95
## 2000 Q1        435 409.3535 460.6465 395.7771 474.2229
## 2000 Q2        383 357.3535 408.6465 343.7771 422.2229
## 2000 Q3        424 398.3535 449.6465 384.7771 463.2229
## 2000 Q4        521 495.3535 546.6465 481.7771 560.2229
## 2001 Q1        435 398.7304 471.2696 379.5304 490.4696
## 2001 Q2        383 346.7304 419.2696 327.5304 438.4696
## 2001 Q3        424 387.7304 460.2696 368.5304 479.4696
## 2001 Q4        521 484.7304 557.2696 465.5304 576.4696
## 2002 Q1        435 390.5790 479.4210 367.0639 502.9361
## 2002 Q2        383 338.5790 427.4210 315.0639 450.9361
## 2002 Q3        424 379.5790 468.4210 356.0639 491.9361
## 2002 Q4        521 476.5790 565.4210 453.0639 588.9361
## 2003 Q1        435 383.7070 486.2930 356.5542 513.4458
## 2003 Q2        383 331.7070 434.2930 304.5542 461.4458
## 2003 Q3        424 372.7070 475.2930 345.5542 502.4458
## 2003 Q4        521 469.7070 572.2930 442.5542 599.4458
## 2004 Q1        435 377.6527 492.3473 347.2949 522.7051
## 2004 Q2        383 325.6527 440.3473 295.2949 470.7051
## 2004 Q3        424 366.6527 481.3473 336.2949 511.7051
## 2004 Q4        521 463.6527 578.3473 433.2949 608.7051
## 2005 Q1        435 372.1792 497.8208 338.9239 531.0761
## 2005 Q2        383 320.1792 445.8208 286.9239 479.0761
## 2005 Q3        424 361.1792 486.8208 327.9239 520.0761
## 2005 Q4        521 458.1792 583.8208 424.9239 617.0761
## 2006 Q1        435 367.1458 502.8542 331.2259 538.7741
## 2006 Q2        383 315.1458 450.8542 279.2259 486.7741
## 2006 Q3        424 356.1458 491.8542 320.2259 527.7741
## 2006 Q4        521 453.1458 588.8542 417.2259 624.7741
```

```
## 2007 Q1        435 362.4608 507.5392 324.0608 545.9392
## 2007 Q2        383 310.4608 455.5392 272.0608 493.9392
## 2007 Q3        424 351.4608 496.5392 313.0608 534.9392
## 2007 Q4        521 448.4608 593.5392 410.0608 631.9392
## 2008 Q1        435 358.0605 511.9395 317.3312 552.6688
## 2008 Q2        383 306.0605 459.9395 265.3312 500.6688
## 2008 Q3        424 347.0605 500.9395 306.3312 541.6688
## 2008 Q4        521 444.0605 597.9395 403.3312 638.6688
## 2009 Q1        435 353.8987 516.1013 310.9662 559.0338
## 2009 Q2        383 301.8987 464.1013 258.9662 507.0338
## 2009 Q3        424 342.8987 505.1013 299.9662 548.0338
## 2009 Q4        521 439.8987 602.1013 396.9662 645.0338
## 2010 Q1        435 349.9402 520.0598 304.9123 565.0877
## 2010 Q2        383 297.9402 468.0598 252.9123 513.0877
```

Point forecast for 2010 Q2. Note the similarity in predictions for any given season.

```
seasonal_naive_model$mean
```

```
##          Qtr1   Qtr2   Qtr3   Qtr4
## 2000      435     383     424     521
## 2001      435     383     424     521
## 2002      435     383     424     521
## 2003      435     383     424     521
## 2004      435     383     424     521
## 2005      435     383     424     521
## 2006      435     383     424     521
## 2007      435     383     424     521
## 2008      435     383     424     521
## 2009      435     383     424     521
## 2010      435     383
```

- Use the ausbeer data set

- Recall:

```
train = window(ausbeer, end=c(1999,04))
test = window(ausbeer, start=c(2000,01))
```

```
library(fpp2)
ausbeer
```

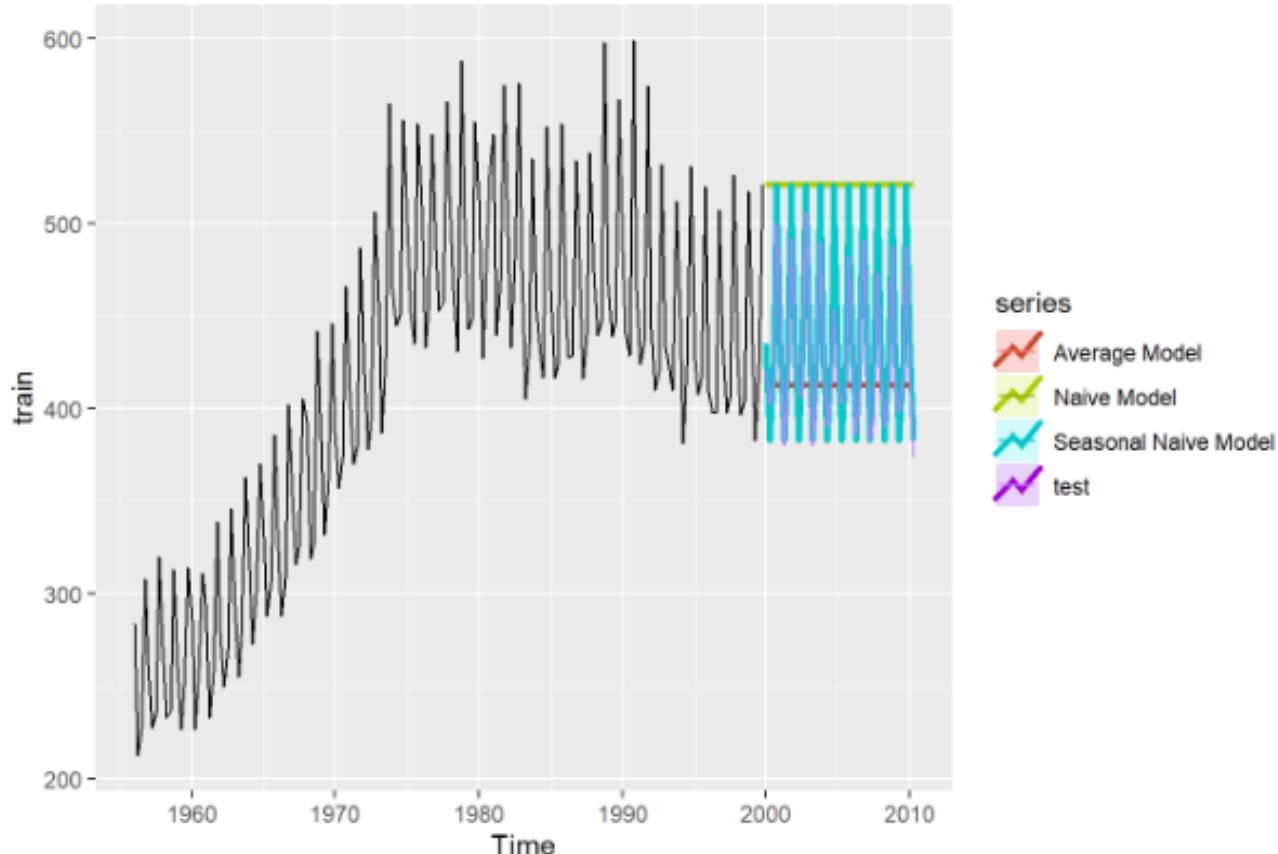
```
##          Qtr1   Qtr2   Qtr3   Qtr4
## 1991      464     424     436     574
## 1992      443     410     420     532
## 1993      433     421     410     512
## 1994      449     381     423     531
## 1995      426     408     416     520
## 1996      409     398     398     507
## 1997      432     398     406     526
## 1998      428     397     403     517
## 1999      435     383     424     521
## 2000      421     402     414     500
## 2001      451     380     416     492
## 2002      428     408     406     506
## 2003      435     380     421     490
## 2004      435     390     412     454
## 2005      416     403     408     482
## 2006      438     386     405     491
## 2007      427     383     394     473
## 2008      420     390     410     488
## 2009      415     398     419     488
## 2010      414     374
```

```
accuracy(seasonal_naive_model,x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 4.25000 20.01206 15.97093 1.148642 3.818071 1.000000 0.02734555
## Test set     -11.54762 21.85286 17.02381 -2.521773 3.863916 1.065925 0.02909800
##               Theil's U
## Training set      NA
## Test set       0.4278746
```

Visualize Forecast

```
autoplot(train)+
  autolayer(average_model,PI = F,size=1.1,series = 'Average Model')+
  autolayer(naive_model,PI=F,size=1.1, series='Naive Model')+
  autolayer(seasonal_naive_model,PI=F,size=1.1,series='Seasonal Naive Model')+
  autolayer(test)
```



Drift Method

Allow forecasts to increase or decrease over time, where the amount of change over time (called the drift) is set to be average change seen in historical data

Model and Forecast

```
drift_model = r wf(train,h=42,drift = T)
drift_model$mean
```

```
##          Qtr1      Qtr2      Qtr3      Qtr4
## 2000 522.3543 523.7086 525.0629 526.4171
## 2001 527.7714 529.1257 530.4800 531.8343
## 2002 533.1886 534.5429 535.8971 537.2514
## 2003 538.6057 539.9600 541.3143 542.6686
## 2004 544.0229 545.3771 546.7314 548.0857
## 2005 549.4400 550.7943 552.1486 553.5029
## 2006 554.8571 556.2114 557.5657 558.9200
## 2007 560.2743 561.6286 562.9829 564.3371
## 2008 565.6914 567.0457 568.4000 569.7543
## 2009 571.1086 572.4629 573.8171 575.1714
## 2010 576.5257 577.8800
```

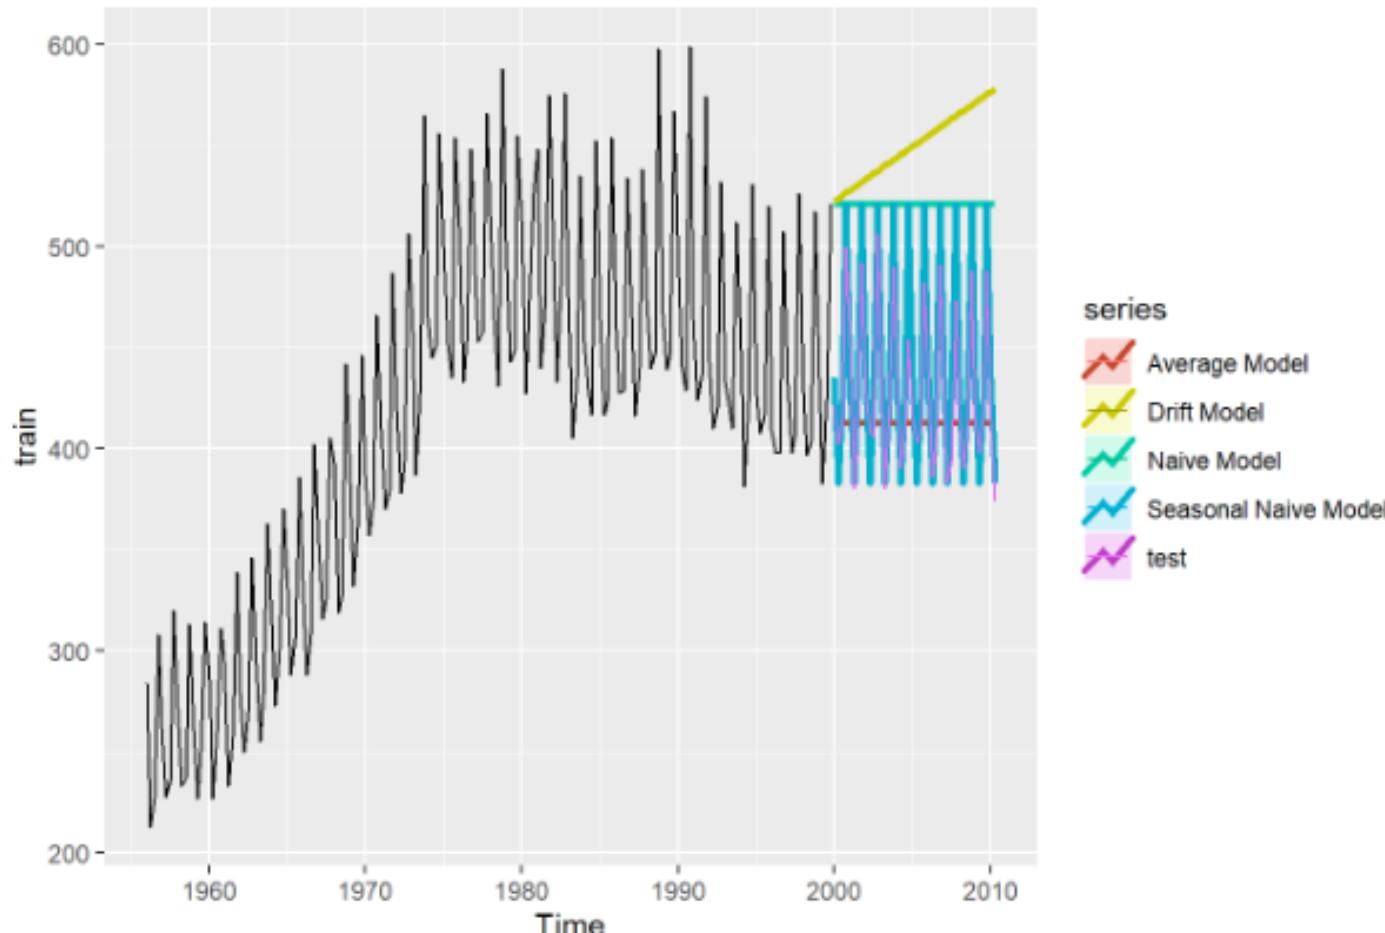
Accuracy

```
accuracy(drift_model,x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -3.572936e-15 70.29513 58.81084 -1.315888 13.97802 3.682368
## Test set     -1.224267e+02 129.44954 122.42667 -29.581531 29.58153 7.665594
##                         ACF1 Theil's U
## Training set -0.2247834        NA
## Test set      0.1520385 2.458927
```

Visualize Forecast

```
autoplot(train)+  
  autolayer(average_model,PI = F,size=1.1,series = 'Average Model')+  
  autolayer(naive_model,PI=F,size=1.1, series='Naive Model')+  
  autolayer(seasonal_naive_model,PI=F,size=1.1,series='Seasonal Naive Model')+  
  autolayer(drift_model,PI=F,size=1.1,series='Drift Model')+  
  autolayer(test)
```



Exponential Smoothing Models

Forecasts are weighted averages of past observations with the weights decaying exponentially such that recent observations get weighted more than distant observations.

Simple exponential smoothing

- Forecasts are calculated using weighted averages, where the weights decrease exponentially. Most recent observations get the heaviest weight.
- Simplest of exponential smoothing methods
- Suitable for forecasting data with no clear trend or seasonal pattern.

Model and Forecast

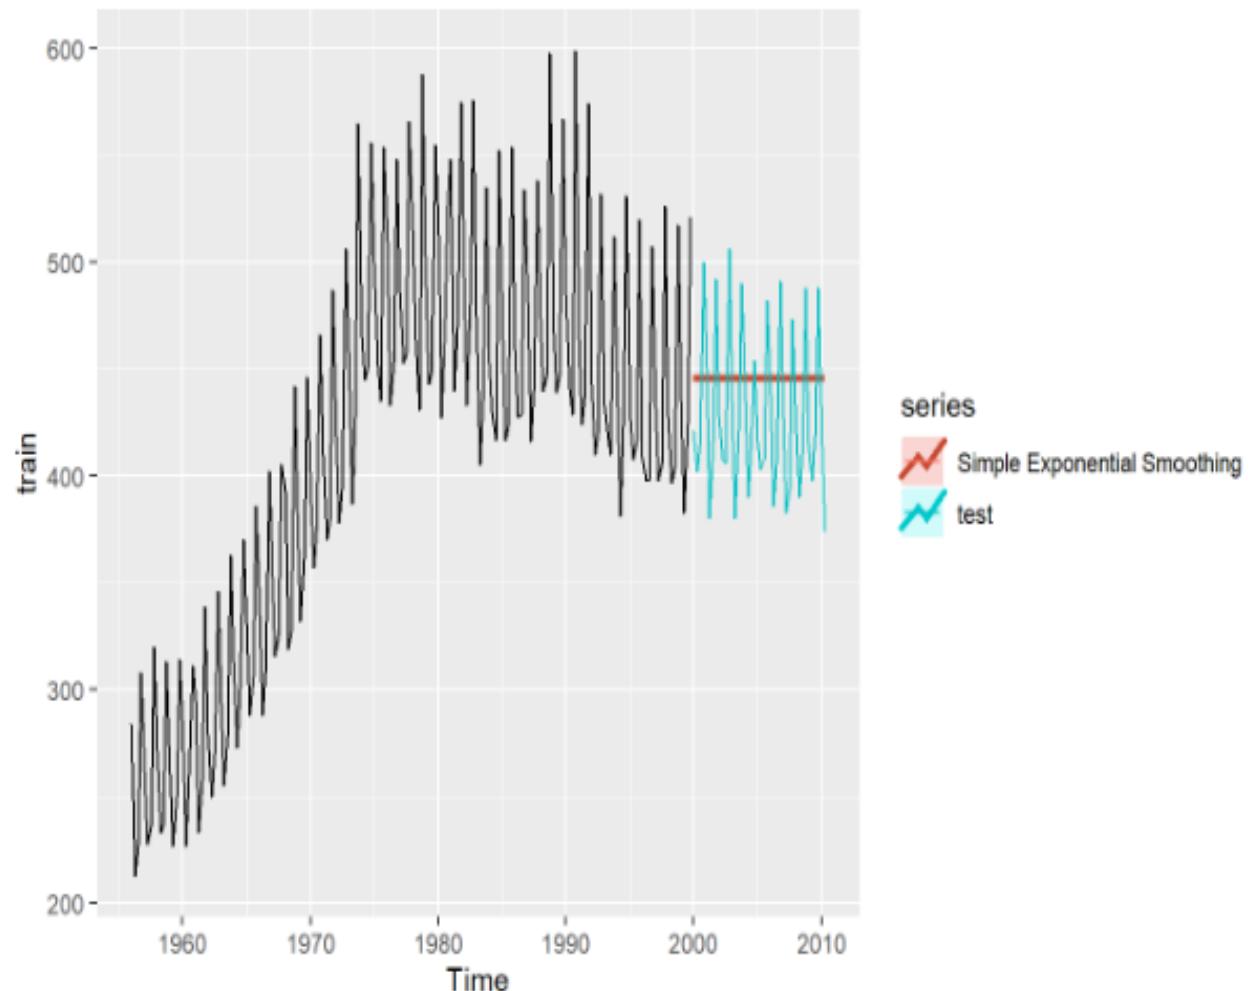
```
ses_model = ses(train,h = 42)
ses_model$mean
```

```
##          Qtr1     Qtr2     Qtr3     Qtr4
## 2000 445.7878 445.7878 445.7878 445.7878
## 2001 445.7878 445.7878 445.7878 445.7878
## 2002 445.7878 445.7878 445.7878 445.7878
## 2003 445.7878 445.7878 445.7878 445.7878
## 2004 445.7878 445.7878 445.7878 445.7878
## 2005 445.7878 445.7878 445.7878 445.7878
## 2006 445.7878 445.7878 445.7878 445.7878
## 2007 445.7878 445.7878 445.7878 445.7878
## 2008 445.7878 445.7878 445.7878 445.7878
## 2009 445.7878 445.7878 445.7878 445.7878
## 2010 445.7878 445.7878
```

- Exponential smoothing method has the same value for all forecast
- It is the estimated mean of the future possible sample path.

Visualize Forecasts

```
autoplot(train)+  
  autolayer(ses_model,series = "Simple Exponential Smoothing",PI = F, size=1.1)+  
  autolayer(test)
```



Accuracy

```
accuracy(ses_model,x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE  
## Training set 6.78376 52.86562 42.65547 0.5391999 10.07522 2.670819  
## Test set     -18.09737 41.26471 37.68469 -4.9785614 8.97363 2.359580  
##  
##                  ACF1 Theil's U  
## Training set -0.05109356       NA  
## Test set     -0.05298108 0.7972014
```

- Exponential smoothing method with trend

Holt's Method

Extends simple exponential smoothing to allow the forecasting of data with a trend

Model and Forecast

```
holt_model = holt(train,h=42)
holt_model$mean
```

```
##          Qtr1     Qtr2     Qtr3     Qtr4
## 2000 443.8257 448.1301 452.4345 456.7389
## 2001 461.0433 465.3477 469.6522 473.9566
## 2002 478.2610 482.5654 486.8698 491.1742
## 2003 495.4786 499.7830 504.0874 508.3918
## 2004 512.6962 517.0006 521.3051 525.6095
## 2005 529.9139 534.2183 538.5227 542.8271
## 2006 547.1315 551.4359 555.7403 560.0447
## 2007 564.3491 568.6535 572.9579 577.2624
## 2008 581.5668 585.8712 590.1756 594.4800
## 2009 598.7844 603.0888 607.3932 611.6976
## 2010 616.0020 620.3064
```

Accuracy

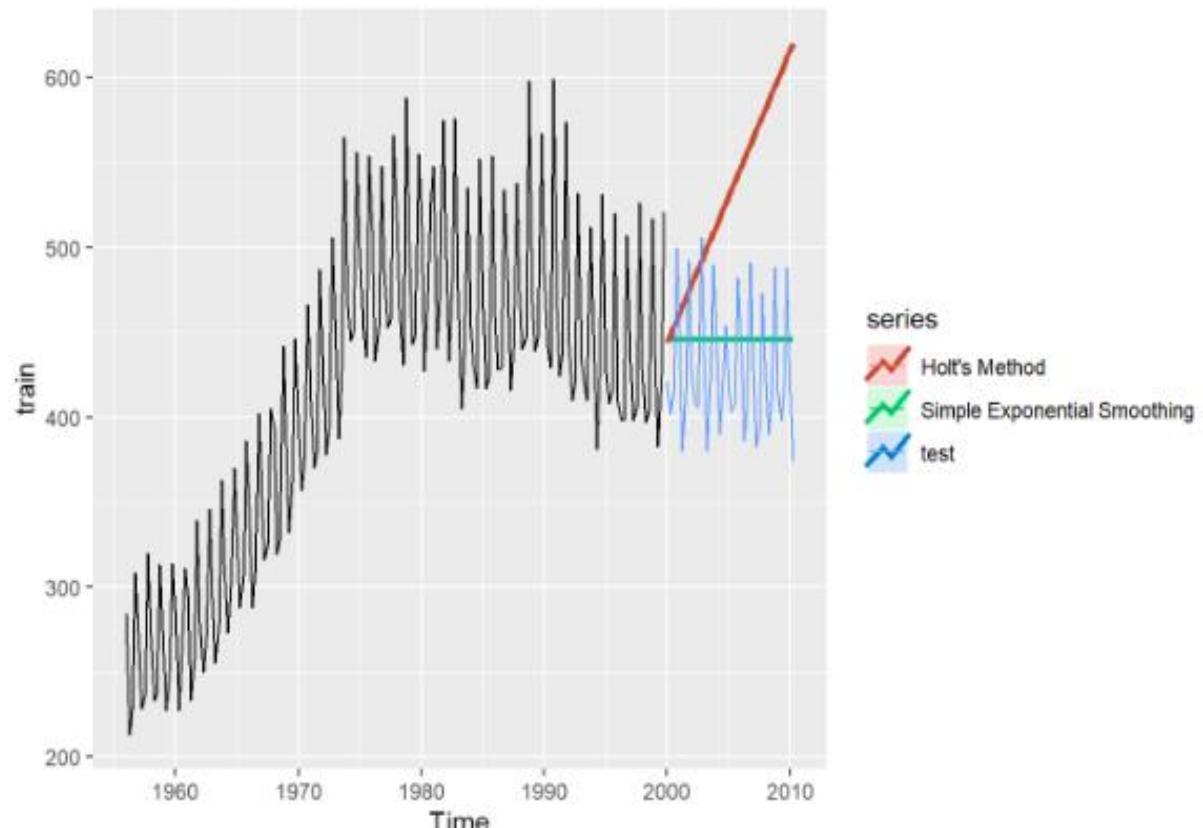
```
accuracy(holt_model,x=ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 1.048297 51.06394 41.34716 -0.8878183 9.914797 2.588901
## Test set    -104.375595 124.04137 108.00085 -25.4057657 26.131936 6.762339
##                  ACF1 Theil's U
## Training set -0.03671365        NA
## Test set      0.61128212 2.375372
```

- Holt's method will produce forecasts where the trend continues at the same slope indefinitely into the future.

Visualize Forecast

```
autoplot(train) +
  autolayer(ses_model,series = "Simple Exponential Smoothing",PI = F, size=1.1) +
  autolayer(holt_model,series="Holt's Method",PI=F,size=1.1) +
  autolayer(test)
```



Holt's Method with Damping

Forecasts generally display a constant trend (increasing or decreasing) indefinitely into the future. For this reason, a damping parameter is usually included

Model and Forecast

```
holt_damped_model = holt(train,h=42,damped = T)  
holt_damped_model$mean
```

```
##          Qtr1     Qtr2     Qtr3     Qtr4  
## 2000 441.6555 443.9208 445.9752 447.8384  
## 2001 449.5282 451.0607 452.4506 453.7112  
## 2002 454.8544 455.8912 456.8315 457.6844  
## 2003 458.4578 459.1593 459.7954 460.3724  
## 2004 460.8957 461.3702 461.8006 462.1910  
## 2005 462.5450 462.8660 463.1572 463.4213  
## 2006 463.6608 463.8780 464.0750 464.2537  
## 2007 464.4157 464.5627 464.6960 464.8168  
## 2008 464.9265 465.0259 465.1161 465.1978  
## 2009 465.2720 465.3393 465.4003 465.4556  
## 2010 465.5058 465.5513
```

Accuracy

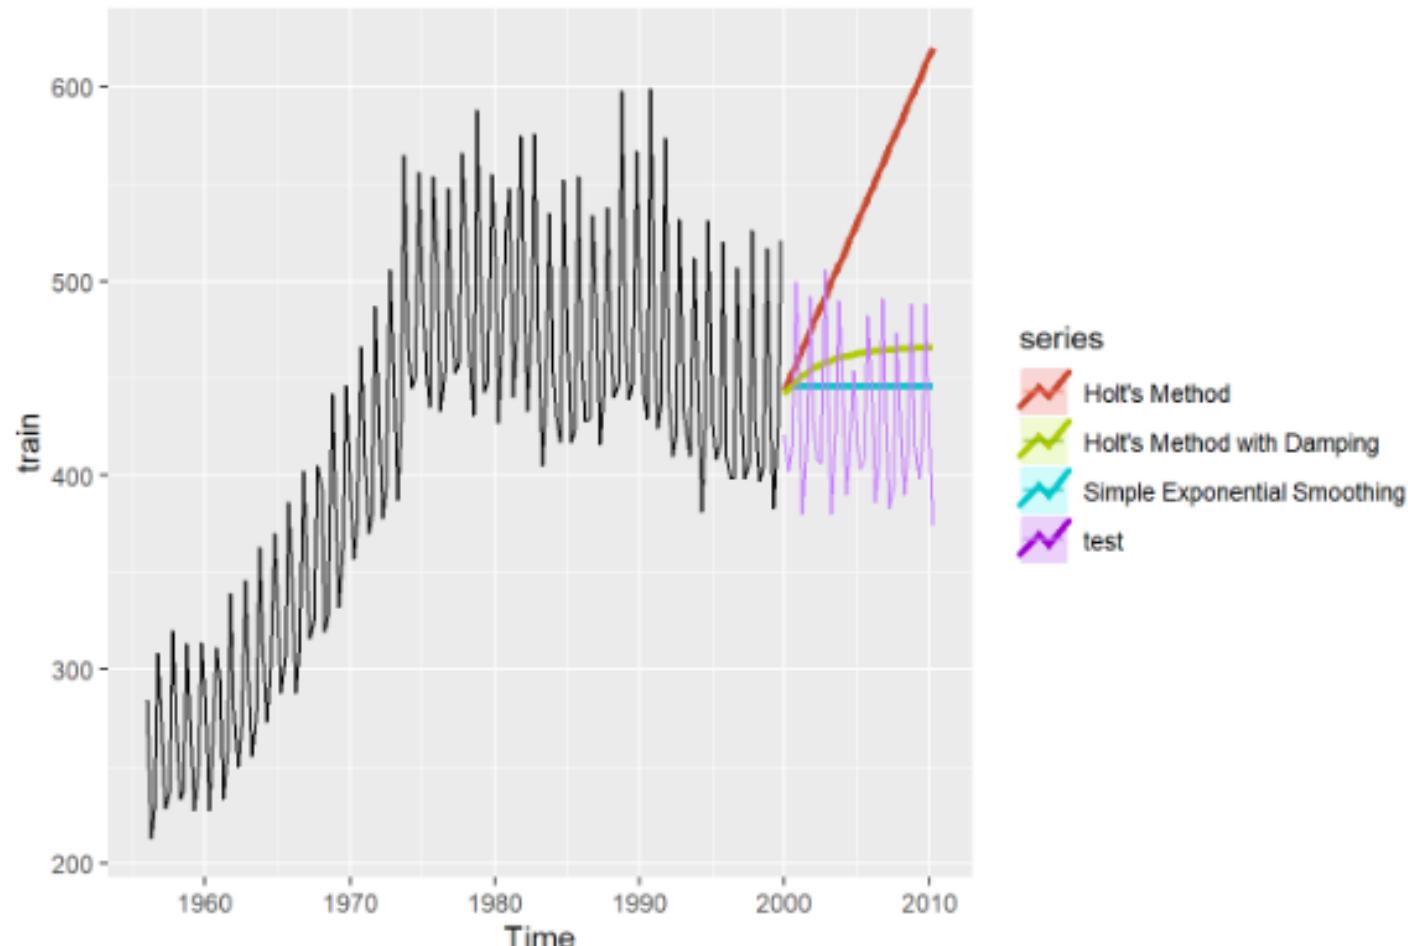
```
accuracy(holt_damped_model,x=ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE  
## Training set 3.113352 50.01115 40.99473 -0.3297766 9.791311 2.566834  
## Test set    -32.180549 49.84832 44.97675 -8.3039388 10.897279 2.816164  
##                         ACF1 Theil's U  
## Training set -0.059527825      NA  
## Test set     -0.005103579 0.9593471
```

- Dampen the trend over time so as to level off to a constant value.
- Under this method, the short run forecast is trended but the long run forecast is constant

Visualize Forecast

```
autoplot(train) +  
  autolayer(ses_model,series = "Simple Exponential Smoothing",PI = F, size=1.1)+  
  autolayer(holt_model,series="Holt's Method",PI=F,size=1.1)+  
  autolayer(holt_damped_model,series="Holt's Method with Damping",PI=F,size=1.1)+  
  autolayer(test)
```



Holt-Winter's seasonal method

Extends Holt's method to capture seasonality. Two types:

- Additive method is used when seasonal variations are roughly constant
- Multiplicative method is used when seasonal variations change in proportion to the level of the series

Holt_Winter's Additive

Model and Forecast

```
hw_additive = hw(train,h=42,seasonal = 'additive', damped=T)  
hw_additive$mean
```

```
##          Qtr1     Qtr2     Qtr3     Qtr4  
## 2000 431.7122 396.8446 412.7531 524.0245  
## 2001 433.3171 398.1286 413.7803 524.8462  
## 2002 433.9745 398.6545 414.2010 525.1828  
## 2003 434.2438 398.8699 414.3733 525.3207  
## 2004 434.3541 398.9581 414.4439 525.3771  
## 2005 434.3992 398.9943 414.4728 525.4003  
## 2006 434.4177 399.0091 414.4847 525.4097  
## 2007 434.4253 399.0151 414.4895 525.4136  
## 2008 434.4284 399.0176 414.4915 525.4152  
## 2009 434.4297 399.0186 414.4923 525.4159  
## 2010 434.4302 399.0190
```

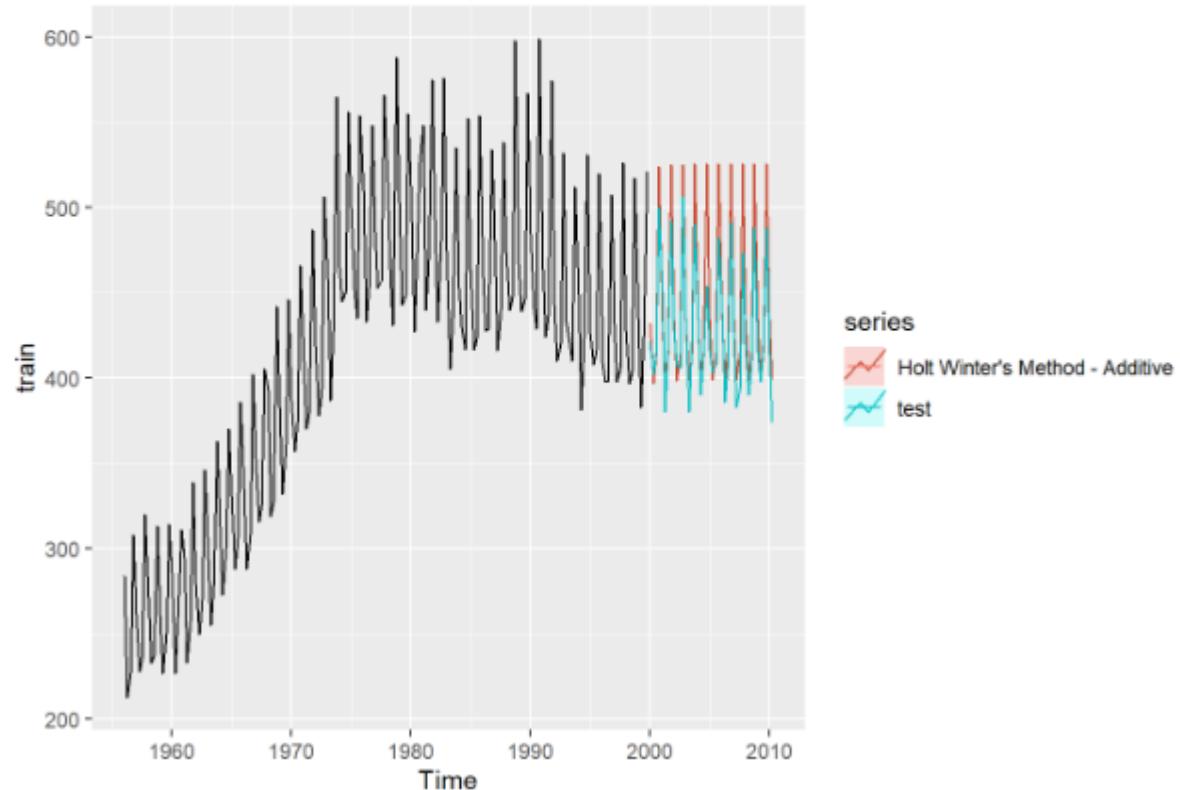
Accuracy

```
accuracy(hw_additive,x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE  
## Training set 1.339248 16.02895 12.41005 0.2859756 3.005731 0.7770402  
## Test set    -14.058340 22.67754 16.71429 -3.1512303 3.775641 1.0465444  
##                  ACF1 Theil's U  
## Training set -0.1816207      NA  
## Test set     -0.1857105 0.4430322
```

Visualize Forecasts

```
autoplot(train)+  
  autolayer(hw_additive,series="Holt Winter's Method - Additive",PI=F)+  
  autolayer(test)
```



- Exponential smoothing method with trend and seasonality (Additive version)
- When seasonal variation increases with the level of the series, multiplicative method should be used instead of additive method.

- Exponential smoothing method with trend and seasonality (Multiplicative version)

Holt_Winter's Multiplicative

Model and Forecast

```
hw_multiplicative = hw(train,h=42,seasonal = 'multiplicative', damped=T)
hw_multiplicative$mean
```

```
##          Qtr1      Qtr2      Qtr3      Qtr4
## 2000 432.4894 398.2671 413.6034 524.9590
## 2001 435.1293 400.2185 415.2329 526.6252
## 2002 436.2614 401.0644 415.9487 527.3690
## 2003 436.7758 401.4576 416.2903 527.7351
## 2004 437.0374 401.6654 416.4787 527.9464
## 2005 437.1954 401.7973 416.6044 528.0945
## 2006 437.3110 401.8981 416.7043 528.2166
## 2007 437.4093 401.9861 416.7938 528.3281
## 2008 437.5005 402.0690 416.8789 528.4353
## 2009 437.5887 402.1497 416.9623 528.5407
## 2010 437.6758 402.2296
```

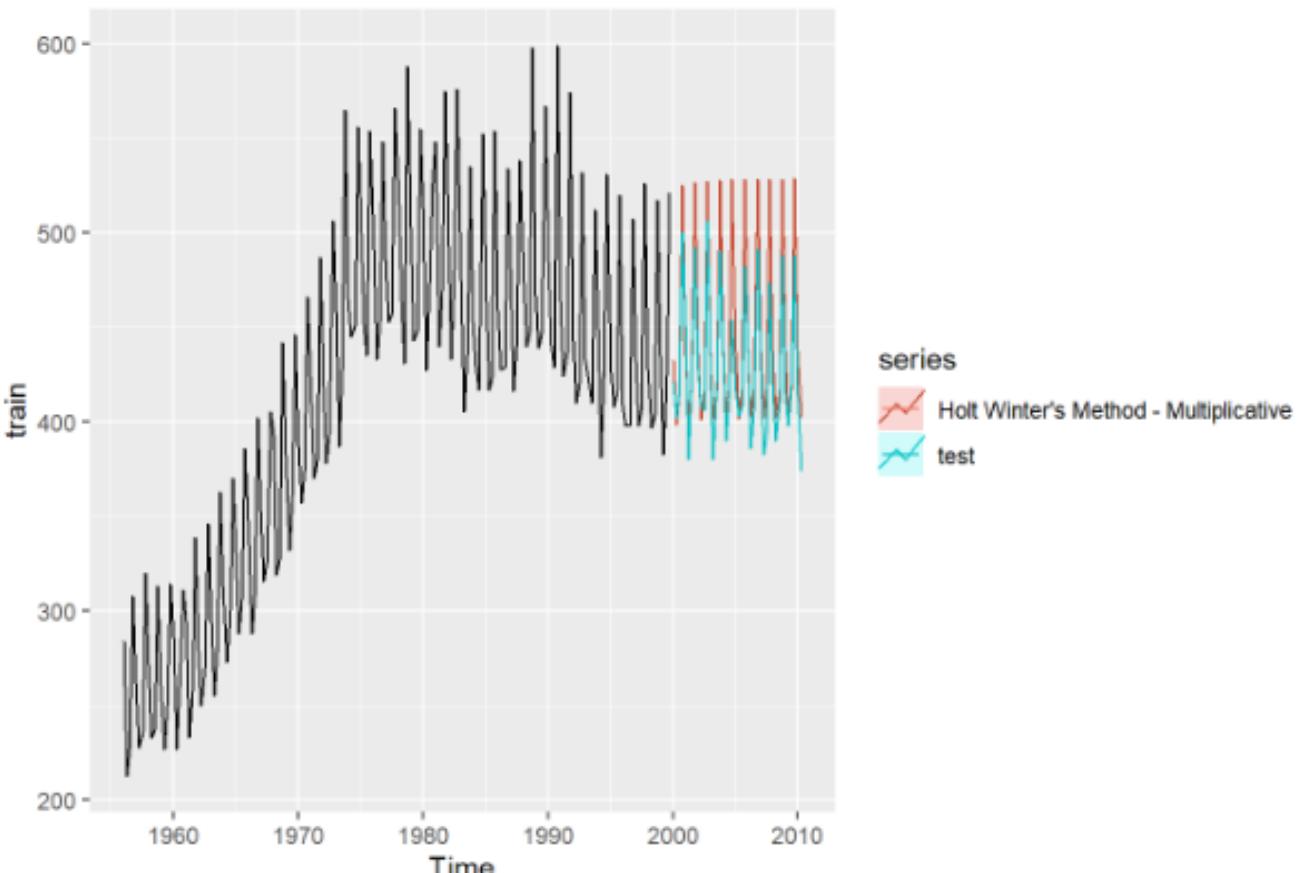
Accuracy

```
accuracy(hw_multiplicative,x=ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 1.259767 15.55545 11.94622 0.314631 2.887841 0.7479979
## Test set     -16.474394 24.38794 18.20496 -3.721500 4.125710 1.1398809
##               ACF1 Theil's U
## Training set -0.2232171      NA
## Test set     -0.1694875 0.475703
```

Visualize Forecast

```
autoplot(train) +
  autolayer(hw_multiplicative, series="Holt Winter's Method - Multiplicative", PI=F) +
  autolayer(test)
```



ETS Models

But, exponential smoothing methods are not limited to the ones examined so far. By considering variations of trend (none, additive, additive damped) and seasonal components (non, additive, multiplicative), there are nine exponential smoothing methods. And, for each of these methods, errors may be additive or multiplicative.

ETS models in R are handled by `ets()` from `library(forecast)`. Unlike functions such as `naive()`, `ses()`, `hw()` functions, the `ets()` function does not produce forecasts. Rather, it estimates the model parameters and returns information on the fitted model.

ETS: AAA

Let us fit an exponential smoothing model where the Errors are Additive, Trend is Additive and Seasonal component is Additive. Note the model argument in `ets()`.

Model

```
ets_aaa = ets(train,model = 'AAA')
```

The ETS model object has methods like other predictive modelling objects such as `lm()`. Let us examine the coefficients of the model

```
coef(ets_aaa)
```

```
##      alpha      beta     gamma      1      b      s0
## 0.2284115 0.0374727 0.2029944 259.2792901 0.2011150 52.7137253
##      s1      s2
## -20.0075675 -41.6237261
```

ETS Model: AAA

Let us fit an exponential smoothing model where the Errors are Additive, Trend is Additive and Seasonal component is Additive. Note the model argument in `ets()`.

```
ets_aaa = ets(train,model = 'AAA')
```

The ETS model object has methods like other predictive modelling objects such as `lm()`. Let us examine the coefficients of the model

```
coef(ets_aaa)
```

```
##      alpha      beta     gamma      1      b      s0
## 0.2284115 0.0374727 0.2029944 259.2792901 0.2011150 52.7137253
##      s1      s2
## -20.0075675 -41.6237261
```

Each exponential smoothing method described previously could be written as in "innovations state space model".

`ets()` provides a way of selecting the best model for a particular time series. It search for model with minimum AICc

`ets()` does not provide forecast, it only estimate the parameters, we will need to pass the estimated results to the `forecast` function

ETS Model: AAA

Table 7.7: State space equations for each of the models in the ETS framework.

ADDITIVE ERROR MODELS

Trend	Seasonal		
	N	A	M
N	$y_t = \ell_{t-1} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \alpha \varepsilon_t$	$y_t = \ell_{t-1} + s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \alpha \varepsilon_t$ $s_t = s_{t-m} + \gamma \varepsilon_t$	$y_t = \ell_{t-1} s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \alpha \varepsilon_t / s_{t-m}$ $s_t = s_{t-m} + \gamma \varepsilon_t / \ell_{t-1}$
A	$y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$ $b_t = b_{t-1} + \beta \varepsilon_t$	$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$ $b_t = b_{t-1} + \beta \varepsilon_t$ $s_t = s_{t-m} + \gamma \varepsilon_t$	$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t / s_{t-m}$ $b_t = b_{t-1} + \beta \varepsilon_t / s_{t-m}$ $s_t = s_{t-m} + \gamma \varepsilon_t / (\ell_{t-1} + b_{t-1})$
A _d	$y_t = \ell_{t-1} + \phi b_{t-1} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha \varepsilon_t$ $b_t = \phi b_{t-1} + \beta \varepsilon_t$	$y_t = \ell_{t-1} + \phi b_{t-1} + s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha \varepsilon_t$ $b_t = \phi b_{t-1} + \beta \varepsilon_t$ $s_t = s_{t-m} + \gamma \varepsilon_t$	$y_t = (\ell_{t-1} + \phi b_{t-1}) s_{t-m} + \varepsilon_t$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha \varepsilon_t / s_{t-m}$ $b_t = \phi b_{t-1} + \beta \varepsilon_t / s_{t-m}$ $s_t = s_{t-m} + \gamma \varepsilon_t / (\ell_{t-1} + \phi b_{t-1})$

MULTIPLICATIVE ERROR MODELS

Trend	Seasonal		
	N	A	M
N	$y_t = \ell_{t-1}(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1}(1 + \alpha \varepsilon_t)$	$y_t = (\ell_{t-1} + s_{t-m})(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1} + \alpha(\ell_{t-1} + s_{t-m})\varepsilon_t$ $s_t = s_{t-m} + \gamma(\ell_{t-1} + s_{t-m})\varepsilon_t$	$y_t = \ell_{t-1} s_{t-m}(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1}(1 + \alpha \varepsilon_t)$ $s_t = s_{t-m}(1 + \gamma \varepsilon_t)$
A	$y_t = (\ell_{t-1} + b_{t-1})(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1})\varepsilon_t$	$y_t = (\ell_{t-1} + b_{t-1} + s_{t-m})(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1} + b_{t-1} + \alpha(\ell_{t-1} + b_{t-1} + s_{t-m})\varepsilon_t$ $b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1} + s_{t-m})\varepsilon_t$ $s_t = s_{t-m} + \gamma(\ell_{t-1} + b_{t-1} + s_{t-m})\varepsilon_t$	$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m}(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1})\varepsilon_t$ $s_t = s_{t-m}(1 + \gamma \varepsilon_t)$
A _d	$y_t = (\ell_{t-1} + \phi b_{t-1})(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + \phi b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = \phi b_{t-1} + \beta(\ell_{t-1} + \phi b_{t-1})\varepsilon_t$	$y_t = (\ell_{t-1} + \phi b_{t-1} + s_{t-m})(1 + \varepsilon_t)$ $\ell_t = \ell_{t-1} + \phi b_{t-1} + \alpha(\ell_{t-1} + \phi b_{t-1} + s_{t-m})\varepsilon_t$ $b_t = \phi b_{t-1} + \beta(\ell_{t-1} + \phi b_{t-1} + s_{t-m})\varepsilon_t$ $s_t = s_{t-m} + \gamma(\ell_{t-1} + \phi b_{t-1} + s_{t-m})\varepsilon_t$	$y_t = (\ell_{t-1} + \phi b_{t-1}) s_{t-m}(1 + \varepsilon_t)$ $\ell_t = (\ell_{t-1} + \phi b_{t-1})(1 + \alpha \varepsilon_t)$ $b_t = \phi b_{t-1} + \beta(\ell_{t-1} + \phi b_{t-1})\varepsilon_t$ $s_t = s_{t-m}(1 + \gamma \varepsilon_t)$

Summary of the model. Note, since we did not specify the damping term, ets() tried both an Additive-Damped trend and just an Additived trend and went with the option that led to a lower AICc.

summary(ets_aaa)

```
## ETS(A,A,A)
## 
## Call:
##   ets(y = train, model = "AAA")
## 
## Smoothing parameters:
##   alpha = 0.2284
##   beta  = 0.0375
##   gamma = 0.203
## 
## Initial states:
##   l = 259.2793
##   b = 0.2011
##   s = 52.7137 -20.0076 -41.6237 8.9176
## 
## sigma: 16.5039
## 
##          AIC    AICc     BIC
## 1906.683 1907.768 1935.218
## 
## Training set error measures:
##          ME    RMSE     MAE      MPE     MAPE     MASE
## Training set 0.0339053 16.12443 12.43517 -0.02612254 2.989648 0.7786129
## 
##          ACF1
## Training set -0.2047428
```

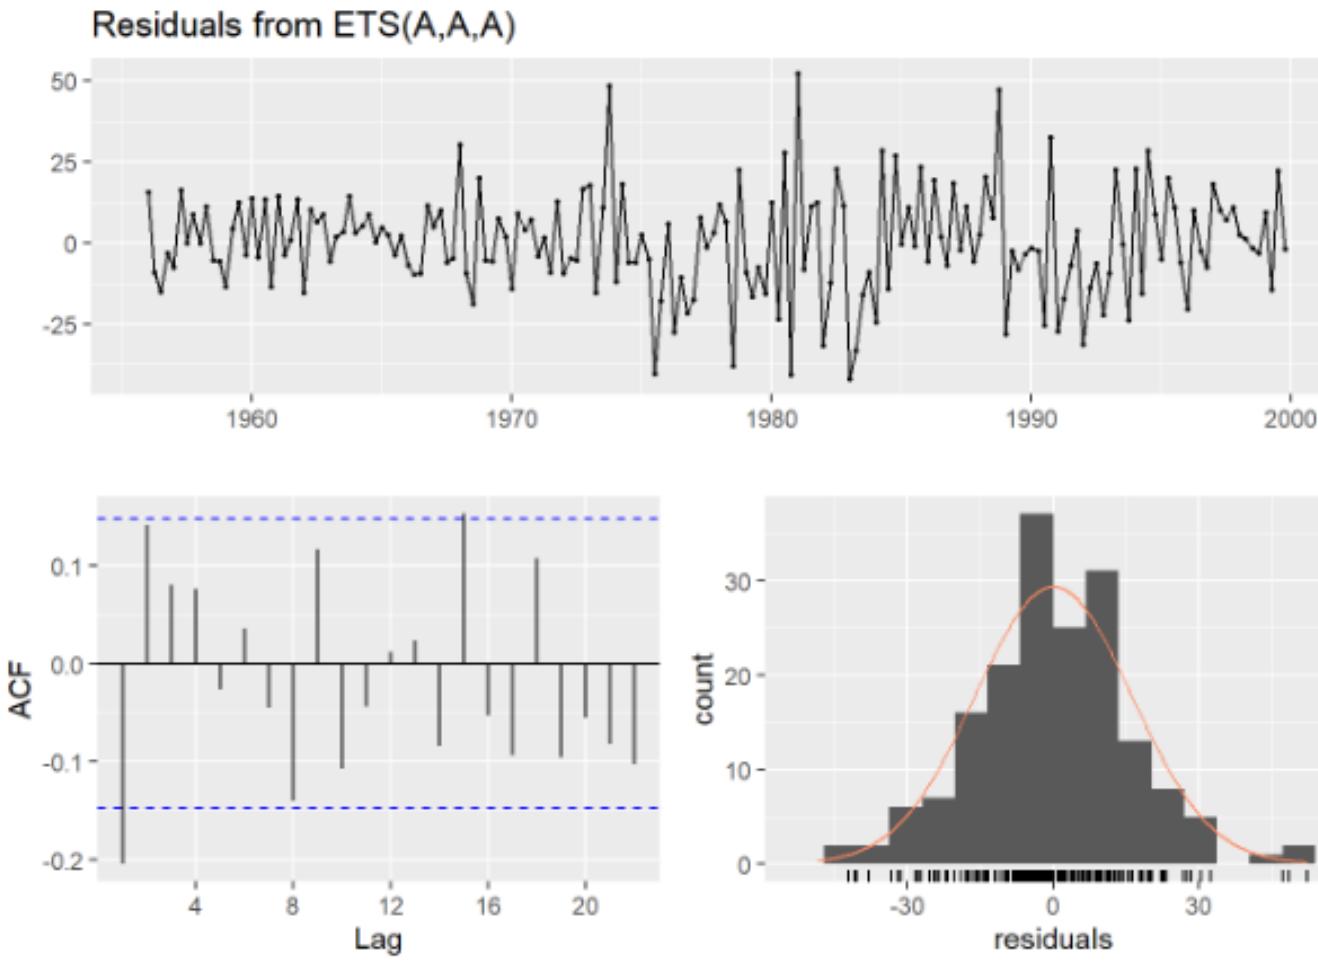
- The model with the minimum value of the AIC (AICc) is often the best model for forecasting.
- The AICc corrected for small sample bias

ETS Model: AAA

Examine the residuals.

Based on the ACF plot and a significant Ljung-Box test, residuals are not like white noise.

```
checkresiduals(ets_aaa)
```



- All of these methods for checking residuals are conveniently packaged into one R function `checkresiduals()`, which will produce a time plot, ACF plot and histogram of the residuals (with an overlaid normal distribution for comparison), and do a Ljung-Box test with the correct degrees of freedom.
- A small p-value (for instance, $p\text{-value} < .05$) indicates the possibility of non-zero autocorrelation within the first m lags. That is, a small p-value indicates that the data are probably not white noise.

H_0 : the autocorrelations up to lag k are all 0
 H_A : the autocorrelations of one or more lags differ from 0.

```
##  
## Ljung-Box test  
##  
## data: Residuals from ETS(A,A,A)  
## Q* = 22.82, df = 3, p-value = 4.403e-05  
##  
## Model df: 8. Total lags used: 11
```

Residuals are not white noise

ETS Model: AAA

Forecast

```
ets_aaa_forecast = forecast(ets_aaa,h=42)
ets_aaa_forecast
```

```
##           Point Forecast    Lo 80     Hi 80    Lo 95     Hi 95
## 2000 Q1      431.7860 410.6354 452.9365 399.4390 464.1330
## 2000 Q2      396.8128 374.9274 418.6982 363.3419 430.2836
## 2000 Q3      412.7249 389.9184 435.5315 377.8453 447.6045
## 2000 Q4      523.7759 499.8572 547.6946 487.1954 560.3564
## 2001 Q1      433.4849 406.5913 460.3785 392.3547 474.6151
## 2001 Q2      398.5117 370.2168 426.8065 355.2385 441.7849
## 2001 Q3      414.4238 384.5491 444.2985 368.7344 460.1133
## 2001 Q4      525.4748 493.8485 557.1011 477.1066 573.8431
## 2002 Q1      435.1838 399.9788 470.3888 381.3424 489.0251
## 2002 Q2      400.2106 363.0282 437.3930 343.3450 457.0762
## 2002 Q3      416.1227 376.8129 455.4326 356.0035 476.2420
## 2002 Q4      527.1737 485.5943 568.7531 463.5835 590.7640
## 2003 Q1      436.8827 391.3182 482.4472 367.1978 506.5676
## 2003 Q2      401.9095 353.8973 449.9217 328.4811 475.3379
## 2003 Q3      417.8217 367.2366 468.4067 340.4586 495.1848
## 2003 Q4      528.8726 475.5961 582.1492 447.3932 610.3520
## 2004 Q1      438.5816 381.0154 496.1478 350.5416 526.6216
## 2004 Q2      403.6084 343.2020 464.0148 311.2248 495.9920
## 2004 Q3      419.5206 356.1673 482.8739 322.6301 516.4111
## 2004 Q4      530.5715 464.1694 596.9737 429.0182 632.1249
```

the point forecast of the stock price for Q4, 2004

```
## 2005 Q1      440.2805 369.3366 511.2244 331.7812 548.7799
## 2005 Q2      405.3073 331.1844 479.4303 291.9460 518.6686
## 2005 Q3      421.2195 343.8240 498.6149 302.8533 539.5856
## 2005 Q4      532.2705 451.5126 613.0283 408.7620 655.7789
## 2006 Q1      441.9794 356.4573 527.5015 311.1846 572.7742
## 2006 Q2      407.0062 318.0039 496.0086 270.8888 543.1236
## 2006 Q3      422.9184 330.3522 515.4846 281.3506 564.4862
## 2006 Q4      533.9694 437.7585 630.1803 386.8275 681.1112
## 2007 Q1      443.6783 342.4990 544.8577 288.9379 598.4188
## 2007 Q2      408.7051 303.7720 513.6383 248.2238 569.1865
## 2007 Q3      424.6173 315.8543 533.3803 258.2787 590.9559
## 2007 Q4      535.6683 423.0016 648.3350 363.3594 707.9771
## 2008 Q1      445.3772 327.5506 563.2039 265.1769 625.5775
## 2008 Q2      410.4040 288.5712 532.2369 224.0768 596.7313
## 2008 Q3      426.3162 300.4071 552.2253 233.7548 618.8776
## 2008 Q4      537.3672 407.3135 667.4209 338.4672 736.2672
## 2009 Q1      447.0761 311.6804 582.4719 240.0062 654.1461
## 2009 Q2      412.1030 272.4654 551.7405 198.5457 625.6602
## 2009 Q3      428.0151 284.0704 571.9598 207.8707 648.1595
## 2009 Q4      539.0661 390.7505 687.3816 312.2370 765.8951
## 2010 Q1      448.7751 294.9428 602.6073 213.5089 684.0412
## 2010 Q2      413.8019 255.5059 572.0979 171.7090 655.8947
```

- RMSE of the ets_aaa model on the test sample

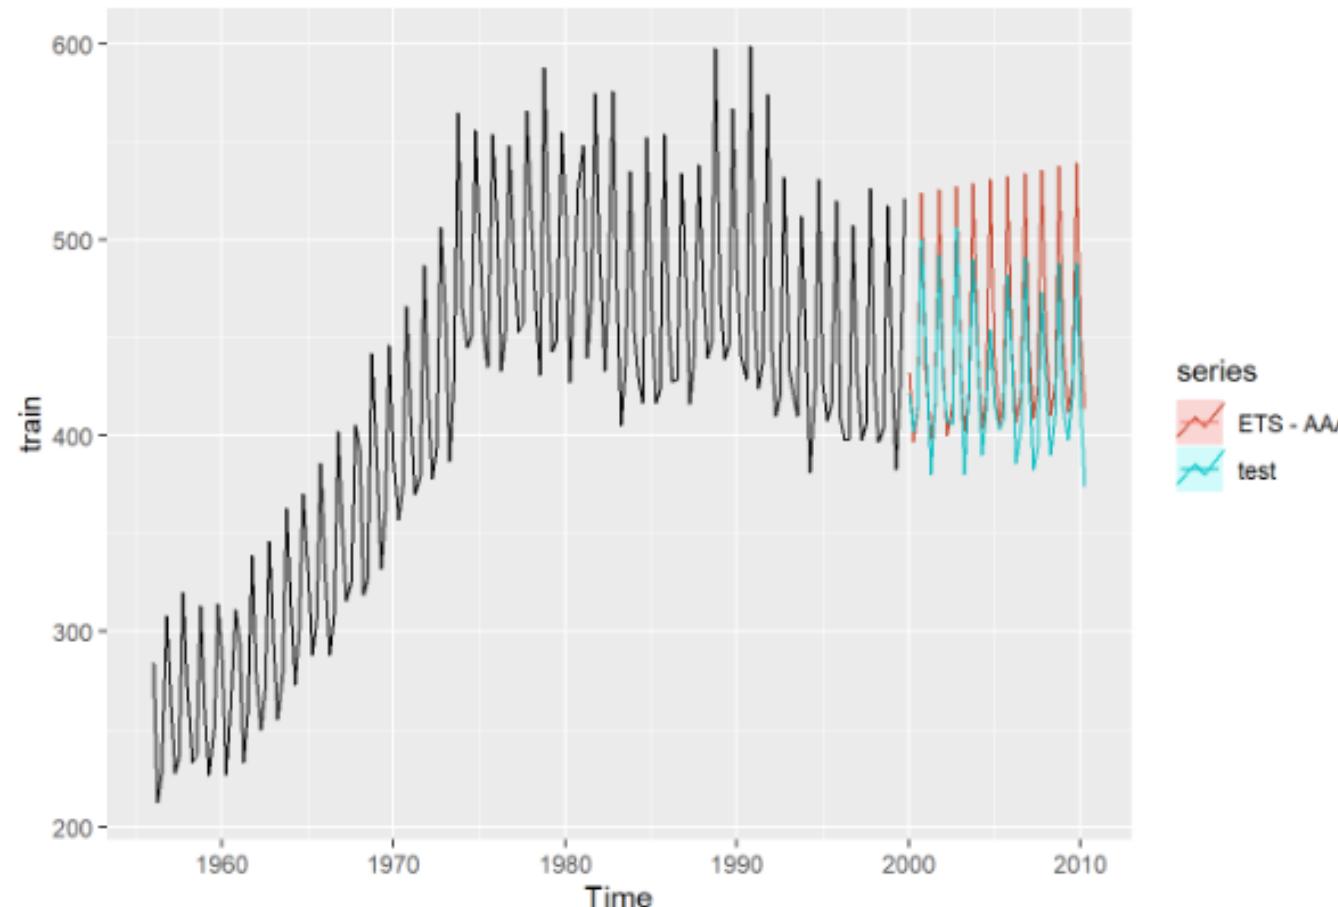
Accuracy

```
accuracy(ets_aaa_forecast,x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  0.0339053 16.12443 12.43517 -0.02612254 2.989648 0.7786129
## Test set      -20.3899116 28.22830 22.12902 -4.65481860 5.060769 1.3855813
##                               ACF1 Theil's U
## Training set -0.20474276          NA
## Test set      0.01545558 0.5492527
```

Visualize Forecast

```
autoplot(train)+  
  autolayer(ets_aaa_forecast, series="ETS - AAA", PI=F)+  
  autolayer(test)
```



ETS: Automatic Selection

When only the time-series is specified, and all other arguments are left at their default values, then `ets()` will automatically select the best model based on AICc. Compare the AICc for this model to ETS(A,A,A) above. Note that the chosen model has Multiplicative Errors, Additive Trend, and Multiplicative Seasonal.

Model

```
ets_auto = ets(train)
summary(ets_auto)
```

```
## ETS(M,A,M)
##
## Call:
##   ets(y = train)
##
## Smoothing parameters:
##   alpha = 0.2263
##   beta  = 0.035
##   gamma = 0.1766
##
## Initial states:
##   l = 255.5524
##   b = 0.8861
##   s = 1.1819 0.9125 0.8628 1.0428
##
## sigma: 0.0371
##
##      AIC    AICc      BIC
## 1869.636 1870.720 1898.170
##
## Training set error measures:
##           ME    RMSE      MAE      MPE      MAPE      MASE
## Training set -0.2090563 15.8621 12.12096 -0.04973312 2.91519 0.758939
##           ACF1
## Training set -0.1999151
```

ETS Model: Automatic Selection

Forecast

```
ets_auto_forecast = forecast(ets_auto,h=42)
```

Accuracy

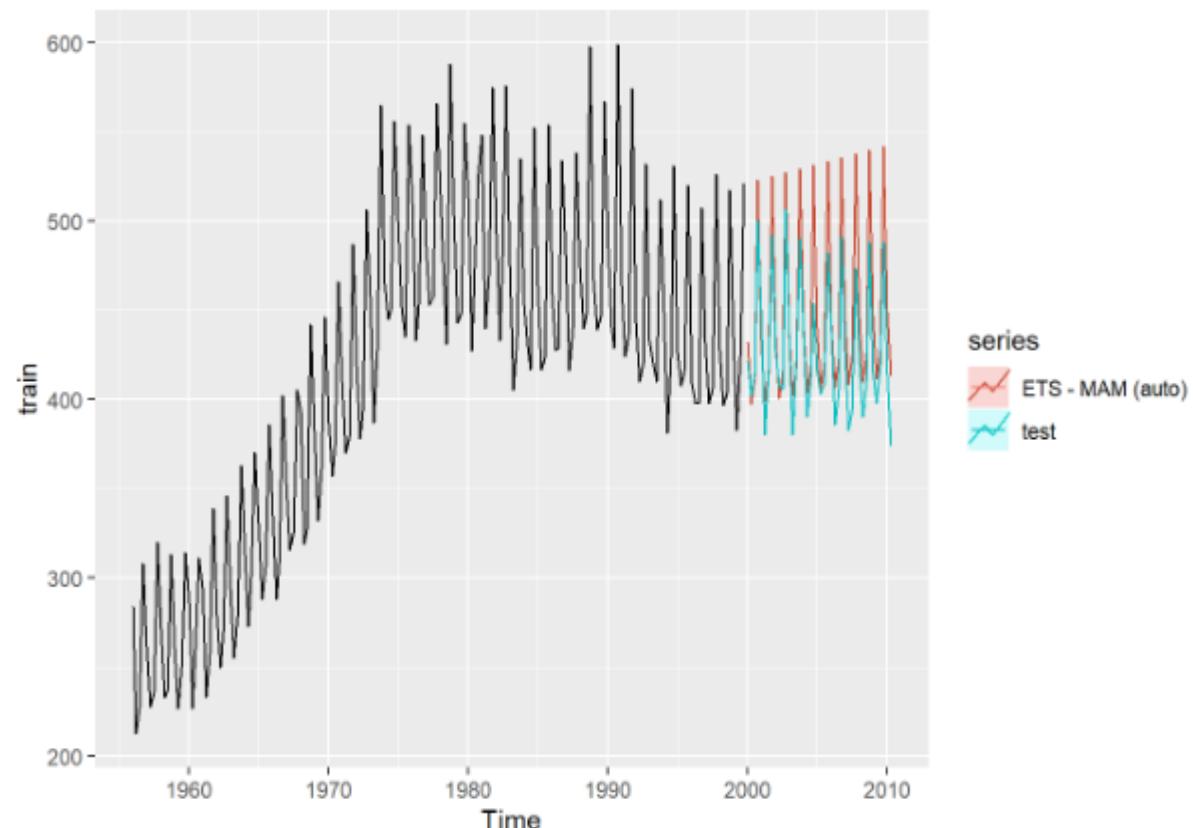
Even though ets_auto has a lower AICc, its RMSE is worse than that of ETS(A,A,A)

```
accuracy(ets_auto_forecast,x = ausbeer)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.2090563 15.86210 12.12096 -0.04973312 2.915190 0.758939
## Test set     -20.4860097 28.61288 22.23155 -4.66935195 5.076505 1.392001
##               ACF1 Theil's U
## Training set -0.1999151      NA
## Test set      0.0102163  0.556942
```

Visualize Forecast

```
autoplot(train) +
  autolayer(ets_auto_forecast,series="ETS - MAM (auto)",PI=F) +
  autolayer(test)
```



ARIMA

- Exponential Smoothing and ARIMA are the two most widely used approaches to time-series forecasting and provide complementary approaches to the problem.
- While Exponential Smoothing models are based on a description of trend and seasonality in the data, ARIMA models aim to describe autocorrelations in the data

Stationary Process

Since an ARMA process assumes the data is stationary, let us examine the assumption of stationarity and how to satisfy it.

- Stationary Process has constant mean, variance and covariance (between identically spaced data points) over time. Stationarity is an assumption of many time-series analysis procedures. However, most time-series are non-stationary
- Fortunately, a non-stationary process can be transformed into a (weakly) stationary process through transformations that remove trend and stabilize variance.
- We will stabilize variance using a Box-Cox Transformation and remove seasonality and trend using differencing

ARIMA: Stationary Process

- Here, you will use a Box-Cox transformation to stabilize the variance of the pre-loaded a10 series, which contains monthly anti-diabetic drug sales in Australia from 1991-2008.

Making a Series Stationary: Example 1

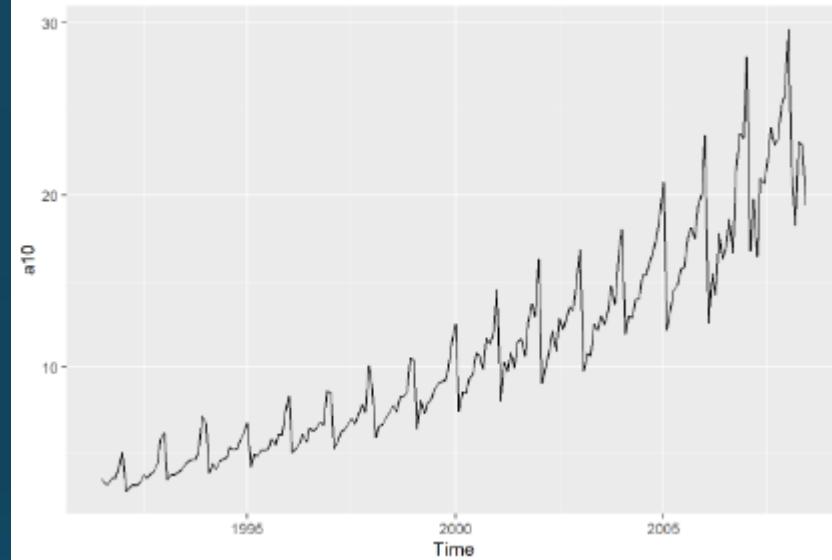
Original Series

a10



##	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1991							3.526591
## 1992	5.088335	2.814520	2.985811	3.204780	3.127578	3.270523	3.737851
## 1993	6.192068	3.450857	3.772307	3.734303	3.905399	4.049687	4.315566
## 1994	6.731473	3.841278	4.394076	4.075341	4.540645	4.645615	4.752607
## 1995	6.749484	4.216067	4.949349	4.823045	5.194754	5.170787	5.256742
## 1996	8.329452	5.069796	5.262557	5.597126	6.110296	5.689161	6.486849
## 1997	8.524471	5.277918	5.714303	6.214529	6.411929	6.667716	7.050831
## 1998	8.798513	5.918261	6.534493	6.675736	7.064201	7.383381	7.813496
## 1999	10.391416	6.421535	8.062619	7.297739	7.936916	8.165323	8.717420
## 2000	12.511462	7.457199	8.591191	8.474000	9.386803	9.560399	10.834295
## 2001	14.497581	8.049275	10.312891	9.753358	10.850382	9.961719	11.443601
## 2002	16.300269	9.053485	10.002449	10.788750	12.106705	10.954101	12.844566
## 2003	16.828350	9.800215	10.816994	10.654223	12.512323	12.161210	12.998046
## 2004	18.003768	11.938030	12.997900	12.882645	13.943447	13.989472	15.339097
## 2005	20.778723	12.154552	13.402392	14.459239	14.795102	15.705248	15.829550
## 2006	23.486694	12.536987	15.467018	14.233539	17.783058	16.291602	16.980282
## 2007	28.038383	16.763869	19.792754	16.427305	21.000742	20.681002	21.834890
## 2008	29.665356	21.654285	18.264945	23.107677	22.912510	19.431740	

autoplot(a10)



##	Aug	Sep	Oct	Nov	Dec
## 1991	3.180891	3.252221	3.611003	3.565869	4.306371
## 1992	3.558776	3.777202	3.924490	4.386531	5.810549
## 1993	4.562185	4.608662	4.667851	5.093841	7.179962
## 1994	5.350605	5.204455	5.301651	5.773742	6.204593
## 1995	5.855277	5.490729	6.115293	6.088473	7.416598
## 1996	6.300569	6.467476	6.828629	6.649078	8.606937
## 1997	6.704919	7.250988	7.819733	7.398101	10.096233
## 1998	7.431892	8.275117	8.260441	8.596156	10.558939
## 1999	9.070964	9.177113	9.251887	9.933136	11.532974
## 2000	10.643751	9.908162	11.710041	11.340151	12.079132
## 2001	11.659239	10.647060	12.652134	13.674466	12.965735
## 2002	12.196500	12.854748	13.542004	13.287640	15.134918
## 2003	12.517276	13.268658	14.733622	13.669382	16.503966
## 2004	15.370764	16.142005	16.685754	17.636728	18.869325
## 2005	17.554701	18.100864	17.496668	19.347265	20.031291
## 2006	18.612189	16.623343	21.430241	23.575517	23.334206
## 2007	23.930204	22.930357	23.263340	25.250030	25.806090
## 2008					

ARIMA: Stationary Process

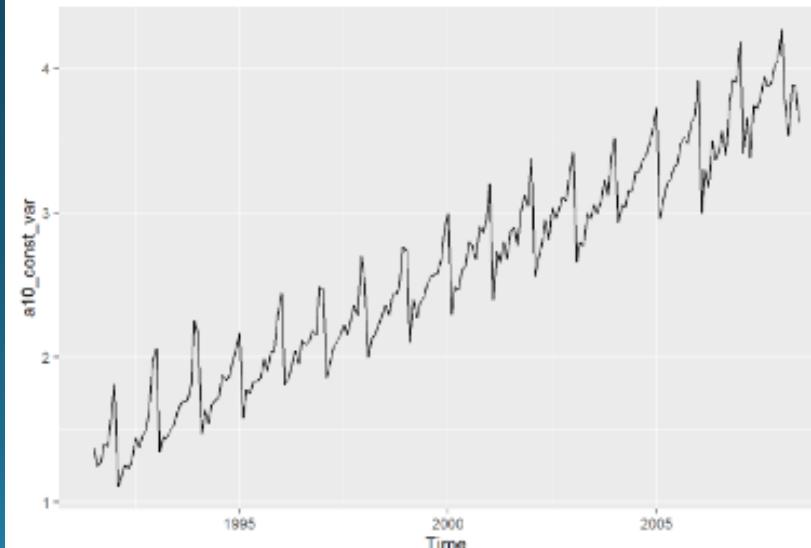
Stabilize Variance ←

```
a10_const_var = BoxCox(a10,lambda = BoxCox.lambda(a10))
a10_const_var
```

```
##          Jan      Feb      Mar      Apr      May      Jun      Jul      Aug
## 1991 1.370640 1.249719
## 1992 1.813838 1.108403 1.176346 1.258433 1.230064 1.282127 1.439555 1.381367
## 1993 2.060083 1.345059 1.450472 1.438425 1.491844 1.535336 1.612066 1.679652
## 1994 2.166791 1.472067 1.633938 1.542927 1.673877 1.701798 1.729698 1.876275
## 1995 2.170224 1.583846 1.779608 1.747770 1.839502 1.833762 1.854241 1.989287
## 1996 2.444272 1.809319 1.855616 1.932585 2.043207 1.953056 2.119355 2.082179
## 1997 2.474891 1.859242 1.958598 2.064685 2.104516 2.154573 2.226514 2.161715
## 1998 2.516905 2.002791 2.128714 2.156116 2.228963 2.286257 2.360153 2.294776
## 1999 2.740742 2.106427 2.401353 2.271098 2.380705 2.418016 2.504592 2.557562
## 2000 2.996340 2.299201 2.485227 2.467026 2.603388 2.628007 2.797657 2.773422
## 2001 3.203656 2.399174 2.730431 2.654922 2.799686 2.683470 2.872744 2.898487
## 2002 3.371455 2.554986 2.688989 2.791898 2.950612 2.812706 3.033019 2.960871
## 2003 3.417552 2.661387 2.795472 2.774764 2.996436 2.956847 3.049642 2.996988
## 2004 3.515807 2.931163 3.049626 3.037159 3.148428 3.153087 3.284117 3.287069
## 2005 3.727330 2.956087 3.092632 3.199894 3.232556 3.317933 3.329258 3.478946
## 2006 3.911279 2.999181 3.296010 3.177573 3.497791 3.370688 3.430582 3.564495
## 2007 4.182556 3.411991 3.655150 3.382661 3.743172 3.720310 3.801421 3.939630
## 2008 4.270269 3.788975 3.536880 3.886679 3.873874 3.627938
```

```
##          Sep      Oct      Nov      Dec
## 1991 1.275573 1.398595 1.383719 1.609482
## 1992 1.452016 1.497678 1.631851 1.979620
## 1993 1.692032 1.707642 1.815177 2.249997
## 1994 1.841819 1.864826 1.971617 2.062651
## 1995 1.908552 2.044244 2.038670 2.292095
## 1996 2.115532 2.185216 2.150983 2.487656
## 1997 2.262758 2.361199 2.288847 2.701625
## 1998 2.435630 2.433288 2.485993 2.762514
## 1999 2.573116 2.583979 2.679584 2.883464
## 2000 2.676181 2.904492 2.860244 2.947449
## 2001 2.773846 3.011932 3.120929 3.046157
## 2002 3.034127 3.107214 3.080547 3.264953
## 2003 3.078540 3.226626 3.120405 3.389388
## 2004 3.357387 3.405230 3.485739 3.584657
## 2005 3.523672 3.474122 3.621507 3.672894
## 2006 3.399808 3.773410 3.916994 3.901423
## 2007 3.875049 3.896824 4.021381 4.054717
## 2008
```

```
autoplot(a10_const_var)
```

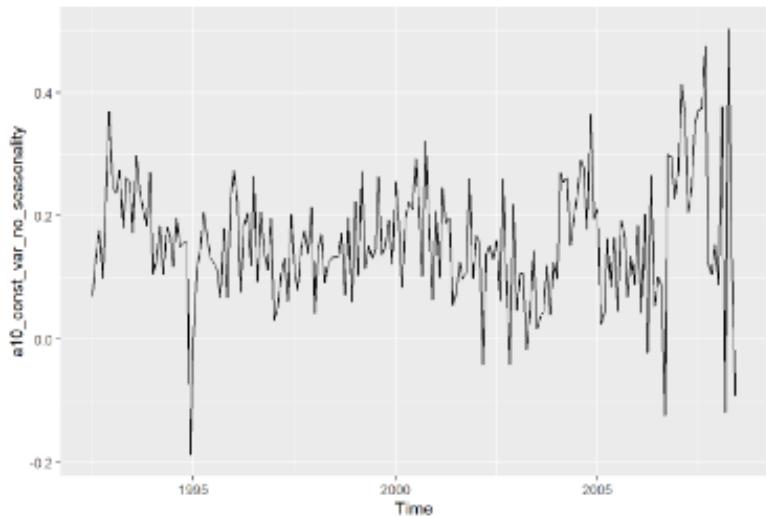


- BoxCox.lambda() will give best value of lambda

```
> BoxCox.lambda(a10)
[1] 0.1313326
```

ARIMA: Stationary Process

```
autoplot(a10_const_var_no_seasonality)
```



Remove Seasonality ←

```
a10_const_var_no_seasonality = diff(x = a10_const_var, lag = 12)  
a10_const_var_no_seasonality
```

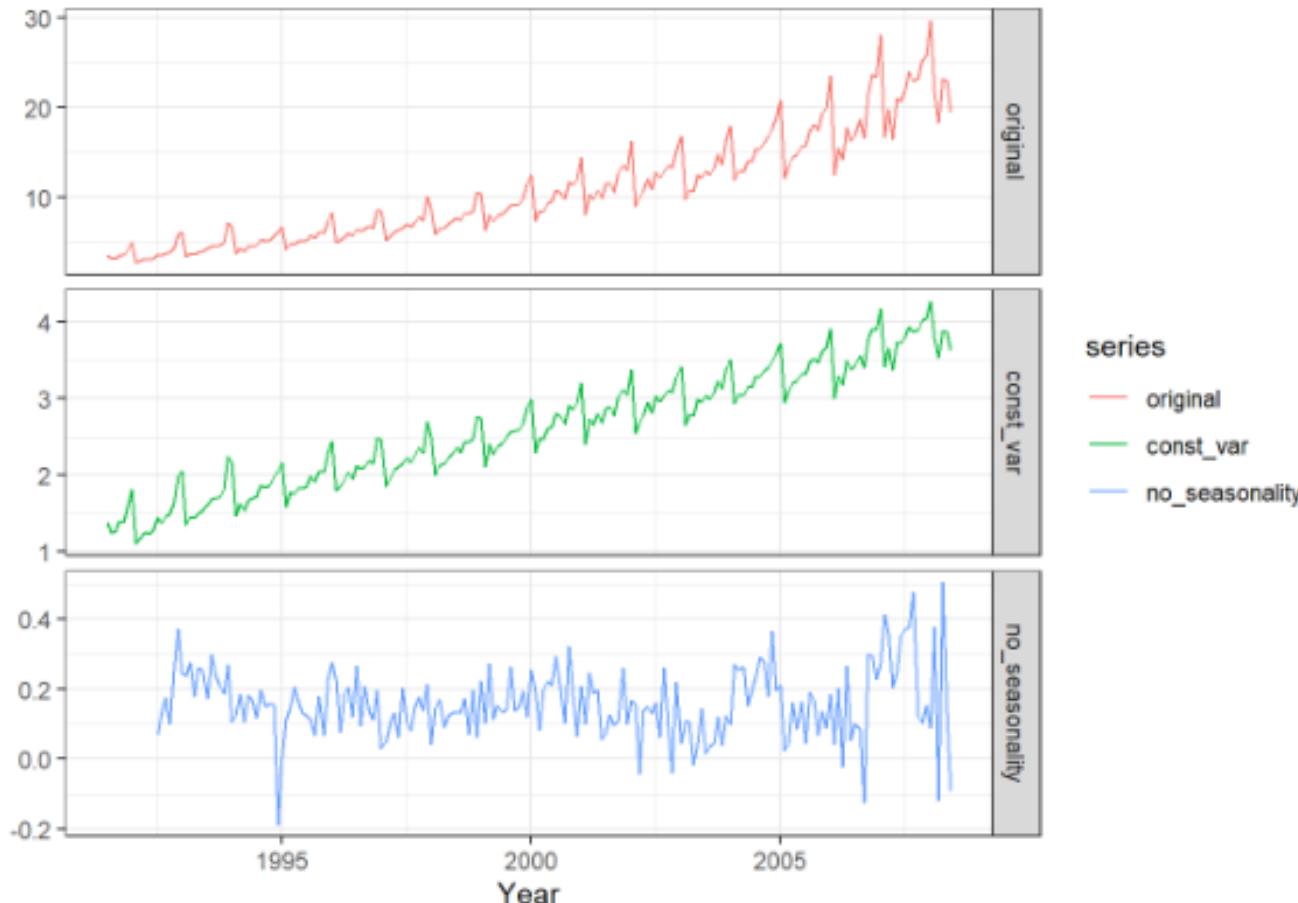
```
##           Jan        Feb        Mar        Apr        May  
## 1992 0.246245444 0.236655903 0.274125392 0.179992291 0.261780298  
## 1993 0.106707317 0.127008025 0.183466224 0.104501404 0.182032776  
## 1994 0.003433011 0.111778361 0.145670228 0.204842989 0.165625047  
## 1995 0.274048279 0.225473639 0.076007964 0.184814821 0.203705337  
## 1996 0.030619086 0.049922844 0.102981627 0.132100224 0.061308706  
## 1997 0.042013647 0.143548717 0.170116222 0.091430829 0.124446414  
## 1998 0.223837120 0.103636140 0.272638470 0.114982618 0.151741933  
## 1999 0.255598717 0.192773618 0.083874135 0.195927270 0.222683128  
## 2000 0.207315424 0.099973328 0.245204433 0.187895960 0.196298560  
## 2001 0.167799083 0.155812045 -0.041441862 0.136976673 0.150926216  
## 2002 0.046097124 0.106401350 0.106482654 -0.017134319 0.045823888  
## 2003 0.098254769 0.269775707 0.254154218 0.262395495 0.151991548  
## 2004 0.211523067 0.024923542 0.043005821 0.162734647 0.084128115  
## 2005 0.183948732 0.043094279 0.203377923 -0.022321033 0.265235151  
## 2006 0.271276841 0.412810508 0.359139697 0.205088323 0.245380862  
## 2007 0.087713881 0.376983677 -0.118270012 0.504017315 0.130702472
```

```
##           Jun        Jul        Aug        Sep        Oct  
## 1992 0.068914872 0.131647730 0.176442705 0.099083250  
## 1993 0.253208354 0.172511836 0.298285240 0.240016156 0.209963892  
## 1994 0.166462225 0.117631639 0.196623325 0.149786828 0.157183802  
## 1995 0.131964341 0.124543290 0.113011767 0.066732755 0.179418490  
## 1996 0.119293319 0.265113644 0.092892385 0.206980636 0.140971432  
## 1997 0.201517801 0.107158949 0.079535272 0.147225882 0.175982731  
## 1998 0.131683621 0.133639241 0.133061334 0.172872087 0.072088989  
## 1999 0.131759208 0.144438879 0.262786380 0.137485989 0.150691433  
## 2000 0.209991223 0.293065044 0.215860117 0.103065352 0.320512593  
## 2001 0.055462118 0.075086949 0.125064642 0.097664934 0.107440890  
## 2002 0.129236467 0.160275193 0.062383518 0.260280947 0.095281699  
## 2003 0.144140785 0.016622688 0.036117330 0.044412227 0.119411497  
## 2004 0.196240086 0.234474574 0.290080889 0.278847563 0.178604272  
## 2005 0.164846565 0.045141607 0.191876814 0.166284639 0.068892416  
## 2006 0.052754167 0.101323581 0.085549371 -0.123863817 0.299287705  
## 2007 0.349622833 0.370839029 0.375135514 0.475241292 0.123414028  
## 2008 -0.092372655  
##           Nov        Dec  
## 1992 0.248131960 0.370137637  
## 1993 0.183326086 0.270377420  
## 1994 0.156439404 -0.187346253  
## 1995 0.067053891 0.229444081  
## 1996 0.112312454 0.195560601  
## 1997 0.137864109 0.213969434  
## 1998 0.197146054 0.060888612  
## 1999 0.193591168 0.120950198  
## 2000 0.180659912 0.063985039  
## 2001 0.260685014 0.098707526  
## 2002 -0.040381829 0.218796465  
## 2003 0.039857509 0.124434429  
## 2004 0.365334465 0.195269583  
## 2005 0.135767731 0.088236836  
## 2006 0.295486720 0.228529224  
## 2007 0.104387656 0.153293541  
## 2008
```

ARIMA: Stationary Process

See the changes

```
dat = cbind(original = a10,
            const_var = a10_const_var,
            no_seasonality = a10_const_var_no_seasonality)
library(ggthemes)
autoplot(dat, facets=T, colour = T)+ylab('')+xlab('Year')+theme_bw()
```



There are automatic functions to estimate the level of differencing required to make a series non-stationary

```
nsdiffs(a10_const_var)
```

```
## [1] 1
```

```
ndiffs(a10_const_var_no_seasonality)
```

```
## [1] 0
```

nsdiffs()

- uses to determine the appropriate number of seasonal differences required to make a given time series stationary

ndiffs()

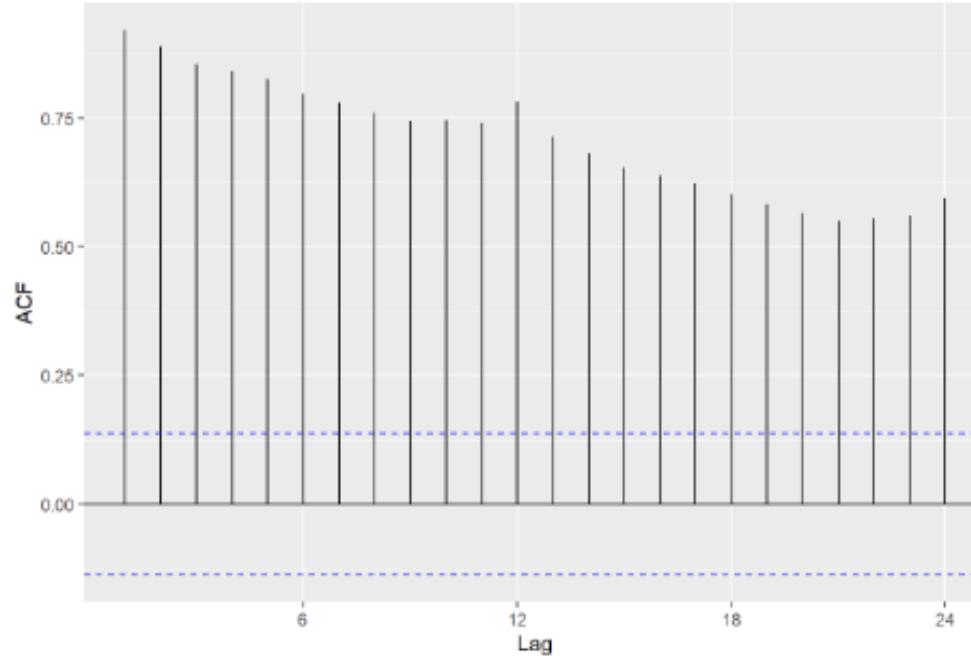
- use to determine the appropriate number of first differences required to make a given time series stationary

ARIMA: Stationary Process: ACF Plot and KPSS Test

Or, one can examine ACF plots to see if the data is stationary

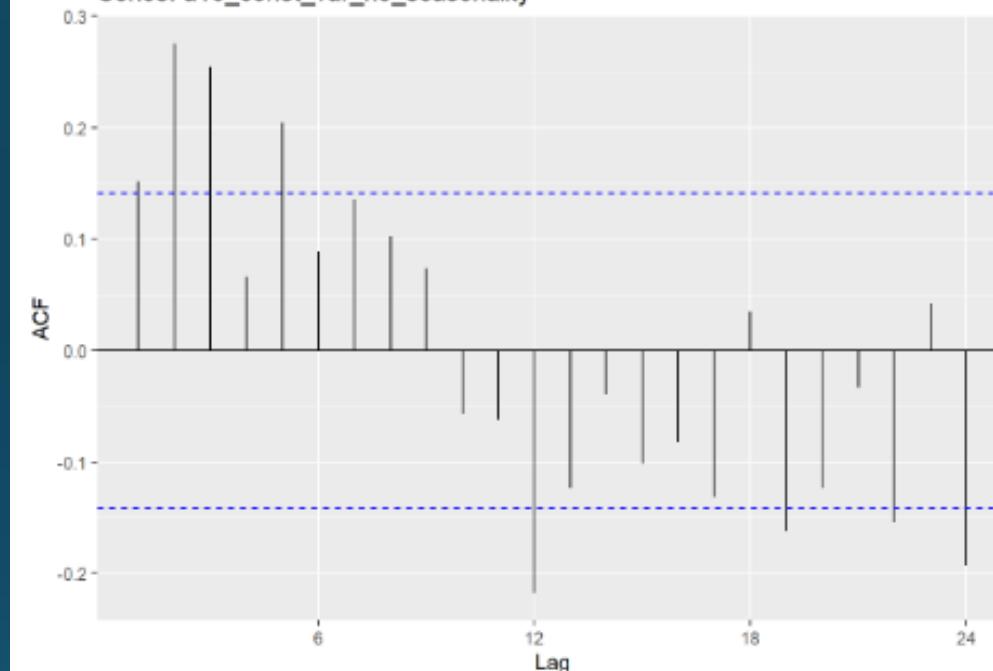
```
ggAcf(a10)
```

Series: a10



```
ggAcf(a10_const_var_no_seasonality)
```

Series: a10_const_var_no_seasonality



KPSS Test for Stationarity Non-significant test statistic indicates the differenced data is stationary.

```
kpss.test(a10_const_var_no_seasonality)
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: a10_const_var_no_seasonality  
## KPSS Level = 0.14853, Truncation lag parameter = 4, p-value = 0.1
```

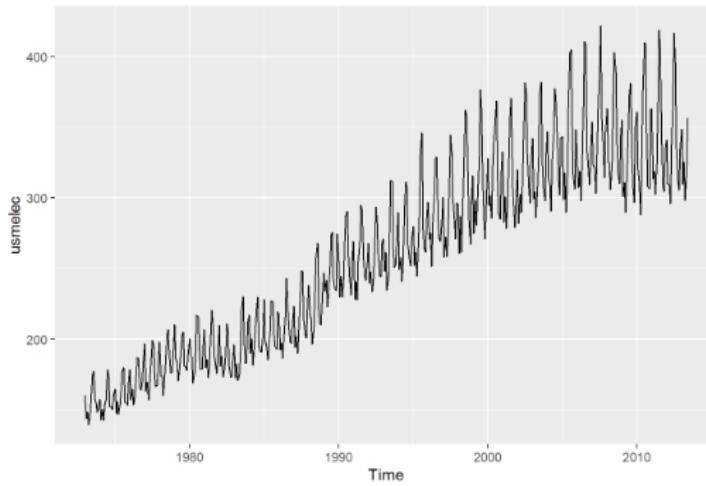
H_0 : data is stationary

Making a Series Stationary: Example 2

Let us use another example to illustrate.

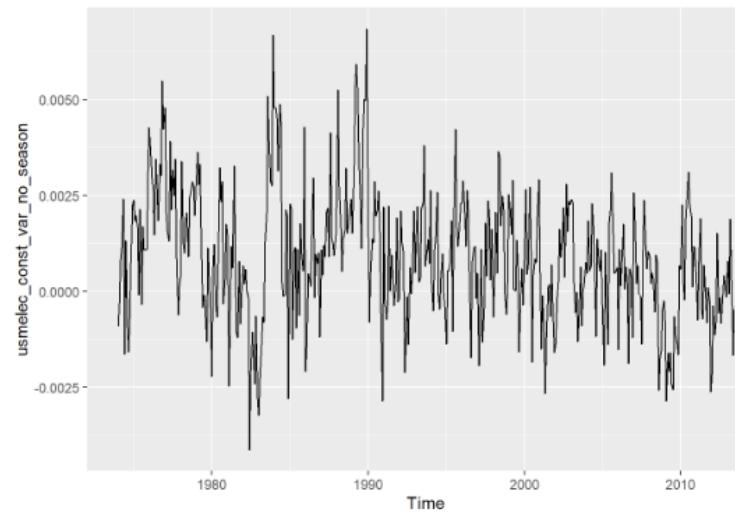
Original Series

```
autoplot(usmlelc)
```



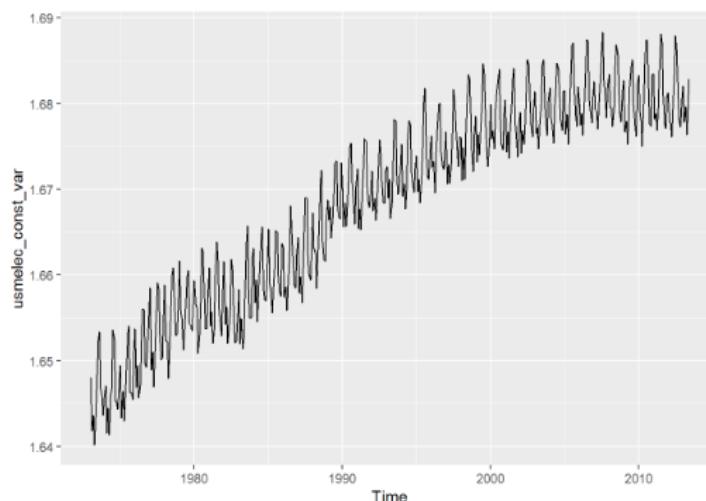
Seasonal Differencing

```
usmlelc_const_var_no_season = diff(usmlelc_const_var, lag = 12)
autoplot(usmlelc_const_var_no_season)
```



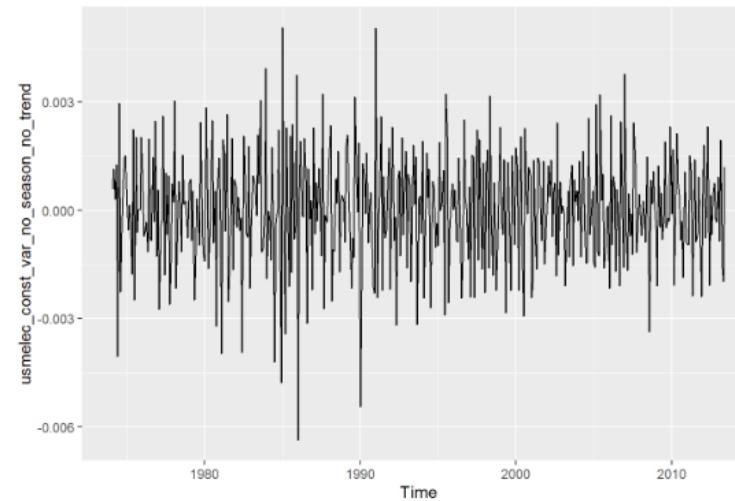
Stabilize Variance

```
usmlelc_const_var = BoxCox(usmlelc, lambda = BoxCox.lambda(usmlelc))
autoplot(usmlelc_const_var)
```



Lag-1 Differencing

```
usmlelc_const_var_no_season_no_trend = diff(usmlelc_const_var_no_season, 1)
autoplot(usmlelc_const_var_no_season_no_trend)
```



usmlelc (fpp2)

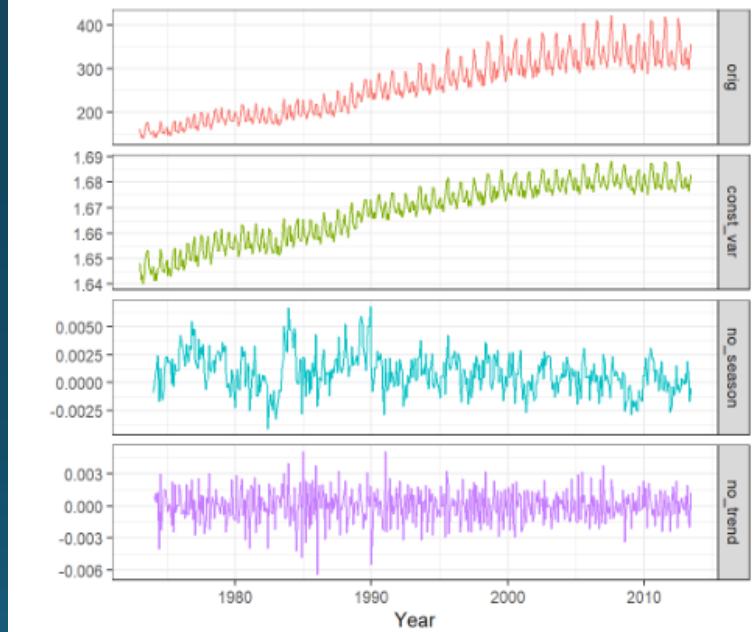
Electricity monthly total net generation. January 1973 - June 2013.

Description

Electricity net generation measured in billions of kilowatt hours (kWh).

See the change

```
dat = cbind(orig = usmlelc,
            const_var = usmlelc_const_var,
            no_season = usmlelc_const_var_no_season,
            no_trend = usmlelc_const_var_no_season_no_trend)
autoplot(dat, colour = T, facets = T)+xlab('Year')+ylab('') +theme_bw()
```



KPSS Test for Stationarity

Non-significant test statistic indicates series is now stationary and differencing is no longer needed.

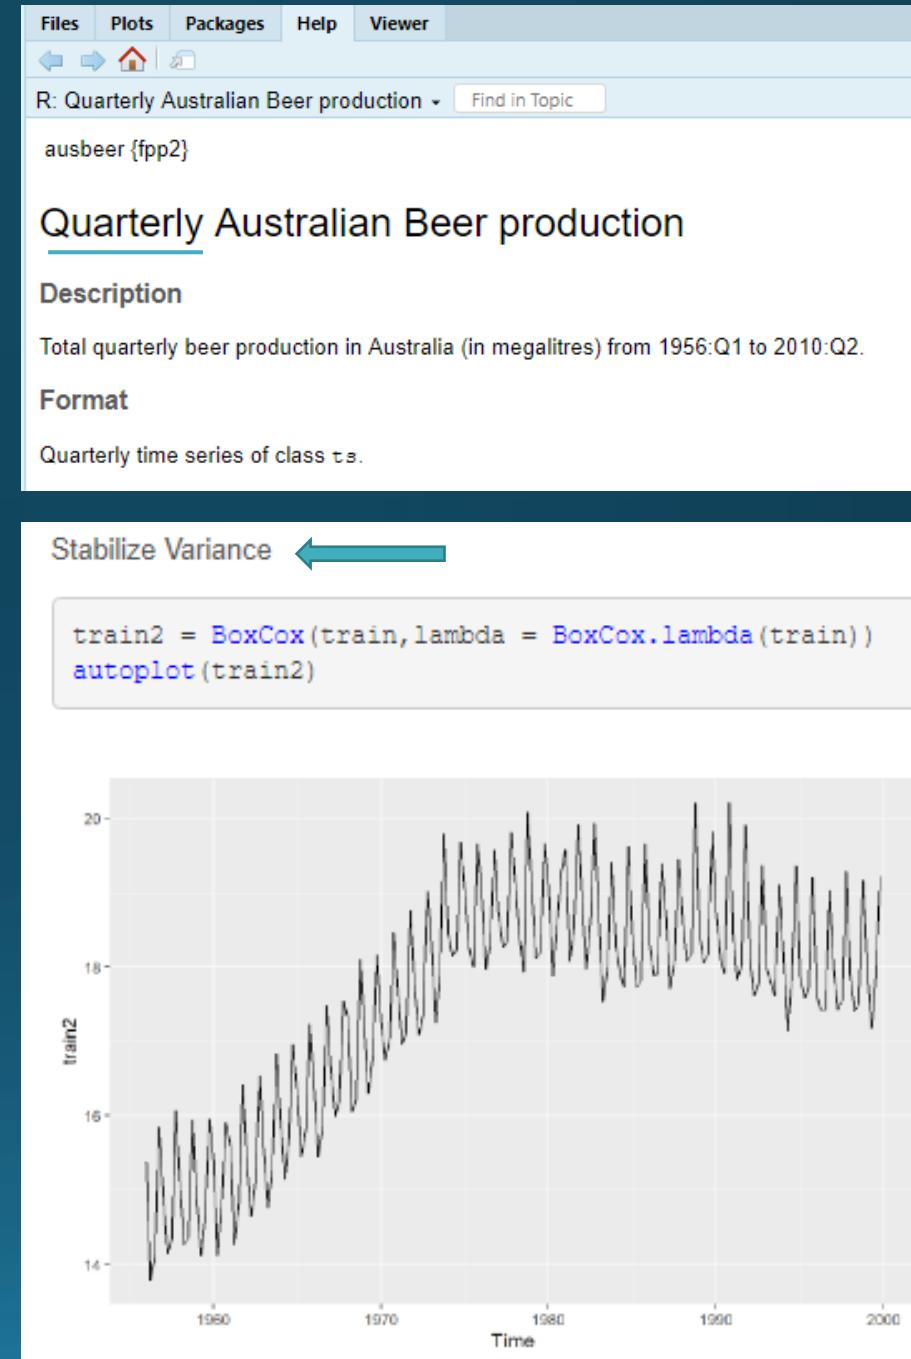
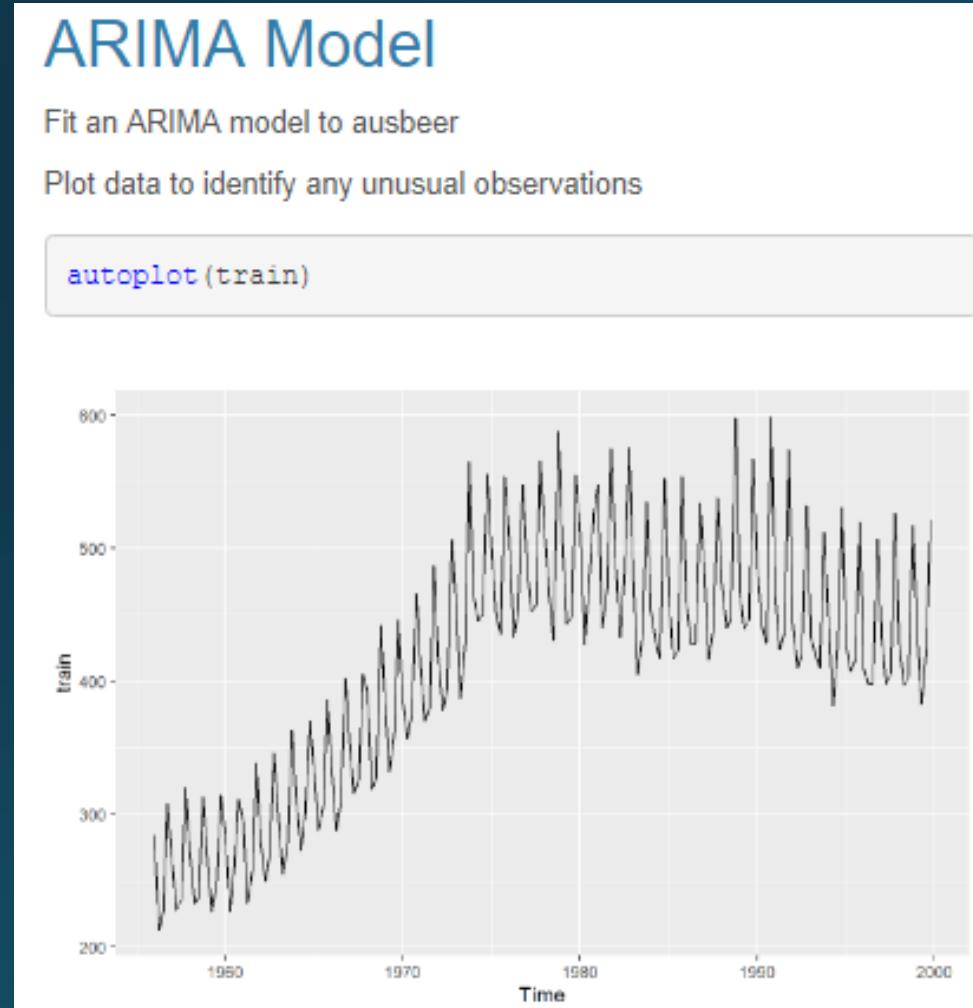
```
kpss.test(usmlelc_const_var_no_season_no_trend)
```

```
##
## KPSS Test for Level Stationarity
##
## data: usmlelc_const_var_no_season_no_trend
## KPSS Level = 0.01669, Truncation lag parameter = 5, p-value = 0.1
```

H₀: data is stationary

ARIMA: Fitting a Model

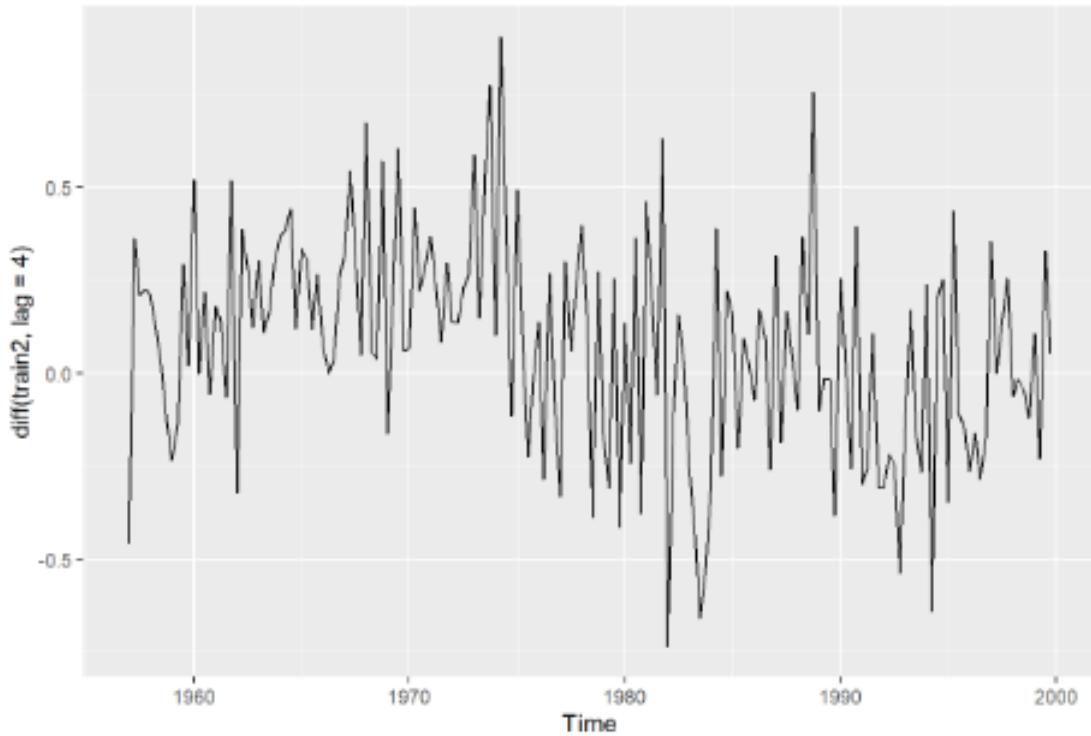
- Using ausbeer data set



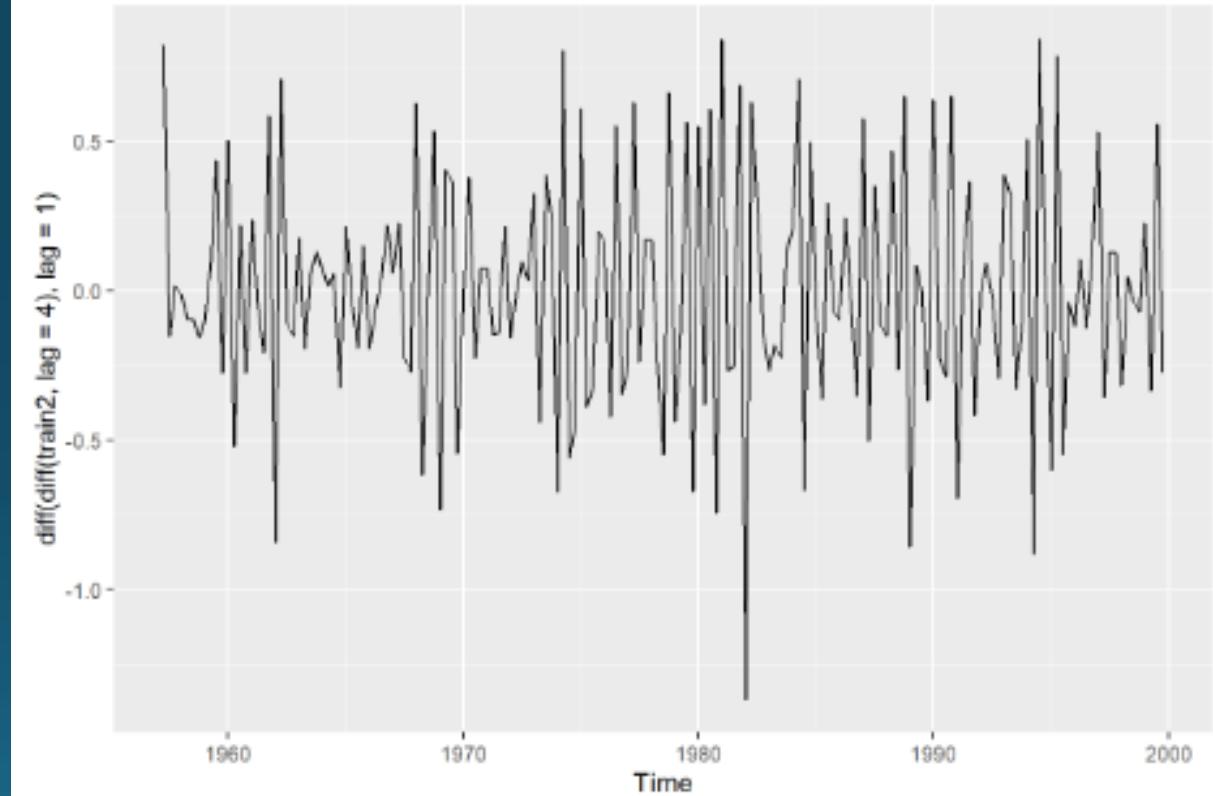
ARIMA: Fitting a Model

If the data are non-stationary, take first differences until the data are stationary

```
autoplot(diff(train2,lag = 4))
```



```
autoplot(diff(diff(train2,lag = 4),lag=1))
```



```
nsdiffs(train)  
  
## [1] 1  
  
nsdiffs(diff(train,lag = 4))  
  
## [1] 0  
  
ndiffs(diff(train,lag = 4))  
  
## [1] 1  
  
ndiffs(diff(diff(train,lag = 4),lag=1))  
  
## [1] 0
```

nsdiffs()

- uses to determine the appropriate number of seasonal differences required to make a given time series stationary

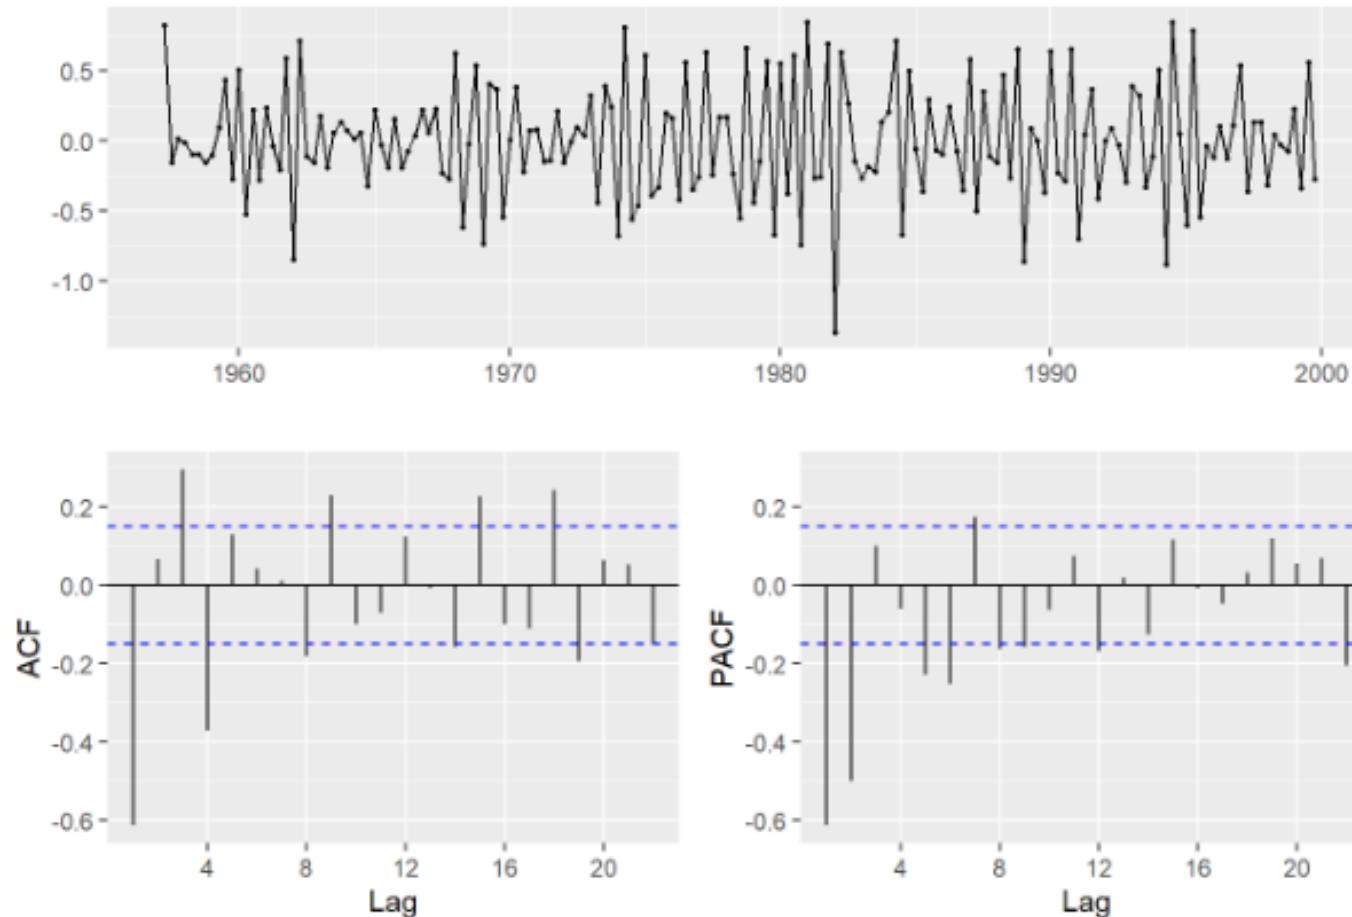
ndiffs()

- use to determine the appropriate number of first differences required to make a given time series stationary

ARIMA: Fitting a Model

Examine ACF and PACF to decide on AR (ARIMA(p,d,0) and/or MA terms (ARIMA(0,d,q))

```
train2 %>%
  diff(lag=4) %>%
  diff(lag=1) %>%
  ggtsdisplay()
```



ARIMA: Fitting a Model

Try chosen model and use AICc to search for a better model

```
modell = Arima(y = train,order = c(0,1,2),seasonal = c(0,1,1),lambda = BoxCox.lambda(train))  
modell  
  
## Series: train  
## ARIMA(0,1,2) (0,1,1) [4]  
## Box Cox transformation: lambda= 0.3102148  
##  
## Coefficients:  
##          ma1      ma2     sma1  
##        -0.9665  0.4085 -0.8268  
## s.e.    0.0713  0.0729  0.0499  
##  
## sigma^2 estimated as 0.05151: log likelihood=9.93  
## AIC=-11.86  AICc=-11.62  BIC=0.71
```

- The model is ARIMA (0,1,2)(0,1,1)
 - p = 0 indicates that no ordinary AR lag is being used.
 - d = 1 indicates 1 ordinary difference is being used.
 - q = 2 indicates 2 ordinary MA lags are being used.
 - P = 0 indicates that no seasonal AR lag is being used
 - D = 1 indicates 1 seasonal difference is being used.
 - Q = 1 indicates 1 seasonal MA lag is being used.

Seasonal ARIMA

- Non-Seasonal ARIMA + seasonal terms = Seasonal ARIMA

- ARIMA(p,d,q)(P,D,Q)m

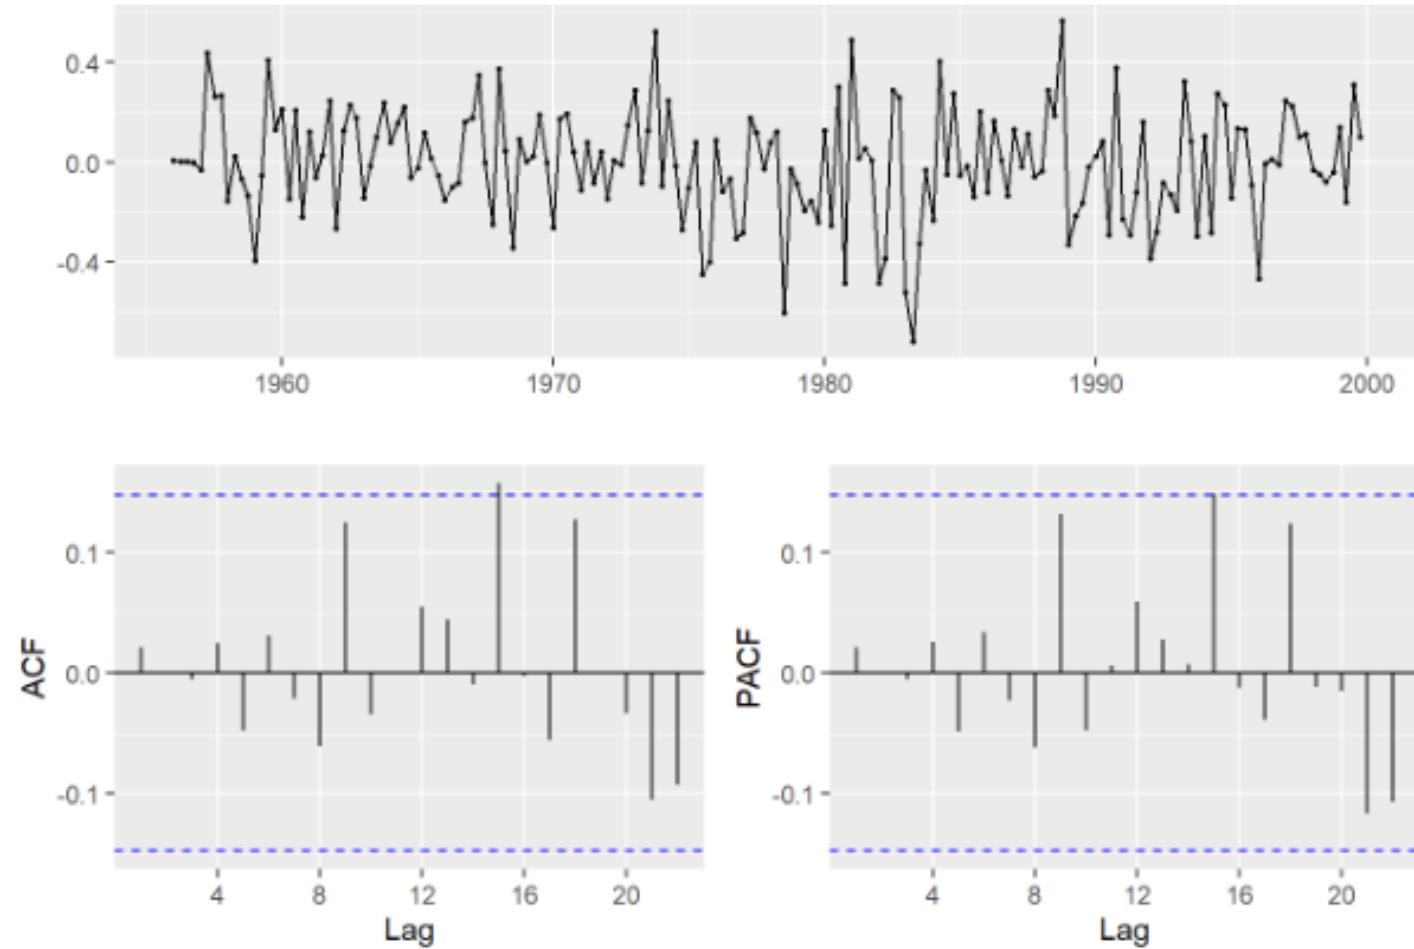
- (P,D,Q) is seasonal component
- m is seasonal period

where

- p = Number of ordinary Autoregressive lags
- d = Number of ordinary differences
- q = Number of ordinary Moving average lags
- P = Number of seasonal Autoregressive lags
- D = Number of seasonal differences
- Q = Number of seasonal Moving Average lags
- m = Seasonal period or number of observations per year

ARIMA: Fitting a Model

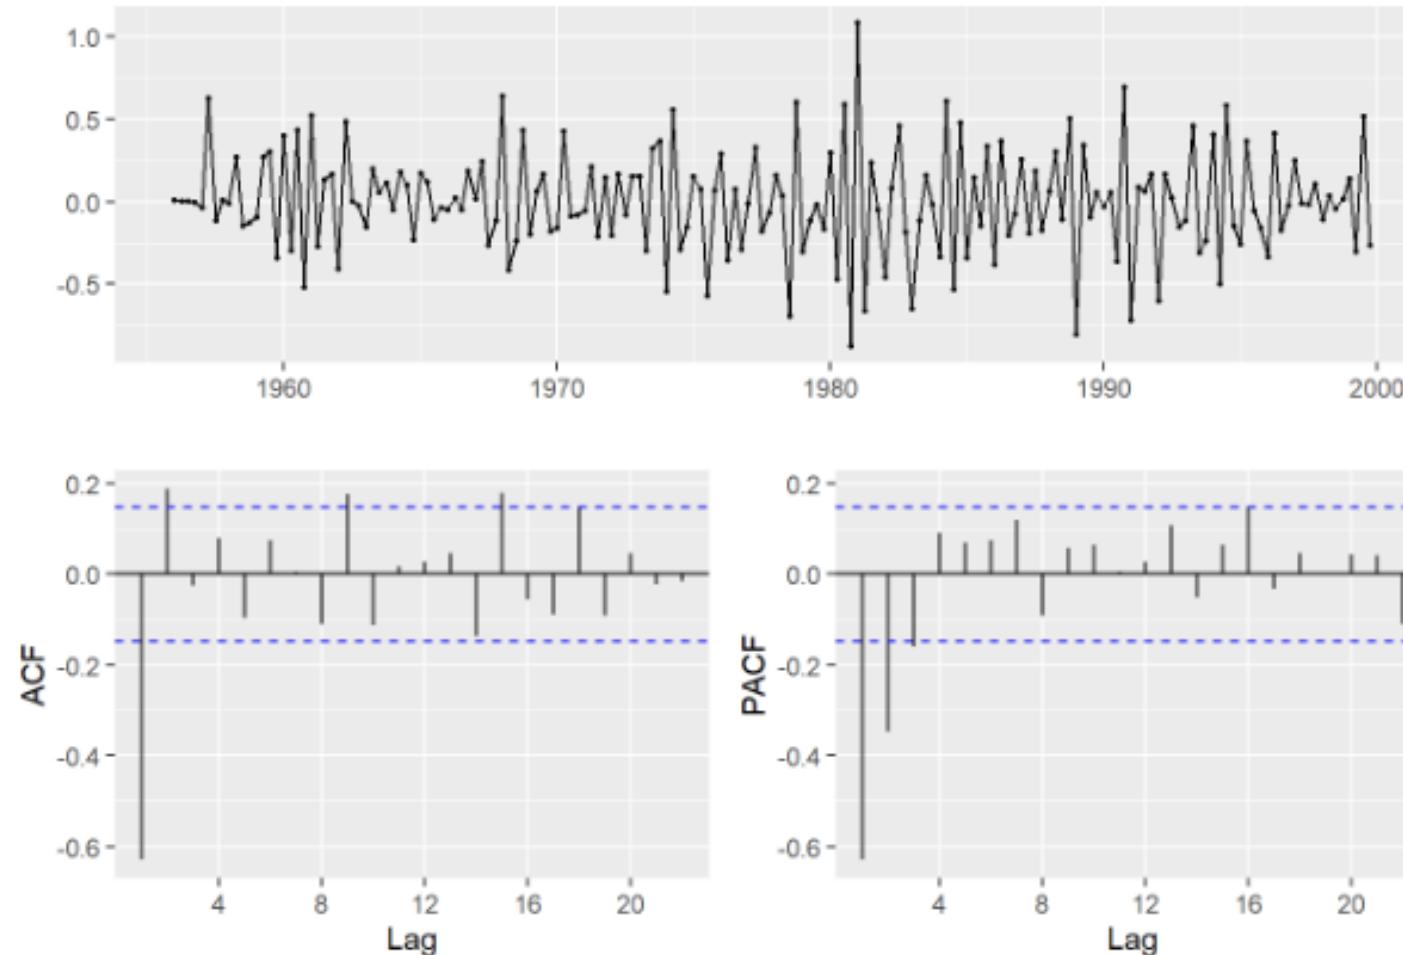
```
ggtsdisplay(residuals(modell))
```



ARIMA: Fitting a Model

Trying out an alternative model ←

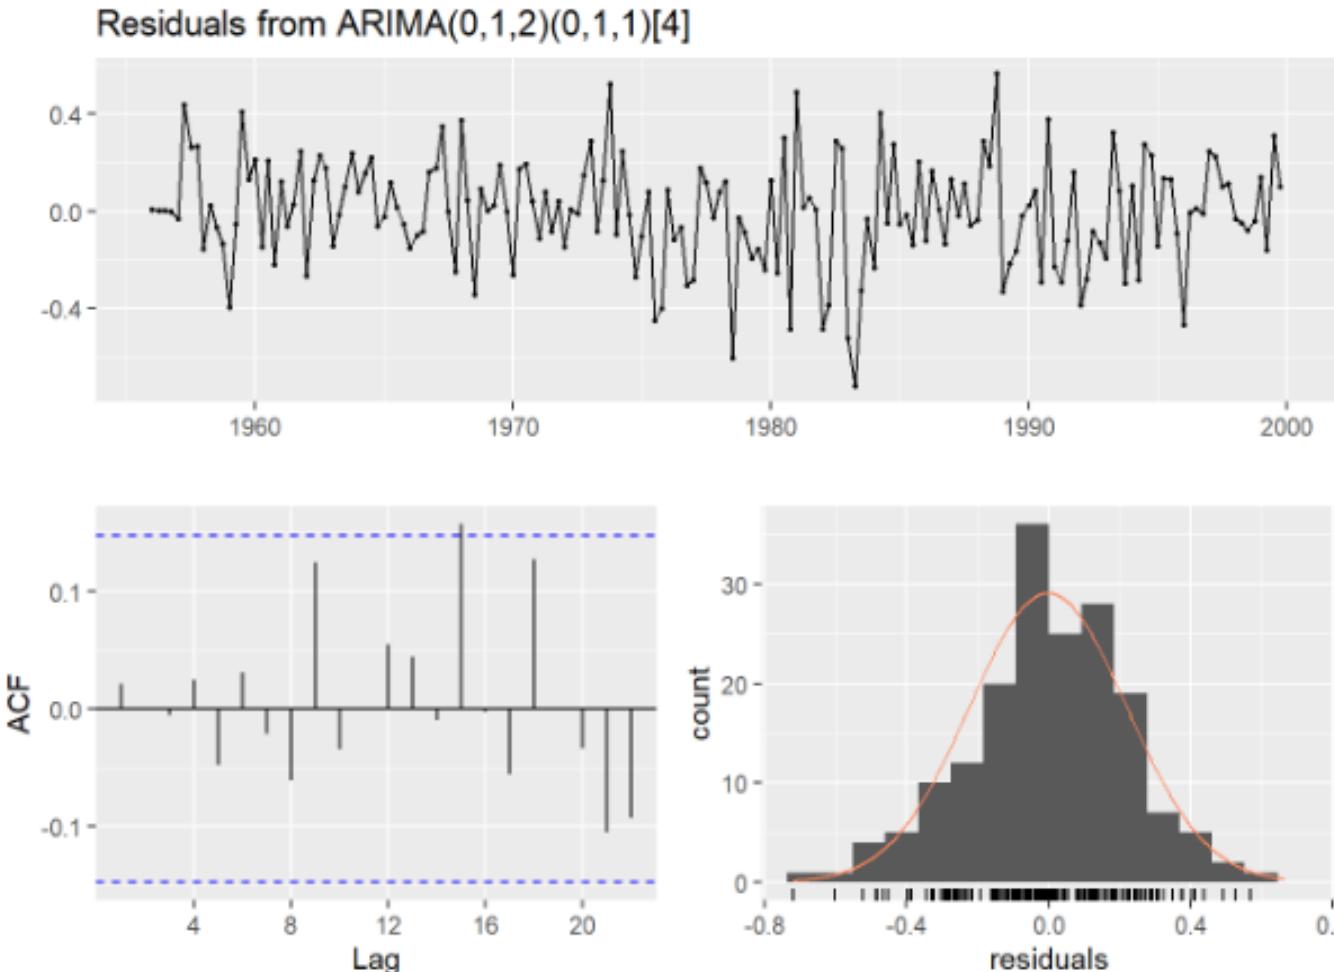
```
model2 = Arima(y = train,order = c(0,1,0),seasonal = c(0,1,1),lambda = BoxCox.lambda(train))  
ggtsddisplay(residuals(model2))
```



ARIMA: Fitting a Model – Check Residuals of Model1

Check residuals by plotting. If they do not look like white noise, try a modified model

```
checkresiduals(model1)
```



H_0 : the autocorrelations up to lag k are all 0
 H_A : the autocorrelations of one or more lags differ from 0.

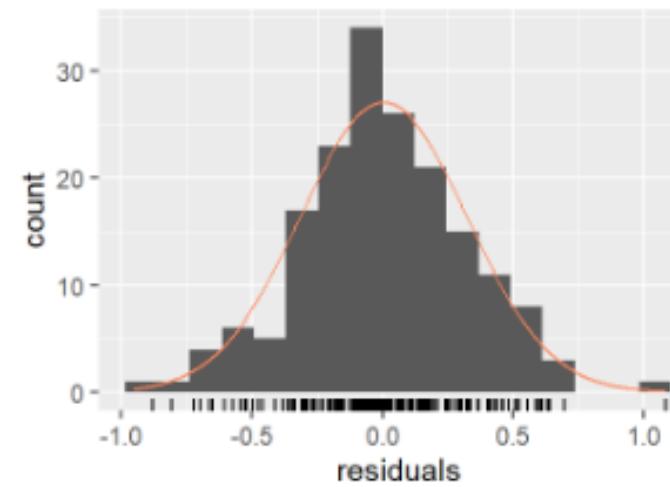
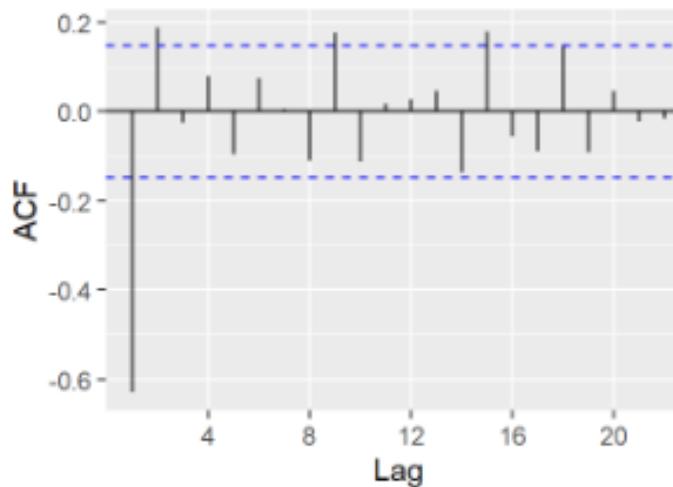
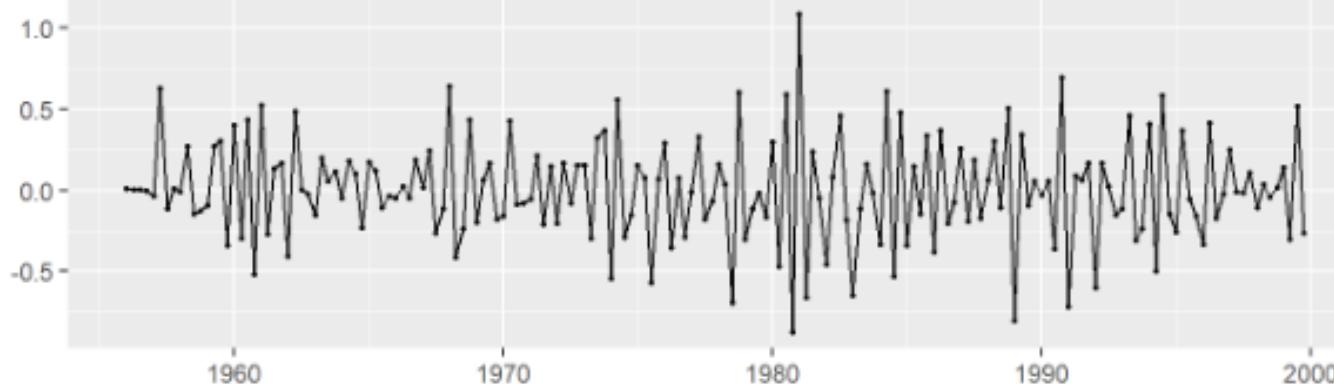
```
##  
## Ljung-Box test  
##  
## data: Residuals from ARIMA(0,1,2) (0,1,1) [4]  
## Q* = 1.5798, df = 5, p-value = 0.9037  
##  
## Model df: 3. Total lags used: 8
```

Residuals resemble white noise

ARIMA: Fitting a Model – Check Residuals of Model2

```
checkresiduals(model2)
```

Residuals from ARIMA(0,1,0)(0,1,1)[4]

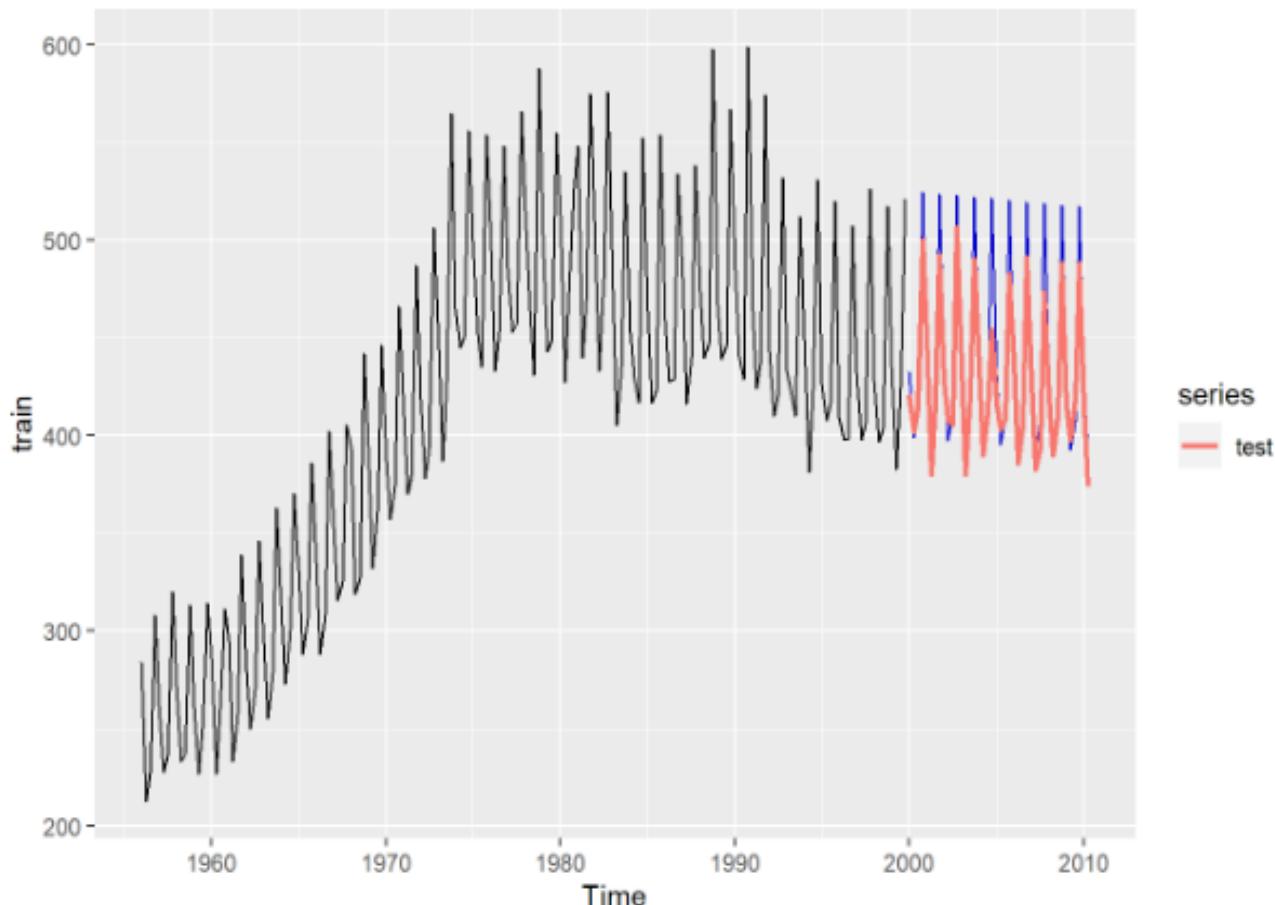


```
##  
## Ljung-Box test  
##  
## data: Residuals from ARIMA(0,1,0) (0,1,1) [4]  
## Q* = 83.58, df = 7, p-value = 2.554e-15  
##  
## Model df: 1. Total lags used: 8
```

Once the residuals look like white noise, calculate forecasts

```
autoplot(forecast(modell,h = 42),PI=F) +
  autolayer(test,size=1)
```

Forecasts from ARIMA(0,1,2)(0,1,1)[4]



- Use model 1 to construct a forecast for Australian beer over the next 42 months of the test sample.

```
modell_forecast = forecast(modell,h=42)
```

```
accuracy(f = modell_forecast,x = ausbeer)
```

	ME	RMSE	MAE	MPE	MAPE	MASE
## Training set	-0.4443474	14.79078	11.11090	-0.06714964	2.659077	0.6956951
## Test set	-10.3114618	19.88872	14.49978	-2.26976852	3.261576	0.9078855
## ACF1 Theil's U						
## Training set	-0.004102863		NA			
## Test set	-0.269479034	0.3887556				

ARIMA - Automatic Model Selection

Use `auto.arima` to pick the best model based on AICc. Sometimes `auto.arima()` does not yield an optimal solution as it uses computational shortcuts. By setting `stepwise` and `approximation` to `False`, we will ensure a more extensive search.

As it turns out, the automatic process picked the model we manually configured.

```
model_auto = auto.arima(y = train,d = 1,D = 1,stepwise = F,approximation = F)
model_auto
```

```
## Series: train
## ARIMA(0,1,2)(0,1,1)[4]
##
## Coefficients:
##          ma1     ma2     sma1
##        -0.9758  0.3959  -0.7659
## s.e.    0.0722  0.0762   0.0520
##
## sigma^2 estimated as 244.5:  log likelihood=-713.37
## AIC=1434.74  AICc=1434.98  BIC=1447.3
```

Comparing Forecasting Models

Let us compare all the models in terms of accuracy metrics on the test sample.

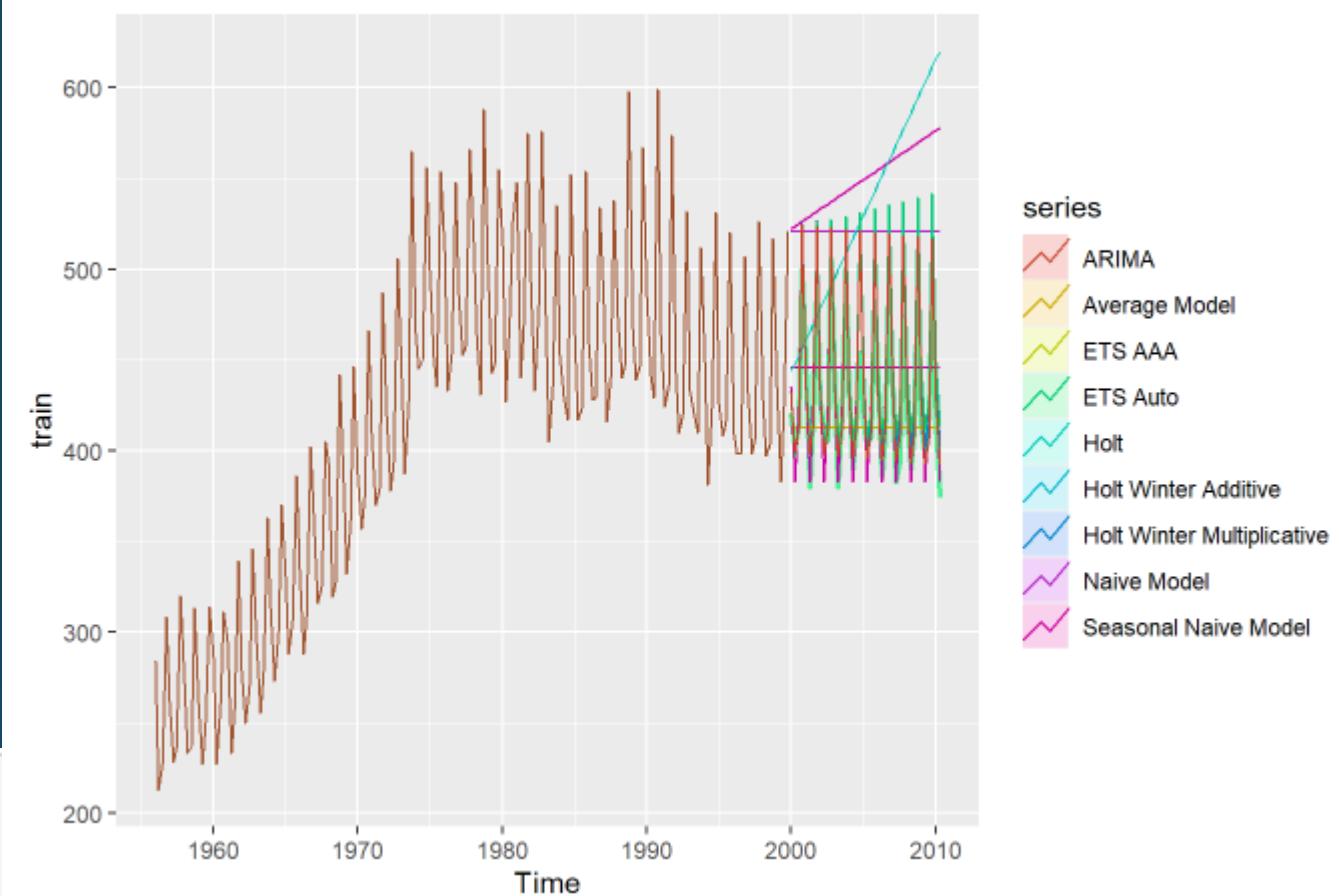
```
rbind(average_model = accuracy(f = average_model,x = ausbeer)[2,],
naive_model = accuracy(f = naive_model,x = ausbeer)[2,],
seasonal_naive_model = accuracy(f = seasonal_naive_model,x = ausbeer)[2,],
drift_model = accuracy(f = drift_model,x = ausbeer)[2,],
ses_model = accuracy(f = ses_model,x = ausbeer)[2,],
holt_model = accuracy(f = holt_model,x = ausbeer)[2,],
holt_damped_model = accuracy(f = holt_damped_model,x = ausbeer)[2,],
hw_additive_model = accuracy(f = hw_additive,x = ausbeer)[2,],
hw_multiplicative = accuracy(f = hw_multiplicative,x = ausbeer)[2,],
ets_aaa = accuracy(ets_aaa_forecast,x = ausbeer)[2,],
ets_auto = accuracy(ets_auto_forecast,x = ausbeer)[2,],
arima = accuracy(arima_forecast,x=ausbeer)[2,]
)
```

	ME	RMSE	MAE	MPE	MAPE
## average_model	15.25866	40.10097	28.70346	2.876448	6.348359
## naive_model	-93.30952	100.40881	93.30952	-22.690269	22.690269
## seasonal_naive_model	-11.54762	21.85286	17.02381	-2.521773	3.863916
## drift_model	-122.42667	129.44954	122.42667	-29.581531	29.581531
## ses_model	-18.09737	41.26471	37.68469	-4.978561	8.973630
## holt_model	-104.37559	124.04137	108.00085	-25.405766	26.131936
## holt_damped_model	-32.18055	49.84832	44.97675	-8.303939	10.897279
## hw_additive_model	-14.05834	22.67754	16.71429	-3.151230	3.775641
## hw_multiplicative	-16.47439	24.38794	18.20496	-3.721500	4.125710
## ets_aaa	-20.38991	28.22830	22.12902	-4.654819	5.060769
## ets_auto	-20.48601	28.61288	22.23155	-4.669352	5.076505
## arima	-10.31146	19.88872	14.49978	-2.269769	3.261576
		MASE		ACF1	Theil's U
## average_model	1.7972318	-0.052981076	0.7832398		
## naive_model	5.8424602	-0.052981076	1.8984090		
## seasonal_naive_model	1.0659247	0.029098002	0.4278746		
## drift_model	7.6655940	0.152038516	2.4589272		
## ses_model	2.3595800	-0.052981076	0.7972014		
## holt_model	6.7623392	0.611282124	2.3753720		
## holt_damped_model	2.8161636	-0.005103579	0.9593471		
## hw_additive_model	1.0465444	-0.185710461	0.4430322		
## hw_multiplicative	1.1398809	-0.169487536	0.4757030		
## ets_aaa	1.3855813	0.015455576	0.5492527		
## ets_auto	1.3920011	0.010216303	0.5569420		
## arima	0.9078855	-0.269479034	0.3887556		

```

autoplot(train, color='sienna')+
  autolayer(test, size=1.05,color='seagreen2')+
  autolayer(average_model,series = 'Average Model',PI=F)+
  autolayer(naive_model,series = 'Naive Model',PI=F)+
  autolayer(seasonal_naive_model,series = 'Seasonal Naive Model',PI=F)+
  autolayer(drift_model,series = 'Seasonal Naive Model',PI=F)+
  autolayer(ses_model,series = 'Seasonal Naive Model',PI=F)+
  autolayer(holt_model,series = 'Holt',PI=F)+
  autolayer(hw_additive,series = 'Holt Winter Additive',PI=F)+
  autolayer(hw_multiplicative,series = 'Holt Winter Multiplicative',PI=F)+
  autolayer(ets_aaa_forecast,series = 'ETS AAA',PI=F)+
  autolayer(ets_auto_forecast,series = 'ETS Auto',PI=F)+
  autolayer(modell_forecast,series = 'ARIMA',PI=F)

```



- In this module we,
 - Examined, explored, and visualized time-series data
 - Reviewed basic/simple forecasting methods
 - Explored exponential smoothing
 - Examined ARIMA models
 - Constructed forecasts using simple methods, exponential smoothing, and ARIMA.